

The Linux OS - Security and Hardening Overview for Developers

***Focus on Buffer Overflow
With a PoC on the ARM Processor***

***kaiwanTECH
<http://kaiwantech.in>***

Linux OS – Security and Hardening



- **Agenda**

- Basic Terminology
- Current State
 - Linux kernel vulnerability stats
 - “Security Vulnerabilities in Modern OS’s” - a few slides
- Tech Preliminary: the process Stack
- BOF Vulnerabilities
 - What is BOF
 - Why is it dangerous?
 - [Demo: a PoC on the ARM processor]



Linux OS – Security and Hardening



- **Agenda (contd.)**

- Modern OS Hardening Countermeasures
 - Using Managed programming languages
 - Compiler protection
 - Libraries
 - Executable space protection
 - ASLR
 - Better Testing
- Concluding Remarks
- Q&A



Basic Terminology



Source - Wikipedia

Vulnerability

In computer security, a vulnerability is a weakness which allows an attacker to reduce a system's information assurance.

Vulnerability is the intersection of three elements: a system susceptibility or flaw, attacker access to the flaw, and attacker capability to exploit the flaw.

A **software vulnerability** is a security flaw, glitch, or weakness found in software or in an operating system (OS) that can lead to security concerns. An example of a software flaw is a buffer overflow.

Basic Terminology

(contd.)



Source - Wikipedia

Exploit

In computing, an exploit is an attack on a computer system, especially one that takes advantage of a particular vulnerability that the system offers to intruders.

Used as a verb, the term refers to the act of successfully making such an attack.

Basic Terminology

(contd.)



Source: CVEdetails

What is an "**Exposure**"?

An information security exposure is a mistake in software that allows access to information or capabilities that can be used by a hacker as a stepping-stone into a system or network.

Aka 'info-leak'.

While the world is now kind of (sadly) used to software vulns and exposures, what about the same but at the hardware level! Recent news stories have the infosec community in quite a tizzy.

-

"The Big Hack: How China Used a Tiny Chip to Infiltrate U.S. Companies", Bloomberg, 04 Oct 2018.

- Side-Channel Attacks & the Importance of Hardware-Based Security, July 2018

Basic Terminology

(contd.)



- **What is CVE?**

[Source]

- *“**Common Vulnerabilities and Exposures** (CVE®) is a dictionary of common names (i.e., CVE Identifiers) for publicly known cybersecurity vulnerabilities. CVE's common identifiers make it easier to share data across separate network security databases and tools, and provide a baseline for evaluating the coverage of an organization's security tools. ...”*
- CVE is
 - One name for one vulnerability or exposure
 - One standardized description for each vulnerability or exposure
 - A dictionary rather than a database
 - How disparate databases and tools can "speak" the same language
 - The way to interoperability and better security coverage
 - A basis for evaluation among tools and databases
 - Free for public download and use
 - Industry-endorsed via the CVE Numbering Authorities, CVE Board, and CVE-Compatible Products

Basic Terminology

(contd.)



What is a CVE Identifier?

CVE Identifiers (also called "CVE names," "CVE numbers," CVE-IDs," and "CVEs") are unique, common identifiers for publicly known information security vulnerabilities.

Each CVE Identifier includes the following:

- CVE identifier number (i.e., "CVE-2014-0160").
- Indication of "entry" or "candidate" status.
- Brief description of the security vulnerability or exposure.
- Any pertinent references (i.e., vulnerability reports and advisories or OVAL-ID).
- CVE Identifiers are used by information security product/service vendors and researchers as a **standard method for identifying vulnerabilities** and for cross-linking with other repositories that also use CVE Identifiers.

The CVEDetails website provides valuable information and a scoring system

CVE FAQs Page

Basic Terminology

(contd.)



CVE Identifier – Old and New Syntax

- *New system, practically
from Jan 2015*

CVE-ID Syntax Change

Old Syntax

CVE-YYYY-NNNN

4 fixed digits, supports
a maximum of 9,999
unique identifiers per
year.

Fixed 4-Digit Examples

CVE-1999-0067
CVE-2005-4873
CVE-2012-0158

New Syntax

CVE-YYYY-NNNN...N

4-digit minimum and no
maximum, provides for
additional capacity each
year when needed.

Arbitrary Digits Examples

CVE-2014-0001
CVE-2014-12345
CVE-2014-7654321

YYYY indicates year the ID is issued to a
CVE Numbering Authority (CNA) or published.

Implementation date: January 1, 2014

Source: <http://cve.mitre.org>

Basic Terminology

(contd.)



A CVE Example

- [CVE-2014-0160](#) *[aka “Heartbleed”]*
- **Description:**
 - The (1) TLS and (2) DTLS implementations in OpenSSL 1.0.1 before 1.0.1g do not properly handle Heartbeat Extension packets, which allows remote attackers to obtain sensitive information from process memory via crafted packets that trigger a buffer over-read, as demonstrated by reading private keys, related to `d1_both.c` and `t1_lib.c`, aka the Heartbleed bug.



Common Vulnerabilities and Exposures

*The Standard for Information Security
Vulnerability Names*

[Home](#) | [CVE IDs](#) | [About CVE](#) | [Compatible Products & More](#) | [Community](#) | [Blog](#) | [News](#) | [Site Search](#)

TOTAL CVE IDs: 82281

[HOME](#) > [CVE](#) > [CVE-2014-0160](#)

Section Menu

CVE IDs

[CVEnew Twitter Feed](#)
[Other Updates & Feeds](#)

Request a CVE ID

[Contact a CVE Numbering Authority \(CNA\)](#)
[Contact Primary CNA \(MITRE\) – CVE Request web form](#)
[Reservation Guidelines](#)

CVE LIST (all existing CVE IDs)

[Downloads](#)
[Search CVE List](#)
[Search Tips](#)
[View Entire CVE List \(html\)](#)
[Reference Key/Maps](#)

NVD Advanced CVE Search

[Printer-Friendly View](#)

CVE-ID

CVE-2014-0160

[Learn more at National Vulnerability Database \(NVD\)](#)

• Severity Rating • Fix Information • Vulnerable Software Versions •

Description

The (1) TLS and (2) DTLS implementations in OpenSSL 1.0.1 before 1.0.1g do not properly handle information from process memory via crafted packets that trigger a buffer over-read, as demonstrated by a proof-of-concept exploit.

References

Note: [References](#) are provided for the convenience of the reader to help distinguish between vulnerable

- BUGTRAQ:20141205 NEW: VMSA-2014-0012 - VMware vSphere product updates address
- [URL:http://www.securityfocus.com/archive/1/archive/1/534161/100/0/threaded](http://www.securityfocus.com/archive/1/archive/1/534161/100/0/threaded)
- EXPLOIT-DB:32745
- [URL:http://www.exploit-db.com/exploits/32745](http://www.exploit-db.com/exploits/32745)
- EXPLOIT-DB:32764
- [URL:http://www.exploit-db.com/exploits/32764](http://www.exploit-db.com/exploits/32764)
- FULLDISC:20140408 Re: heartbleed OpenSSL bug CVE-2014-0160

Basic Terminology

(contd.)



Most software security vulnerabilities fall into one of a small set of categories:

- ***buffer overflows***
- ***unvalidated input***
- ***race conditions***
- ***access-control problems***
- ***weaknesses in authentication, authorization, or cryptographic practices***

[Source](#)

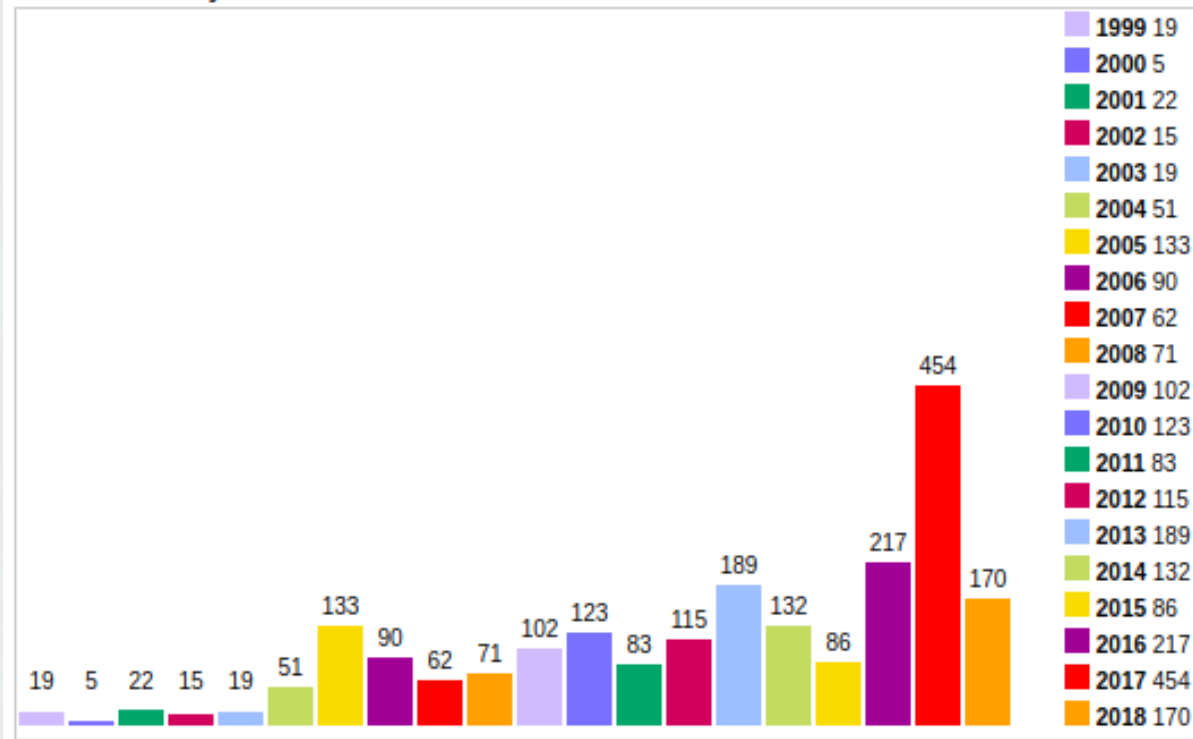
CWE - Common Weakness Enumeration - Types of Exploits

Related Activities		
<ul style="list-style-type: none">The Software Assurance Metrics and Tool Evaluation (SAMATE) Project, NIST.		
NVD CWE Slice		
Name	CWE-ID	Description
Access of Uninitialized Pointer	CWE-824	The program accesses or uses a pointer that has not been initialized.
Algorithmic Complexity	CWE-407	An algorithm in a product has an inefficient worst-case computational complexity that may be detrimental to system performance and can be triggered by an attacker, typically using crafted manipulations that ensure that the worst case is being reached.
Allocation of File Descriptors or Handles Without Limits or Throttling	CWE-774	The software allocates file descriptors or handles on behalf of an actor without imposing any restrictions on how many descriptors can be allocated, in violation of the intended security policy for that actor.
Argument Injection or Modification	CWE-88	The software does not sufficiently delimit the arguments being passed to a component in another control sphere, allowing alternate arguments to be provided, leading to potentially security-relevant changes.
Asymmetric Resource Consumption (Amplification)	CWE-405	Software that does not appropriately monitor or control resource consumption can lead to adverse system performance.
Authentication Issues	CWE-287	When an actor claims to have a given identity, the software does not prove or insufficiently proves that the claim is correct.
Buffer Errors	CWE-119	The software performs operations on a memory buffer, but it can read from or write to a memory location that is outside of the intended boundary of the buffer.

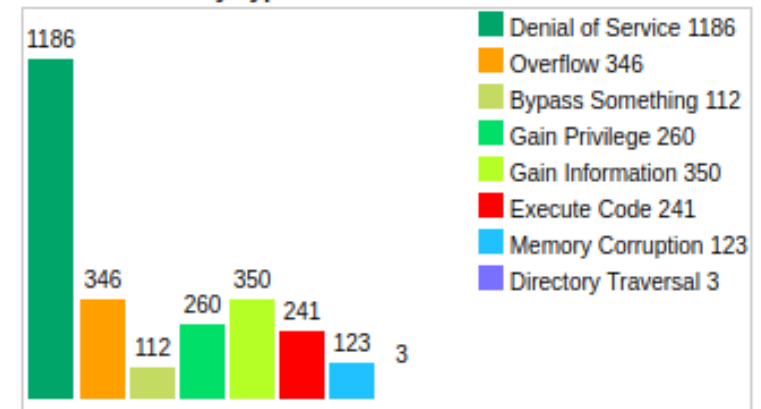
Linux kernel - Vulnerability Stats

Source (CVEdetails)
(1999 to 2018)

Vulnerabilities By Year



Vulnerabilities By Type



"A bunch of links related to Linux kernel exploitation"

Security Vulnerabilities in Modern Operating Systems



***By Cisco, Canada, April 2014.
All rights with Cisco.***

**“The Common Exposures and Vulnerabilities database has over 25 years of data on vulnerabilities in it. In this deck we dig through that database and use it to map out trends and general information on vulnerabilities in software in the last quarter century. For more information please visit our website:
<http://www.cisco.com/web/CA/index.html> ”**

Source: *Security Vulnerabilities in Modern Operating Systems*, Cisco, Apr 2014


OS Security Vulnerabilities



- A look at more than 25 years of past vulnerabilities
 - Based on the CVE/NVD data.
 - CVE started in 1999, but includes historical data going back to 1988.
 - NVD hosts all CVE information in addition to some extra data about vulnerability types, etc.
 - Based on Sourcefire report: <http://www.sourcefire.com/25yearsofvulns>
- Updated (with data from 2013, 2014) and data from other sources

Source: *Security Vulnerabilities in Modern Operating Systems*, Cisco, Apr 2014

Total Critical Vulnerabilities




Source: [Security Vulnerabilities in Modern Operating Systems](#), Cisco, Apr 2014

Vulnerabilities by Type

- Common Weakness Enumeration creates a number of categories for vulnerabilities
- NVD uses a subset of CWE to categorize vulnerabilities:
 - Authentication issues: not properly authenticating users
 - Credentials management: password/credential storage/transmission issues
 - Access Control: permission errors, privilege errors, etc.
 - Buffer error: buffer overflows, etc.
 - CSRF: cross-site request forgery
 - XSS: cross site scripting

Source: [Security Vulnerabilities in Modern Operating Systems](#), Cisco, Apr 2014

Vulnerabilities by Type



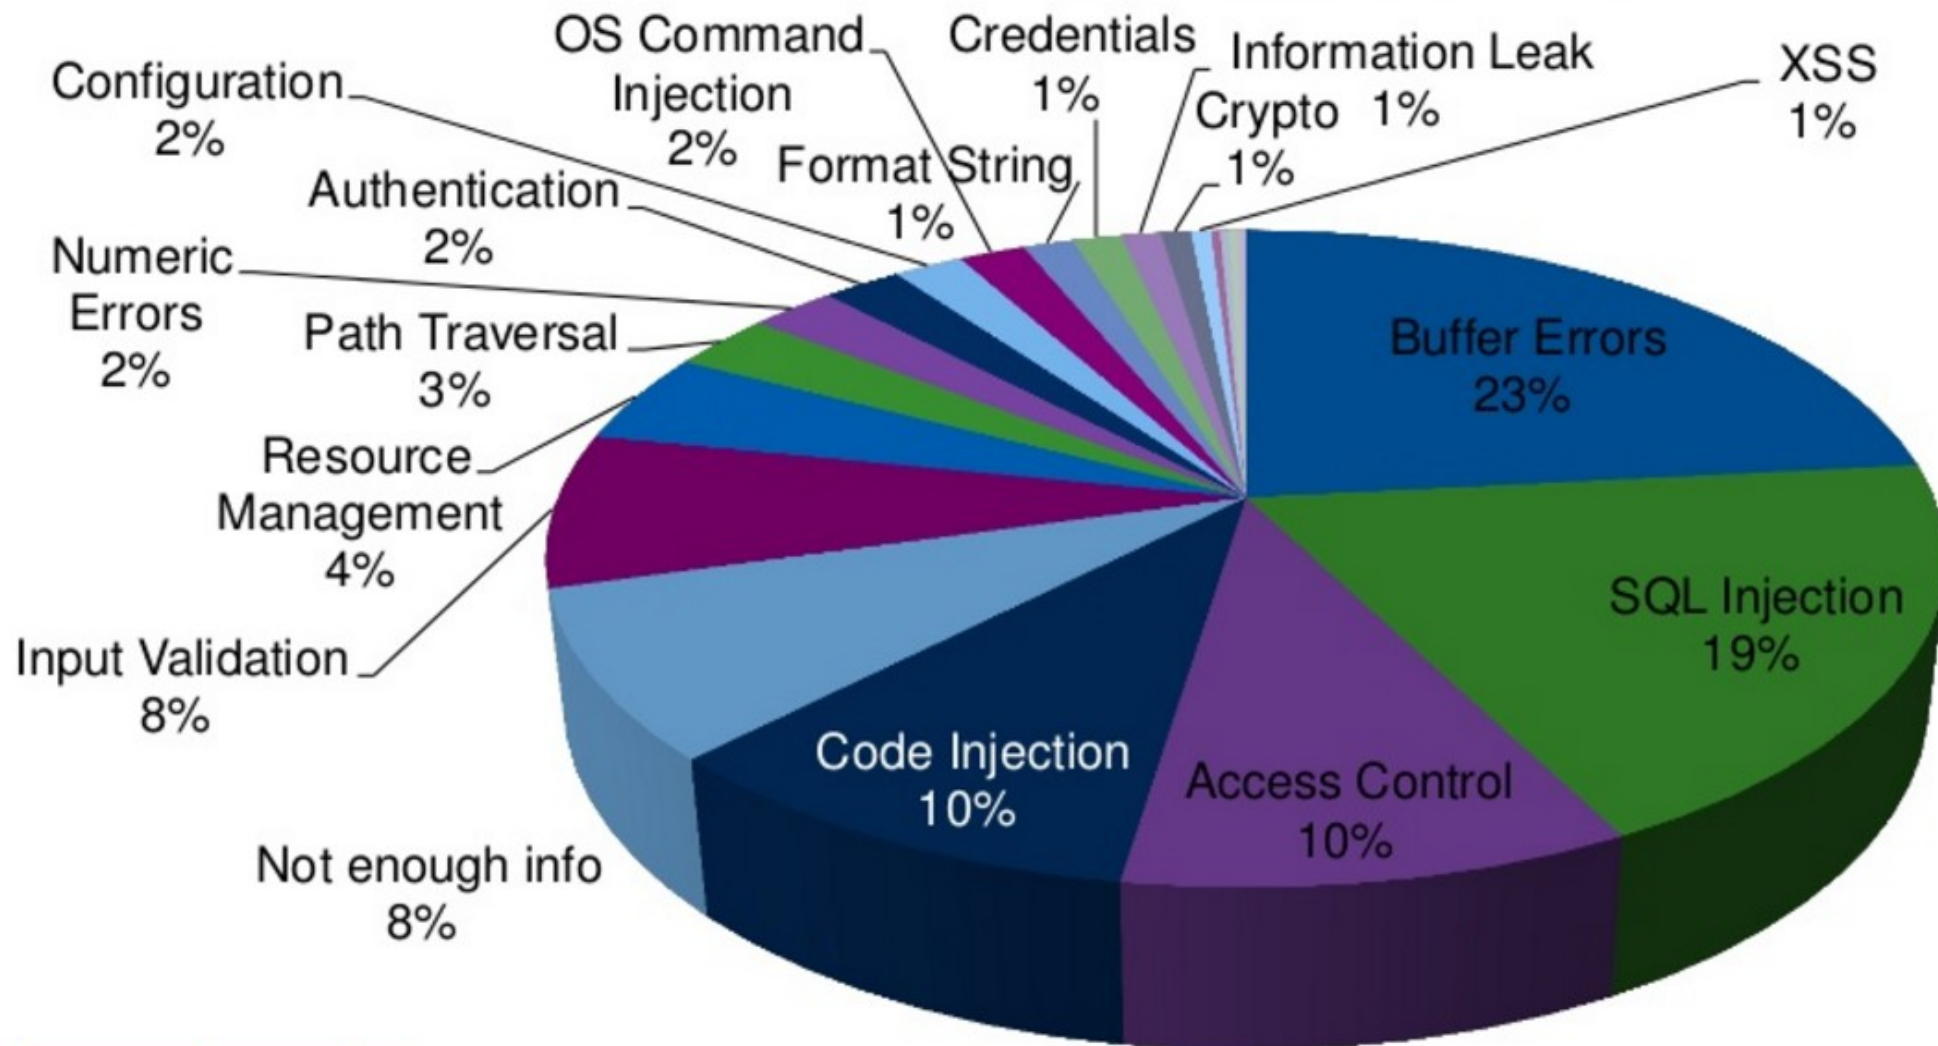
T-SEC-18-B

Cisco and/or its affiliates. All rights reserved.

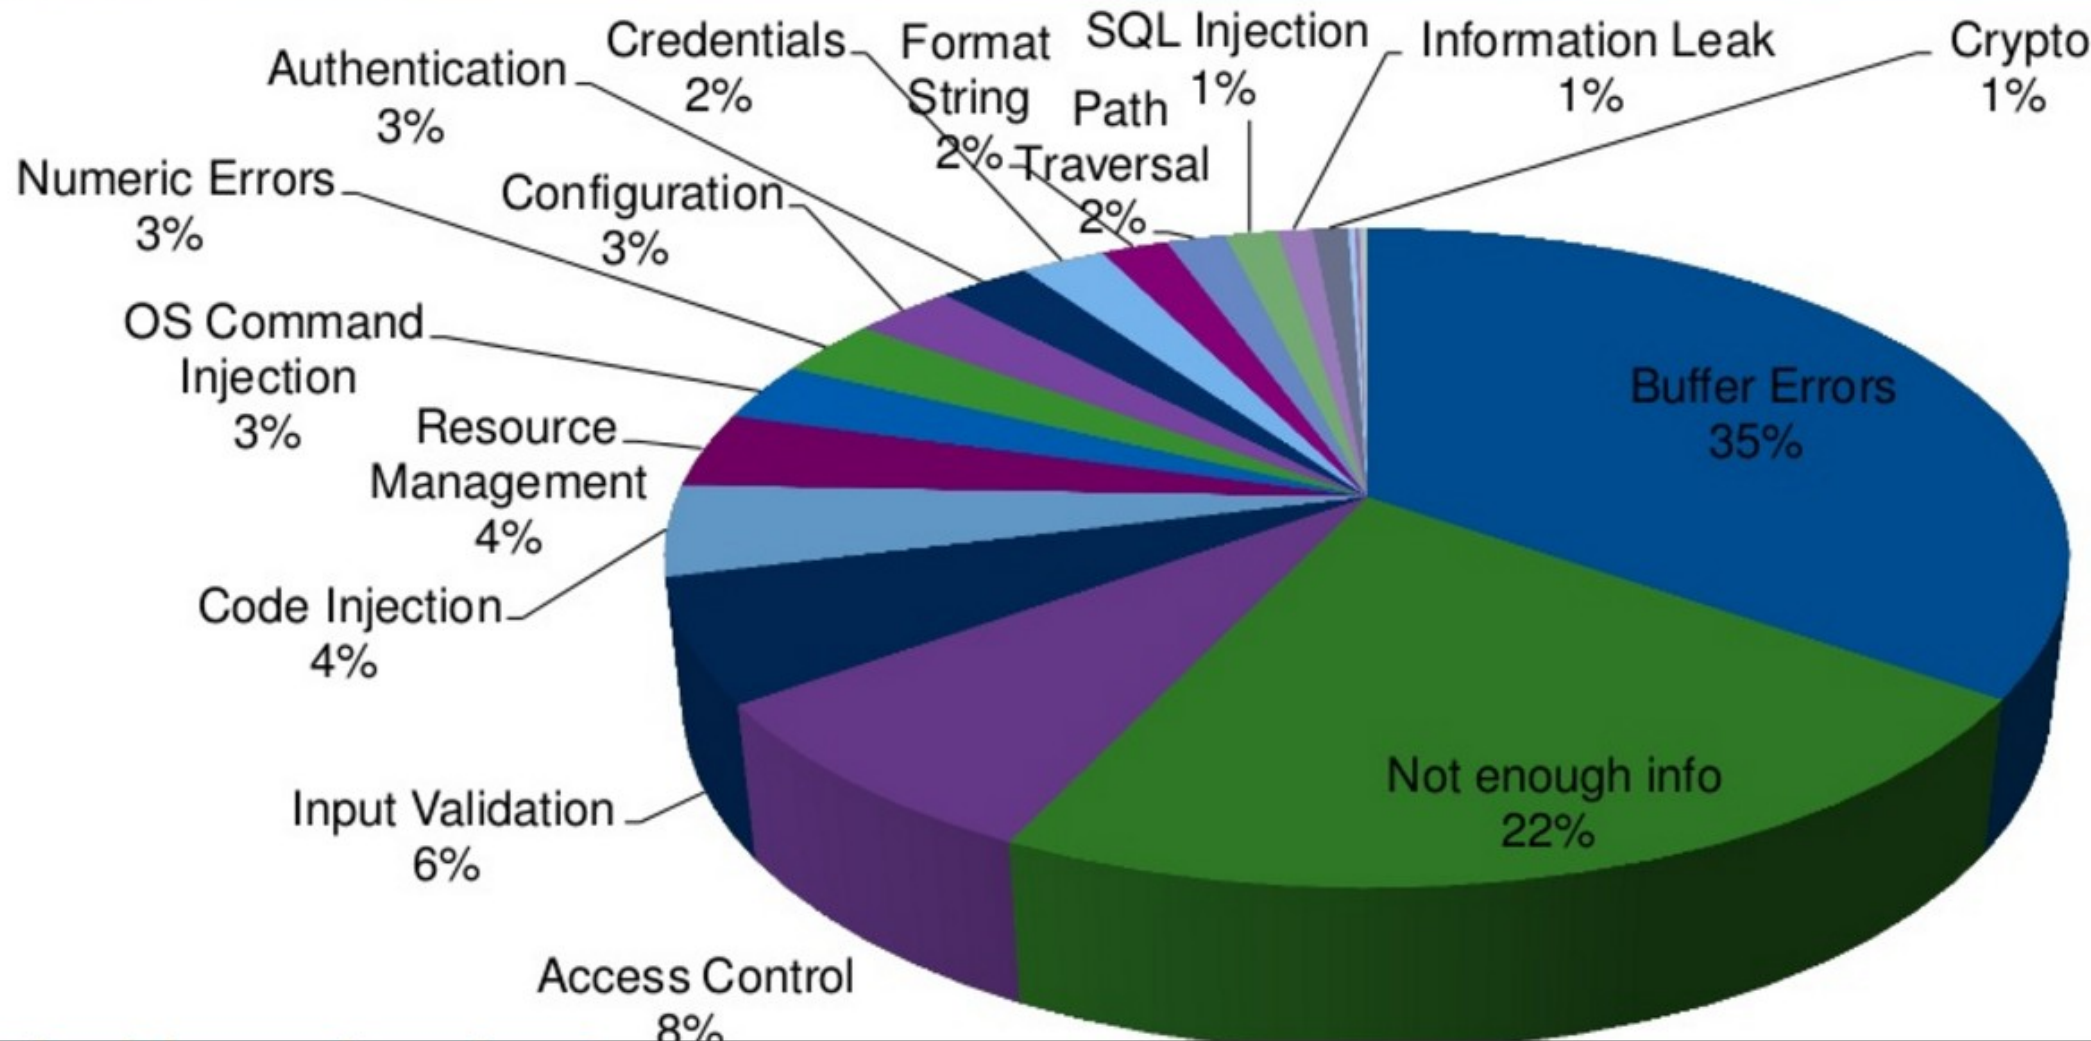
Cisco Public

Source: [Security Vulnerabilities in Modern Operating Systems](#), Cisco, Apr 2014

Serious Vulnerabilities by Type



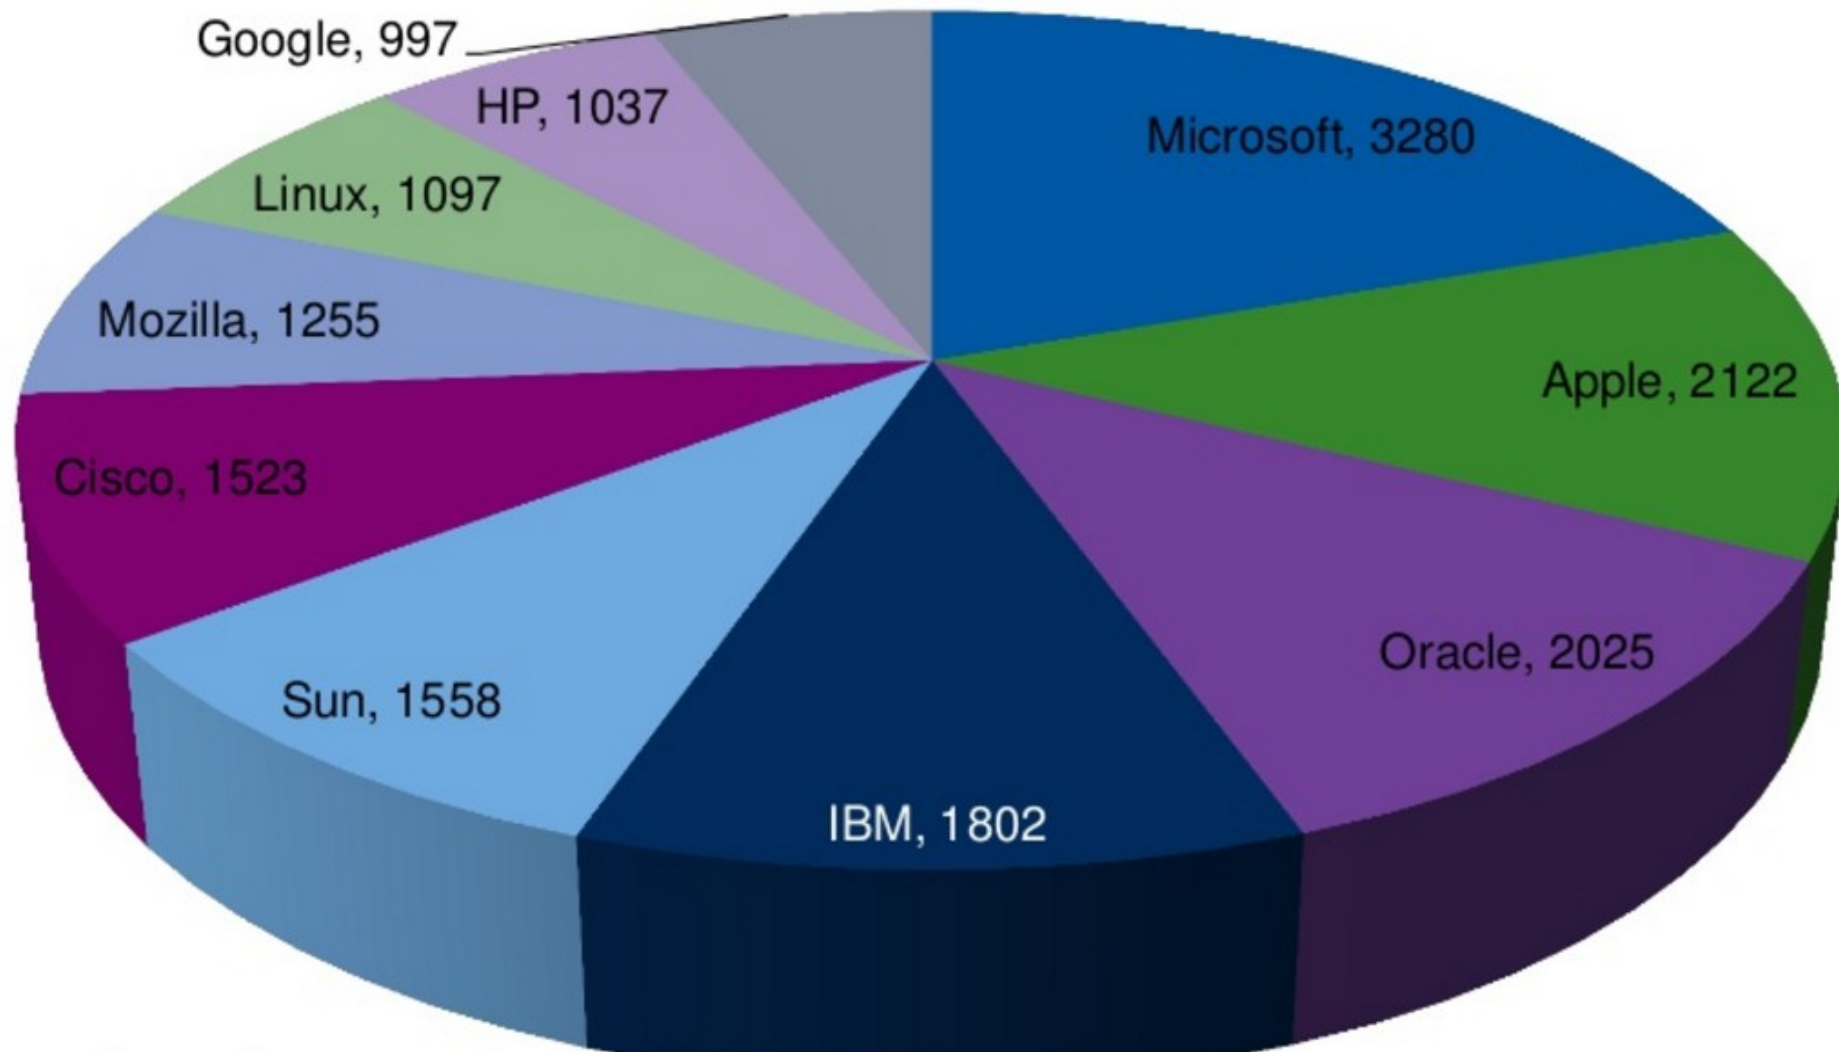
Critical Vulnerabilities by Type



◀ 16 of 55 ▶

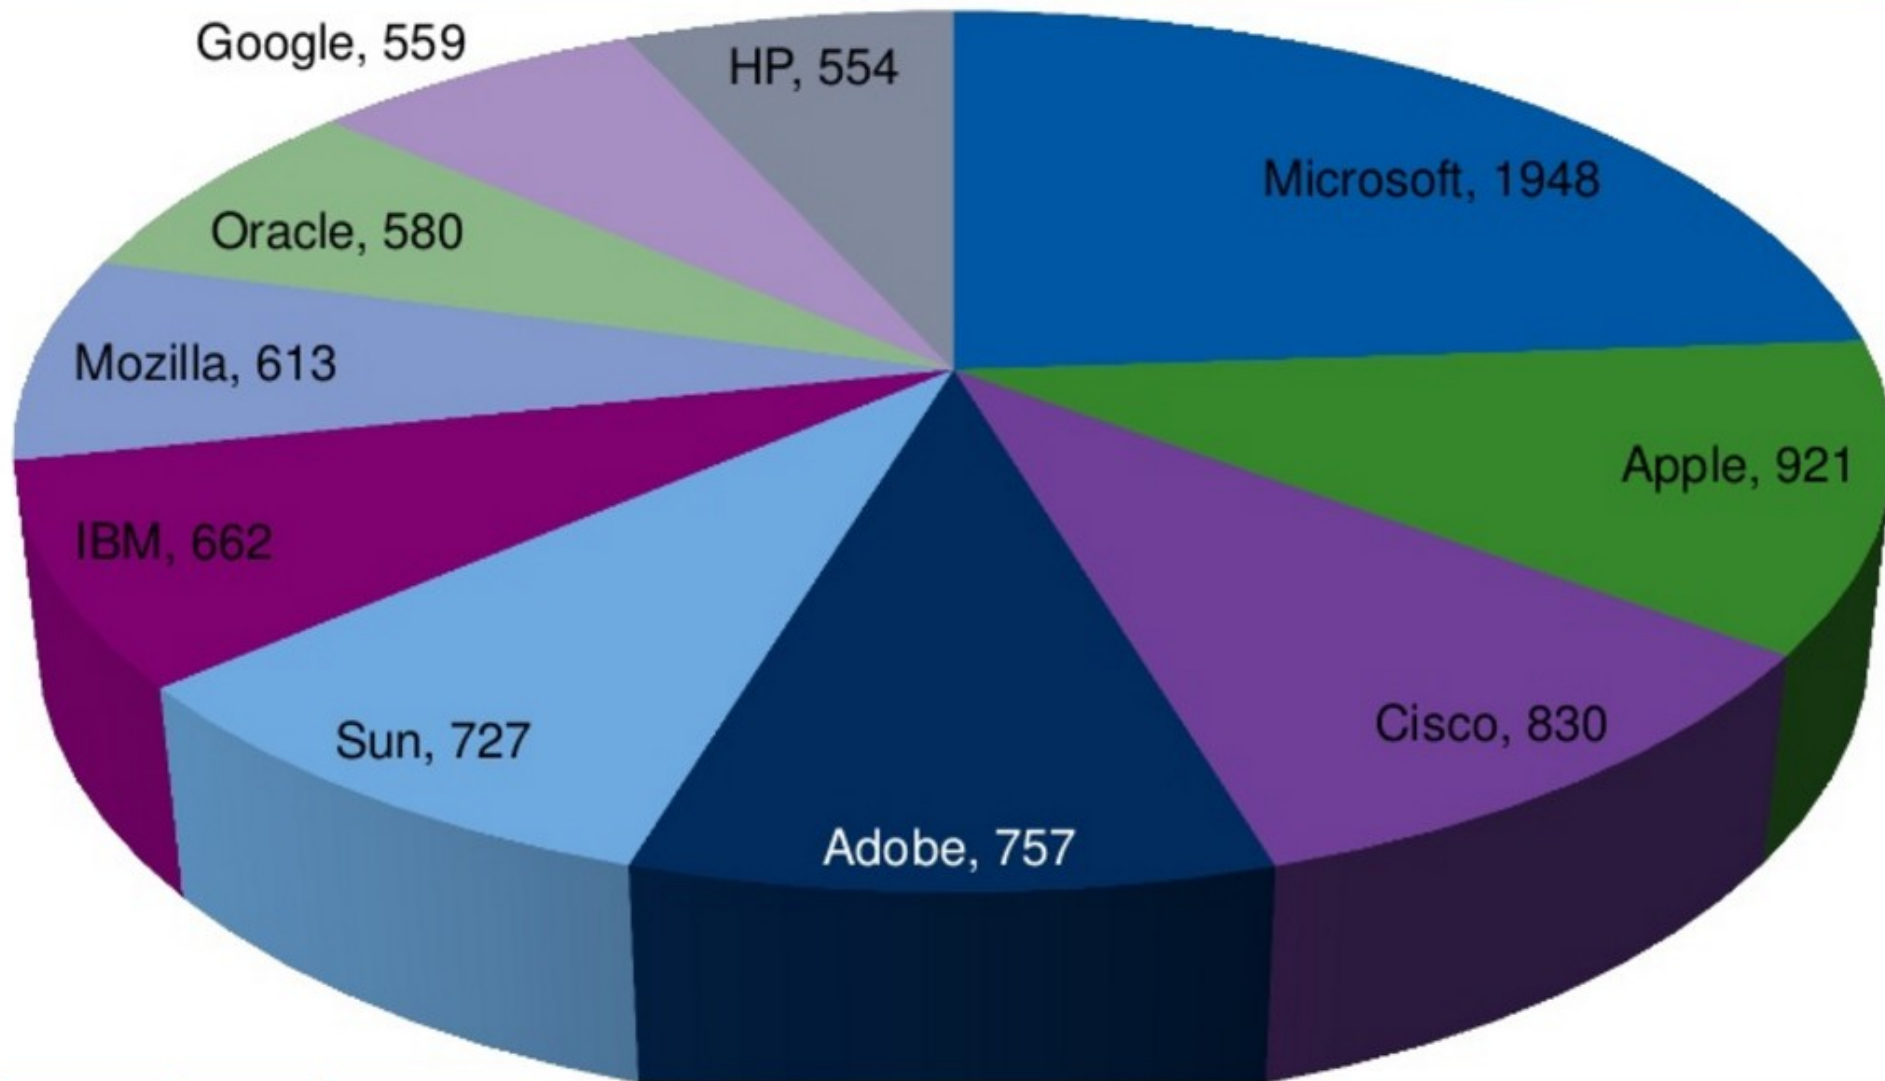
Source: [Security Vulnerabilities in Modern Operating Systems](#), Cisco, Apr 2014

Top 10 Vendors for Total Vulnerabilities

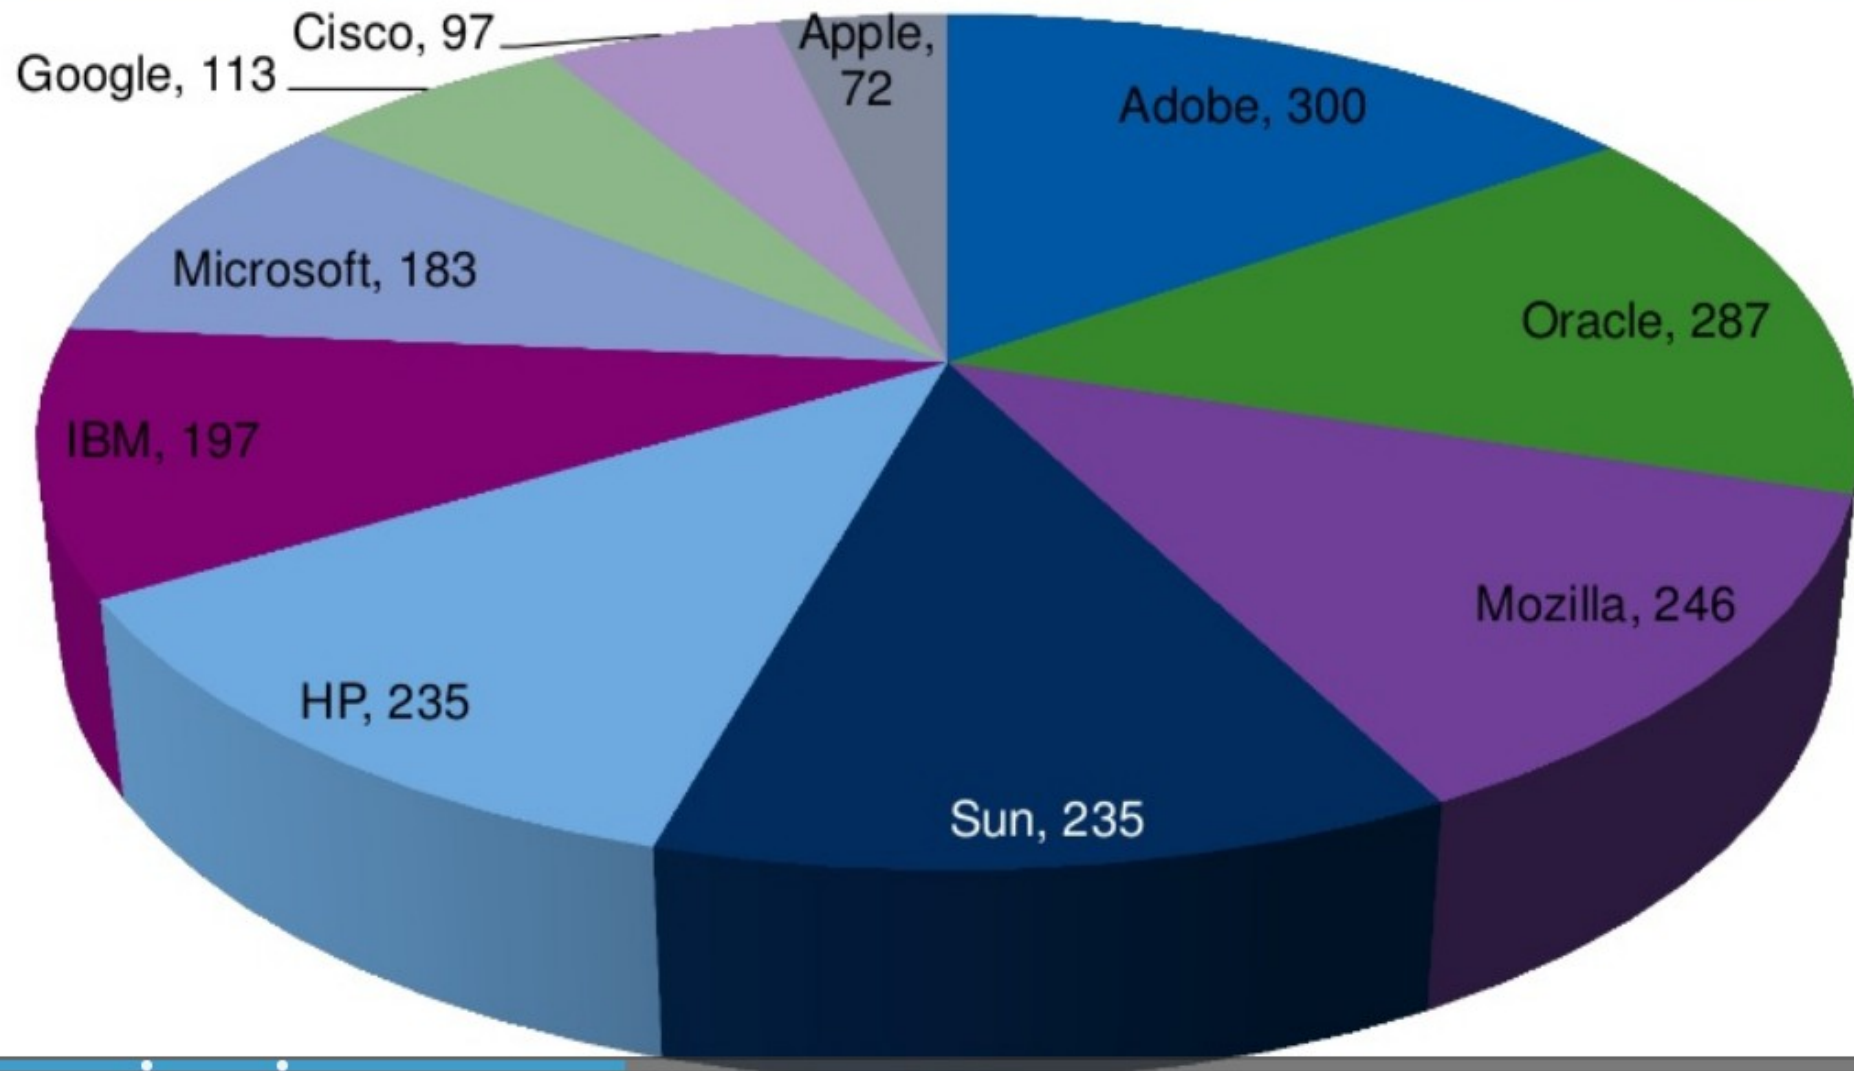


Source: [Security Vulnerabilities in Modern Operating Systems](#), Cisco, Apr 2014

Top 10 Vendors for Serious Vulnerabilities



Top 10 Vendors for Critical Vulnerabilities



T-SEC-18-B

Cisco and/or its affiliates. All rights reserved.

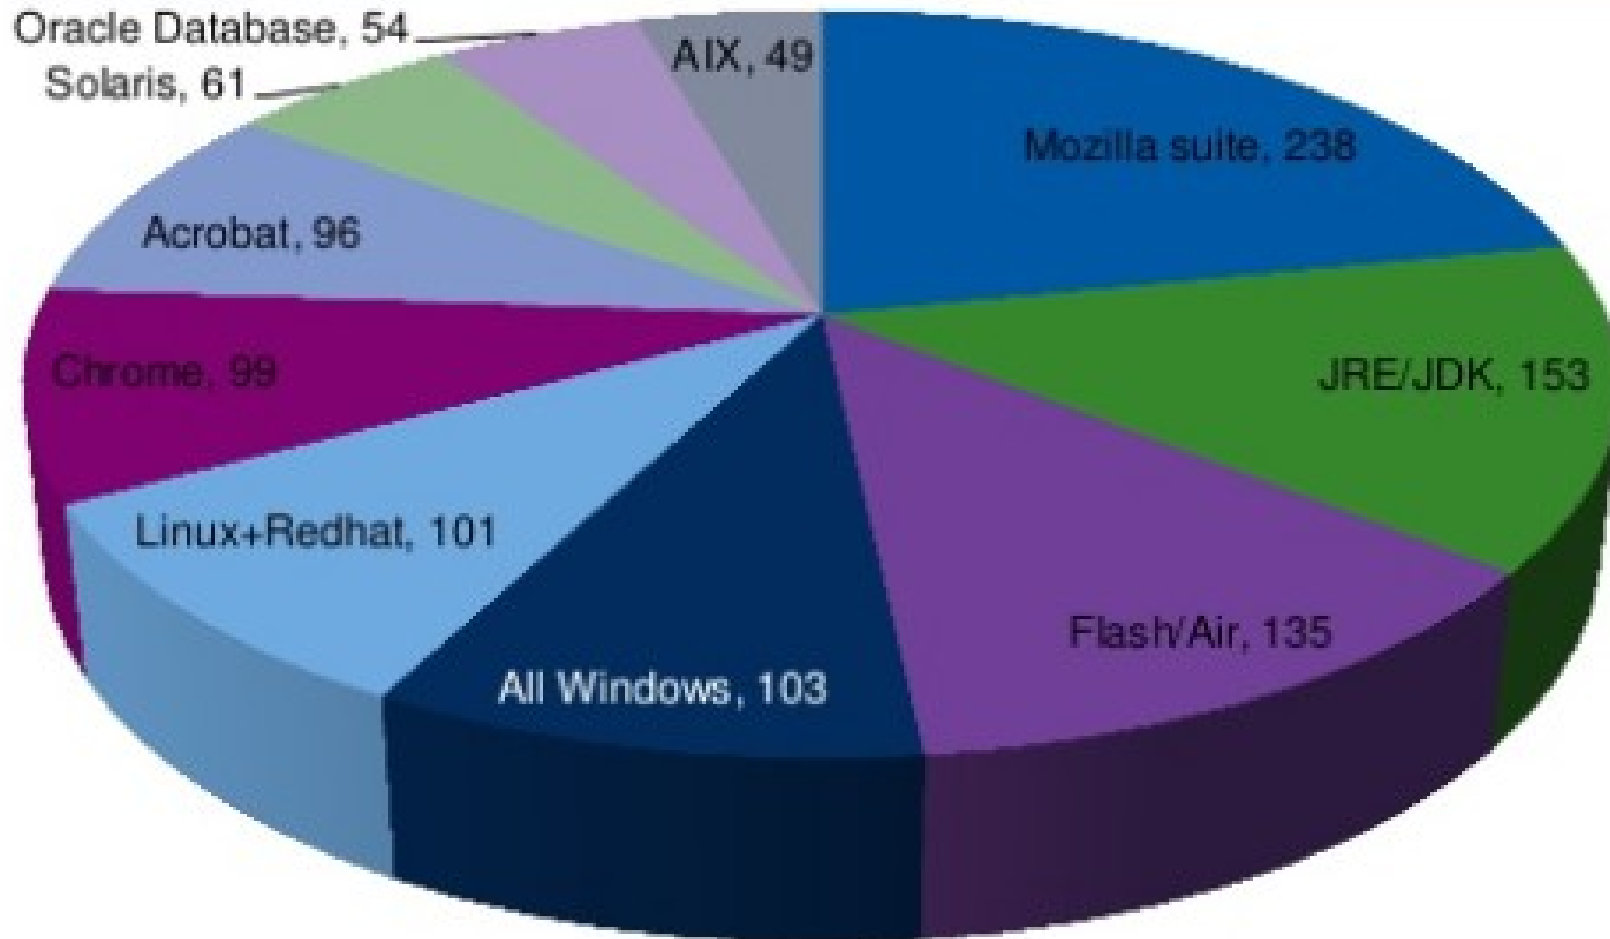
21 of 55

Cisco Public

Source: [Security Vulnerabilities in Modern Operating Systems](#), Cisco, Apr 2014

Top 10 Critically Vulnerable Products, totalled (similar)

Clip slide



Source: [Security Vulnerabilities in Modern Operating Systems](#), Cisco, Apr 2014

OS Security Vulnerabilities

- Vulnerabilities are here to stay
 - While serious vulnerabilities have been in decline, total vulnerabilities are not and neither are critical
 - At some point many vendors thought that hunting for enough vulnerabilities would make software secure
 - New features increase the attack surface or make previously non-exploitable errors exploitable
 - Using several non-serious vulnerabilities in concert could result in a more serious issue
 - Buffer overflows have been around for 25 years yet are still one of the top vulnerabilities
- Full report (up to 2012) available via <http://www.sourcefire.com/25yearsofvulns>

Source: *Security Vulnerabilities in Modern Operating Systems*, Cisco, Apr 2014

END “Security Vulnerabilities in Modern Operating Systems”

CISCO Presentation Slides

Buffer Overflow (BOF)

• • **BoF + Other Attacks in the Real-World**

- **MUST-SEE** [Real Life Examples](#) - gathers a few actual attacks of different kinds- phishing, password, crypto, input, BOFs, etc
- A few 'famous' (public) Buffer Overflow (BOF) Exploits
 - 02 Nov 1988: [Morris Worm](#) – first network 'worm'; exploits a BoF in fingerd (and DEBUG cmd in sendmail). [Article](#) and [Details](#)
 - [24 Sep 2014: ShellShock](#) [serious bug in bash!]
 - [15 July 2001: Code Red](#) and Code Red II ; [CVE-2001-0500](#)
 - [07 Apr 2014: Heartbleed](#) ; [CVE-2014-0160](#)
- [The Risks Digest](#)

Buffer Overflow (BOF)

• *BoF + Other Attacks in the Real-World*

Interesting!

- **Side-Channel attack examples:**
 - Mar 2017: [Hard Drive LED Allows Data Theft From Air-Gapped PCs](#)
 - [Exploiting the DRAM Rowhammer bug](#)
- **Gaming console hacks – due to BOF exploits**
 - Jan 2003: [Hacker breaks Xbox protection without mod-chip](#)
 - [PlayStation 2 Homebrew](#)
 - Wii [Twilight hack](#)
- [10 of the worst moments in network security history](#)

IoT - Attacks in the Real-World

- IoT devices in the real world: iotlineup.com + many more ...

Bitdefender BOX IoT Security Solution



Video Buy Website

Blocks incoming threats and can scan all your devices for vulnerabilities... »read more

Home Security and Safety, Home Automation, Network Security

Google Home Voice controller



Video Website

The connected voice controller from Google. Besides controlling your home it... »read more

Home Appliances, Home Automation, Remote Controls

Amazon Echo (2nd Generation) Voice controller



Video Buy Website

The connected voice controller from Amazon. Can give you information, music... »read more

Home Appliances, Home Automation, Remote Controls, Alexa Enabled

Nest Cam Indoor camera



Video Buy Website

The monitoring tool you've been waiting for. It brings all the benefits... »read more

Mr. Coffee Smart Coffeemaker



Video Website

The 10-Cup Smart Coffeemaker makes it easy to schedule, monitor, and... »read more

SmartMat Intelligent Yoga Mat

The interactive yoga utility that helps you perfect your yoga training through real time pressure sensing technology and smart mobile interfacing.

Video Website

Helps you perfect your yoga training through real time pressure sensing... »read more

IoT - Attacks in the Real-World

- **IoT Security Wiki : One Stop for IoT Security Resources**

Huge number of resources (whitepapers, slides, videos, etc) on IoT security

- **US-CERT
Alert (TA16-288A) - Heightened DDoS Threat Posed by Mirai and Other Botnets**

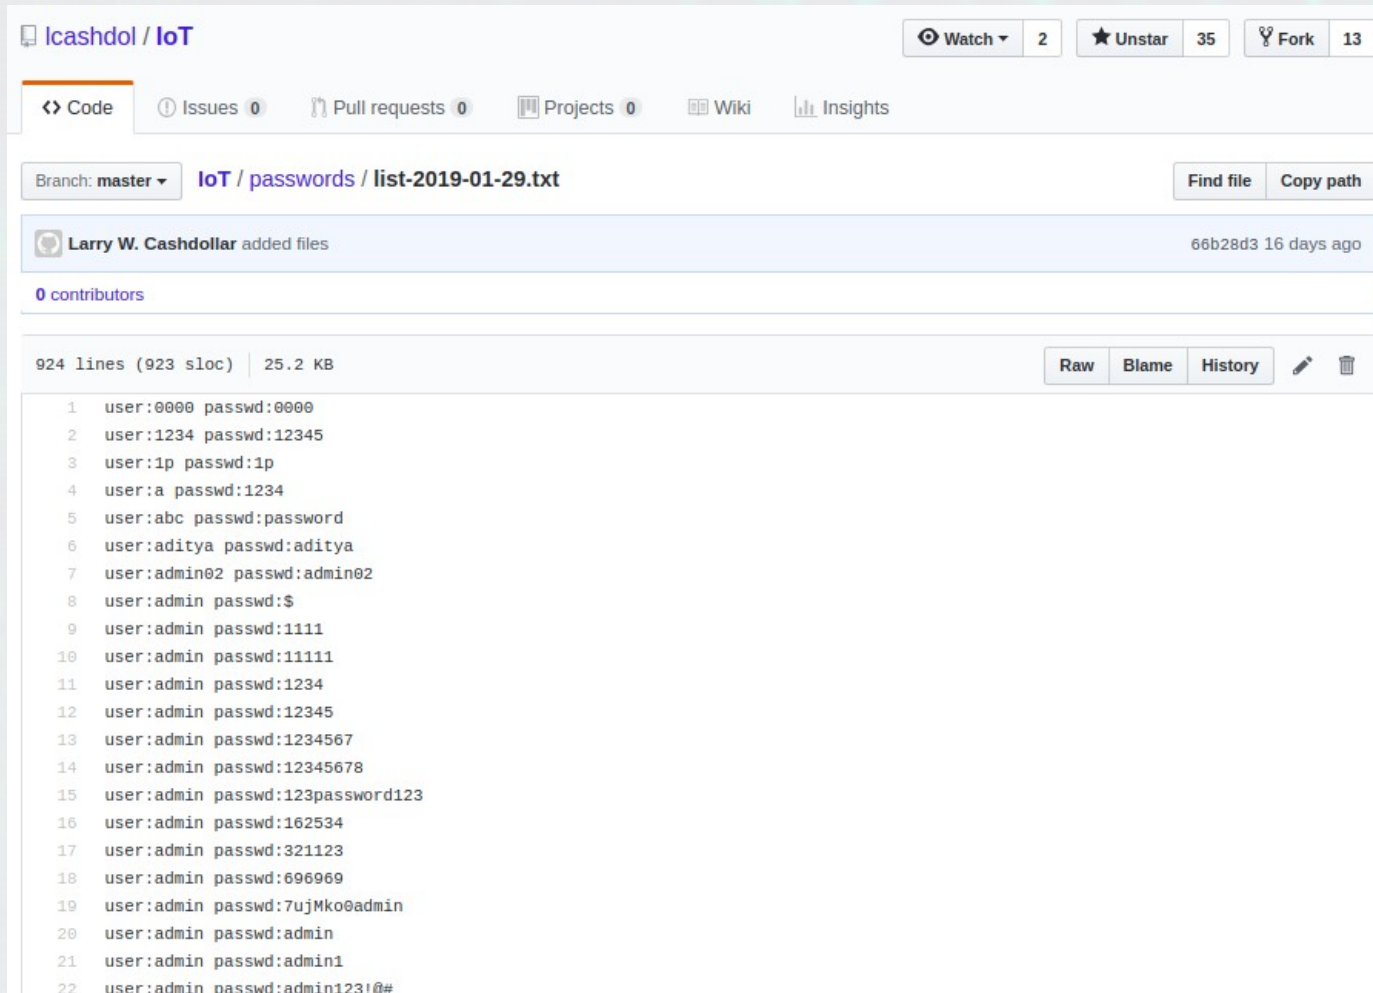
*"On September 20, 2016, Brian Krebs' security blog (krebsonsecurity.com) was targeted by a massive DDoS attack, one of the largest on record, exceeding 620 gigabits per second (Gbps).[1] An IoT botnet powered by **Mirai malware** created the DDoS attack.*

*The Mirai malware continuously scans the Internet for vulnerable IoT devices, which are then infected and used in botnet attacks. **The Mirai bot uses a short list of 62 common default usernames and passwords to scan for vulnerable devices.** Because many IoT devices are unsecured or weakly secured, this short dictionary allows the bot to access hundreds of thousands of devices.[2] The purported Mirai author claimed that over 380,000 IoT devices were enslaved by the Mirai malware in the attack on Krebs' website.[3]*

In late September, a separate Mirai attack on French webhost OVH broke the record for largest recorded DDoS attack. That DDoS was at least 1.1 terabits per second (Tbps), and may have been as large as 1.5 Tbps.[4] ..."

IoT - Attacks in the Real-World

<https://github.com/lcashdol/IoT/blob/master/passwords/list-2019-01-29.txt>



The screenshot shows a GitHub repository page for 'lcashdol / IoT'. The file 'list-2019-01-29.txt' is selected, showing 924 lines of code. The file contains a list of user:password pairs. The repository has 2 watchers, 35 unstars, and 13 forks. The file was added by Larry W. Cashdollar 16 days ago.

```
1 user:0000 passwd:0000
2 user:1234 passwd:12345
3 user:1p passwd:1p
4 user:a passwd:1234
5 user:abc passwd:password
6 user:aditya passwd:aditya
7 user:admin02 passwd:admin02
8 user:admin passwd:$
9 user:admin passwd:1111
10 user:admin passwd:11111
11 user:admin passwd:1234
12 user:admin passwd:12345
13 user:admin passwd:1234567
14 user:admin passwd:12345678
15 user:admin passwd:123password123
16 user:admin passwd:162534
17 user:admin passwd:321123
18 user:admin passwd:696969
19 user:admin passwd:7ujMko0admin
20 user:admin passwd:admin
21 user:admin passwd:admin1
22 user:admin passwd:admin123!@#
```

IoT - Attacks in the Real-World

- Hacking DefCon 23's IoT Village Samsung fridge, Aug 2015 (DefCon 23)

- “HACKING IoT: A Case Study on Baby Monitor Exposures and Vulnerabilities” [[PDF](#)]

“... these devices are marketed and treated as if they are single purpose devices, rather than the general purpose computers they actually are. ...

IoT devices are actually general purpose, networked computers in disguise, running reasonably complex network-capable software. In the field of software engineering, it is generally believed that such complex software is going to ship with exploitable bugs and implementation-based exposures. Add in external components and dependencies, such as cloud-based controllers and programming interfaces, the surrounding network, and other externalities, and it is clear that vulnerabilities and exposures are all but guaranteed.”

<< See the PDF pg 6, ‘Ch 5: COMMON VULNERABILITIES AND EXPOSURES FOR IoT DEVICES’ ; old and new vulnerabilities mentioned;

Pg 9: ‘Disclosures’ - the vulns uncovered in actual products >>

Just too much. Bottom line: critical to outsource or do pentesting yourself!

InfoSec: Focus back on Developers

- **Source: the “DZone Guide to Proactive Security”, Vol 3, Oct 2017**
 - Ransomware & malware attacks up in 2017
 - WannaCry, Apr '17 : \$100,000 in bitcoin
 - NotPetya, June '17 : *not* ransomware, wiper malware
 - *CVEdetails* shows that # vulns in 2017 is 14,714, the highest since 1999! 2018 is on track to overtake that
 - “... how can the global business community counteract these threats? The answer is to catch vulnerabilities sooner in the SDLC ...”
 - “... **Shifting security concerns (left) towards developers**, creates an additional layer of checks and can eliminate common vulnerabilities early on through simple checks like validating inputs and effective assignment of permissions. “

Buffer Overflow (BOF)

● ● ● Preliminaries – the Process VAS

What exactly is a buffer overflow (BOF)?

- Prerequisite – an understanding of the process stack!
- *Soon, we shall see some very simple ‘C’ code to understand this first-hand.*
- *But before that, an IMPORTANT Aside: As we shall soon see, nowadays several mitigations/hardening technologies exist to help prevent BOF attacks. So, sometimes the question (SO InfoSec) arises: “Should I bother teaching buffer overflows any more?”:*
Short answer, “YES”:

*“... Yes. Apart from the systems where buffer overflows lead to successful exploits, full explanations on buffer overflows are always a great way to demonstrate **how you should think about security**. Instead on concentrating on how the application should run, see what can be done in order to make the application derail.*

*Also, regardless of stack execution and how many screaming canaries you install, **a buffer overflow is a bug**. All those security features simply alter the consequences of the bug: instead of a remote shell, you “just” get an immediate application crash. Not caring about application crashes (in particular crashes which can be triggered remotely) is, at best, very sloppy programming. ...”*

Buffer Overflow (BOF)

● ● ● Preliminaries – the Process VAS

On 17 Nov 2017, [Linus wrote on the LKML](#):

"...

As a security person, you need to repeat this mantra:

"security problems are just bugs"

and you need to _internalize_ it, instead of scoff at it.

The important part about "just bugs" is that you need to understand that the patches you then introduce for things like hardening are primarily for DEBUGGING.

..."

Buffer Overflow (BOF)

••• Preliminaries – the Process VAS


- **UNIX philosophy:**
*“Everything is a process;
if it’s not a process, it’s a file”*
- **Process has a Virtual Address Space (VAS); consists of “segments”:**
 - Text (code)
 - Data
 - ‘Other’ mappings
 - Stack

Buffer Overflow (BOF)

●●● Preliminaries – the Process

The Process

Virtual Address Space (VAS)



Buffer Overflow (BOF)

●●● Preliminaries – the Process VAS

Visualizing the Process Virtual Address Space (VAS) with the *vasu_grapher* utility:

git clone
https://github.com/kaiwan/vasu_grapher.git

(still under development)

```
~ $ vasu_grapher 1

[===== V A S U _ G R A P H E R =====]
Virtual Address Space Usermode (VASU) process GRAPHER (via /proc/1/maps)
https://github.com/kaiwan/vasu_grapher
Fri Feb 15 08:10:43 IST 2019

[===== Start memory map PID 1 (systemd) =====]
do_vgraph.sh: Processing, pl wait ...
+-----+ 0000000000000000
| [ NULL trap ] [4 KB] |
+-----+ 0000000000001000
|<< ... Sparse Region ... >> [92145937500 KB 89986267.08 MB 87877.21 GB]
|
|
|
|
|
|
|
|
|
|
|
|
+-----+ 000055d14db18000
|/lib/systemd/systemd [180 KB] |
+-----+ 000055d14db45000
|/lib/systemd/systemd [672 KB] |
+-----+ 000055d14dbed000
|/lib/systemd/systemd [300 KB] |
+-----+ 000055d14dc38000
|/lib/systemd/systemd [256 KB] |
+-----+ 000055d14dc78000
|/lib/systemd/systemd [4 KB] |
+-----+ 000055d14dc79000
|<< ... Sparse Region ... >> [4312 KB 4.21 MB]
|
```


Buffer Overflow (BOF)

● ● ● ● ● ● ● ● ● ● ● ● ● ● ● ● Preliminaries – the STACK

The Classic Case

Lets imagine that here below is part of the (drastically simplified) **Requirement Spec** for a console-based app:

- *Write a function 'foo()' that accepts the user's name, email id and employee number*

Buffer Overflow (BOF)

●●●●● Preliminaries – the STACK

The Classic Case – function foo() implemented below by app developer in ‘C’:

```
[...]  
static void foo(void)  
{  
    char local[128];  
    printf("Name: ");  
    gets(local);  
    [...]  
}
```

Buffer Overflow (BOF)

... Preliminaries – the STACK

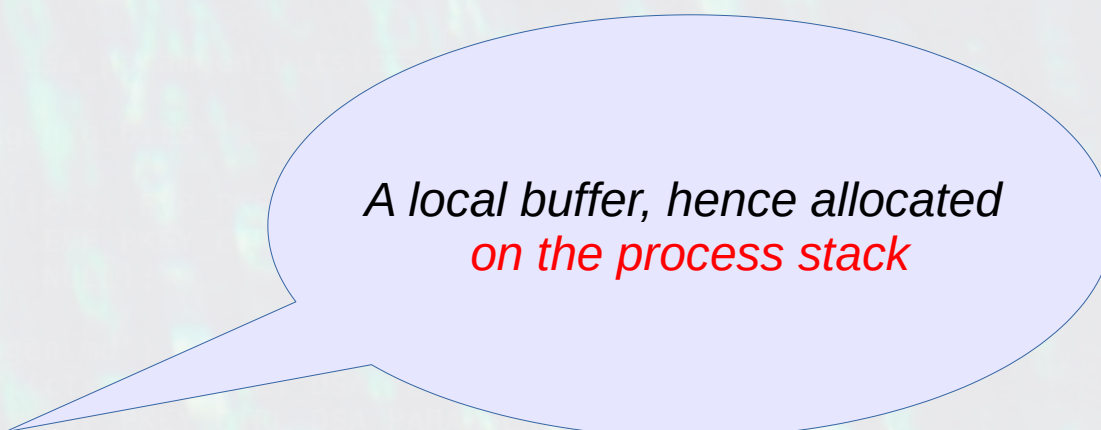
- **Why have a “stack” at all ?**
 - Humans write code using a 3rd or 4th generation high-level language (*well, most of us anyway :-)*
 - The processor hardware does not ‘get’ what a **function** is, what parameters, local variables and return values are!

Buffer Overflow (BOF)

... Preliminaries – the STACK

The Classic Case – function foo() implemented below by app developer in ‘C’:

```
[...]  
static void foo(void)  
{  
    char local[128];  
    printf("Name: ");  
    gets(local);  
    [...]  
}
```



A local buffer, hence allocated
on the process stack

Buffer Overflow (BOF)

..... Preliminaries – the STACK

- ***So, what really, is this process stack ?***
 - it's just memory treated with special semantics
 - Theoretically via a “PUSH / POP” model
 - [Realistically, the OS just allocates pages as required to “grow” the stack and the SP register tracks it]
 - “Grows” towards lower (virtual) addresses; called a “downward-growing stack”
 - this attribute is processor-specific; it's the case for most modern CPUs, including x86, x86_64, ARM, ARM64, MIPS, PPC, etc

Buffer Overflow (BOF)

.....Preliminaries – the STACK

- ***Why have a “stack” at all ?***
 - *The saviour:* the **compiler** generates assembly code which enables the *function-calling-parameter-passing-local-vars-and-return* mechanism
 - By making use of the stack!
 - How exactly?

... Aha ...

Buffer Overflow (BOF)

.....Preliminaries – the STACK

• The Stack

- When a program calls a function, the compiler generates code to setup a call stack, which consists of individual **stack frames**
- A **stack frame** can be visualized as a “block” containing all the metadata necessary for the system to process the “function” call and return
 - Access it's parameters, if any
 - Allocate it's local variables
 - Execute it's code (text)
 - Return a value as required

Buffer Overflow (BOF)

.....Preliminaries – the STACK

- **The Stack Frame**

- Hence, the stack frame will require a means to
 - Locate the previous (caller's) stack frame (achieved via the **SFP** – Stack Frame Pointer) *[technically, the frame pointer is Optional]*
 - Gain access to the function's **parameters** (iff passed via stack, see the processor ABI)
 - Store the address to which control must continue, IOW, the **RETurn address**
 - Allocate storage for the function's **local variables**
- Turns out that the exact stack frame layout is very processor-dependant (depends on it's ABI, calling conventions)
- [In this presentation, we consider the typical IA-32 / ARM-32 stack frame layout]

Buffer Overflow (BOF)

.....*Preliminaries – the STACK*

Recall our simple ‘C’ code:

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}
```


Buffer Overflow (BOF)

.....*Preliminaries – the STACK*

Have it called from main()

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}

void main() {
    foo();
}
```

Buffer Overflow (BOF)

.....Preliminaries – the STACK

When main() calls foo(), a stack frame is setup (as part of the call stack)

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}

void main() {
    foo();
}
```

Buffer Overflow (BOF)

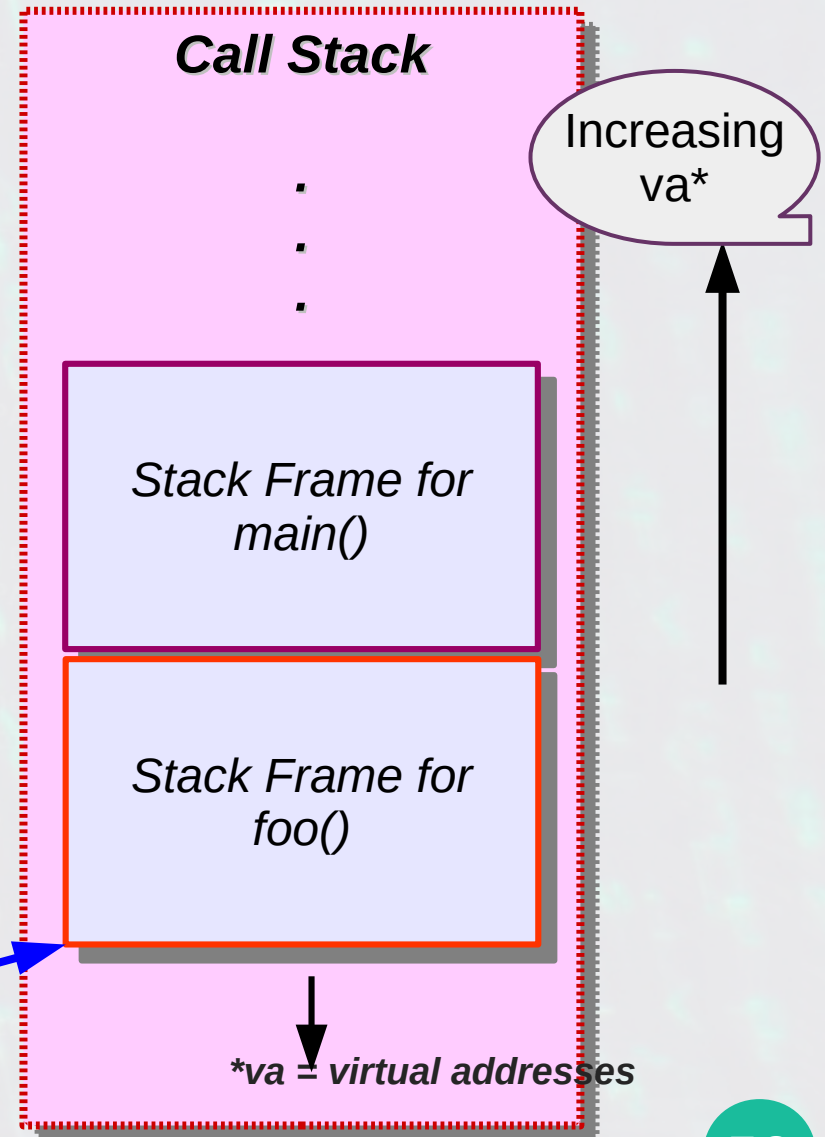
.....Preliminaries – the STACK

- When `main()` calls `foo()`, a stack frame is setup (as part of the call stack)

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}


void main() {
    foo();
}
```

SP (top of the stack)
Lowest address



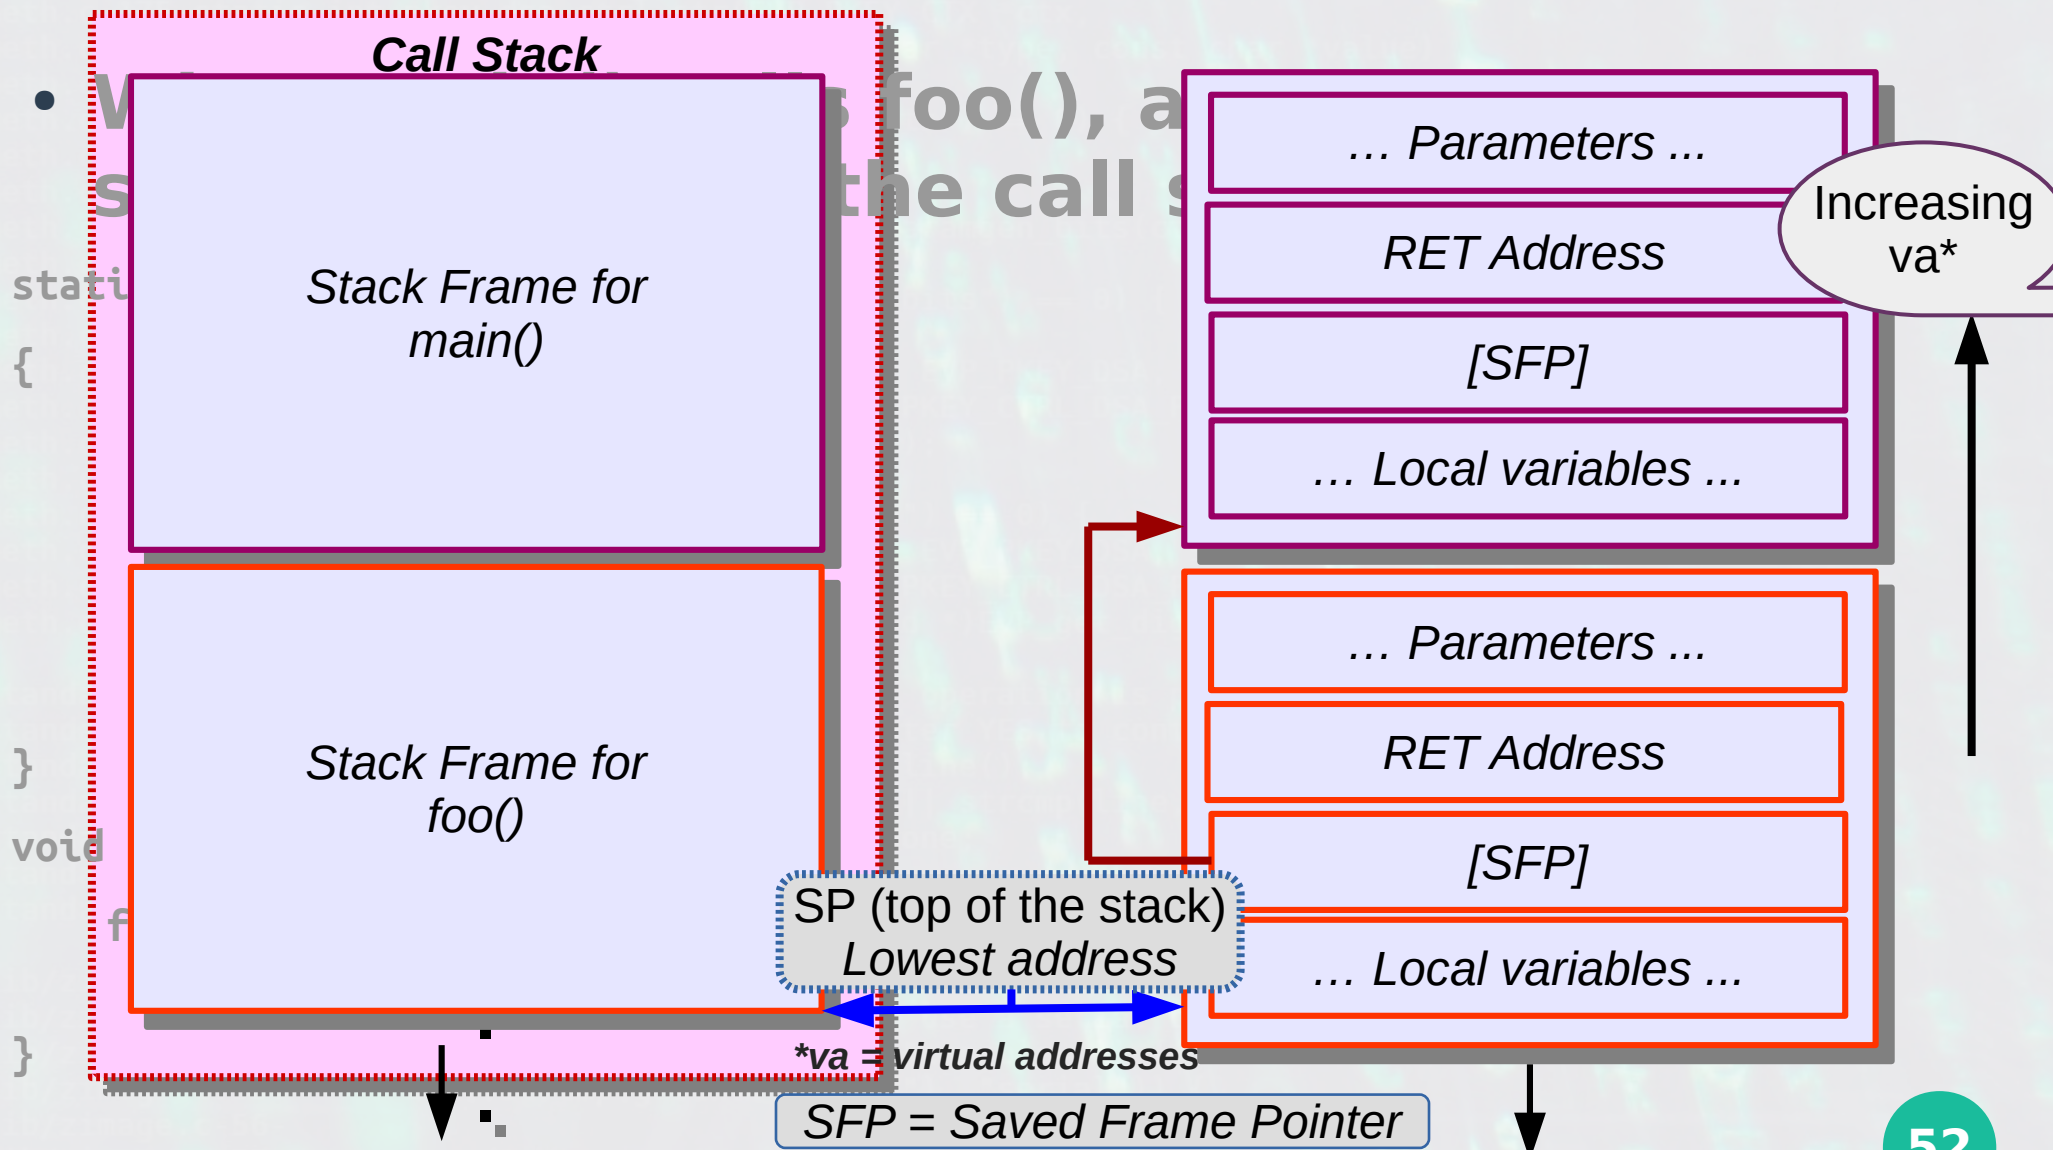
Buffer Overflow (BOF)

.....Preliminaries – the STACK



Buffer Overflow (BOF)

.....Preliminaries – the STACK



Buffer Overflow (BOF)



Wikipedia on BOF

In computer security and programming, a buffer overflow, or buffer overrun, **is an anomaly** where a program, while writing data to a buffer, overruns the buffer's boundary and overwrites adjacent memory locations.

Buffer Overflow (BOF)

● ● ● ● ● ● ● *A Simple BOF*

Recall:

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}

void main() {
    foo();
}
```


Buffer Overflow (BOF)

... A Simple BOF

At runtime

```
static void foo(void)
{
    char local[128];
    printf("Name: ");
    gets(local);
    [...]
}

void main() {
    foo();
}
```



Buffer Overflow (BOF)

... *A Simple BOF*

Lets give it a spin-

```
$ gcc getdata.c -o getdata  
[...]
```

```
$ printf "AAAABBBBCCCCDDDD" | ./getdata  
Name: Ok, about to exit...  
$
```


Buffer Overflow (BOF)

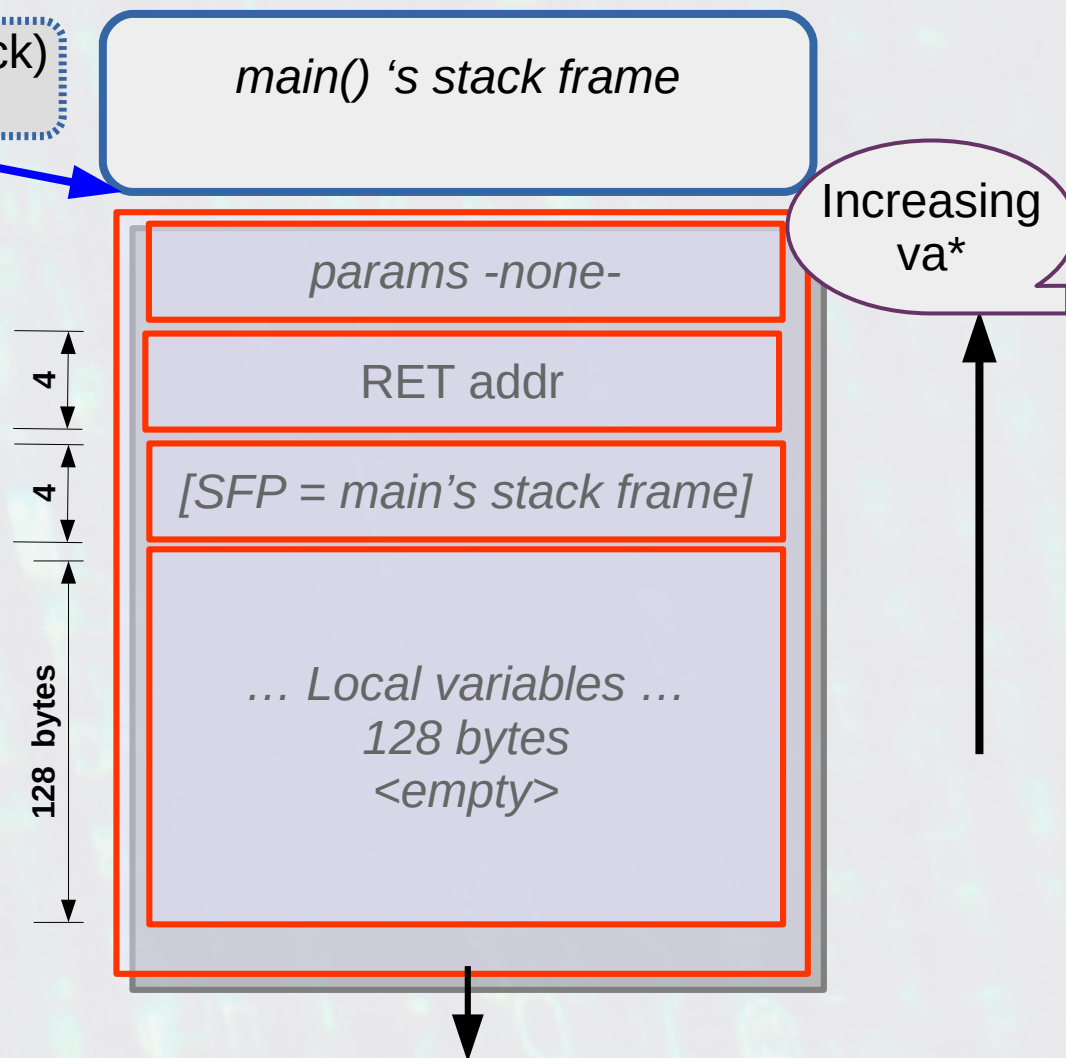
.....A Simple BOF

Okay, step by step:
Step 1 of 4 :

Prepare to execute;
main() is called

```
$ gcc getdata.c -o getdata  
[...]  
$  
$ printf "AAAABBBBCCCCDDDD"  
| ./getdata
```

SP (top of the stack)
Lowest address



Buffer Overflow (BOF)

.....A Simple BOF

Okay, step by step:

Step 2 of 4 :

Input 16 characters into the local buffer via the gets();

```
$ gcc getdata.c -o getdata
```

```
[...]
```

```
$
```

16 chars
written into the stack
@ var 'local'

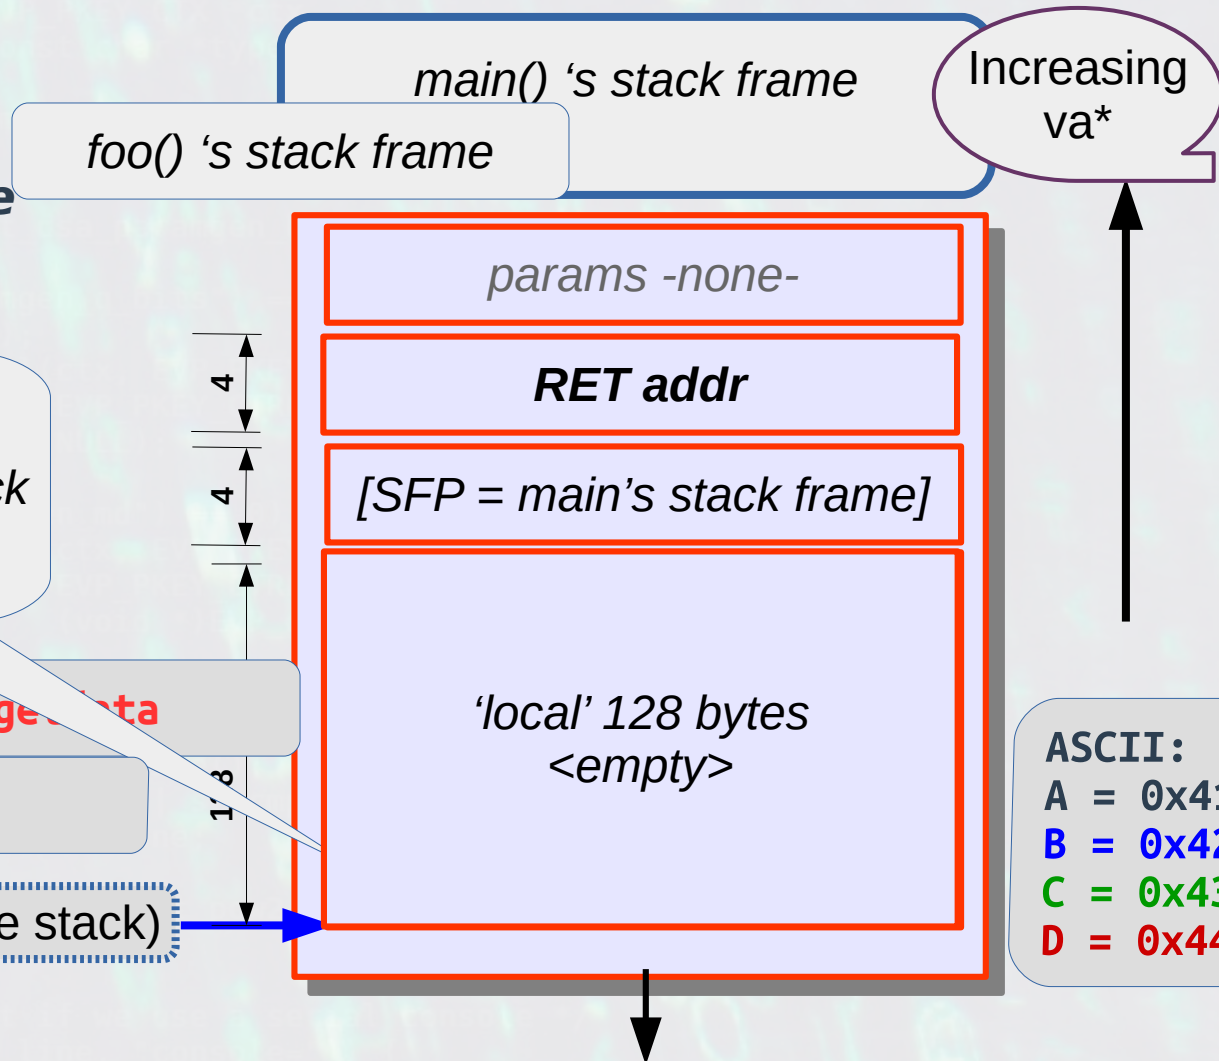
```
$ printf "AAAABBBBCCCCDDDD" | ./getdata
```

```
Name: Ok, about to exit...
```

```
$
```

All okay!

(for the stack)



Buffer Overflow (BOF)

..... A Simple **BOF – Action!**

Okay, step by step:
Step 3 of 4 :

Input **128+4+4 = 136** characters
into the local buffer via the `gets()`;

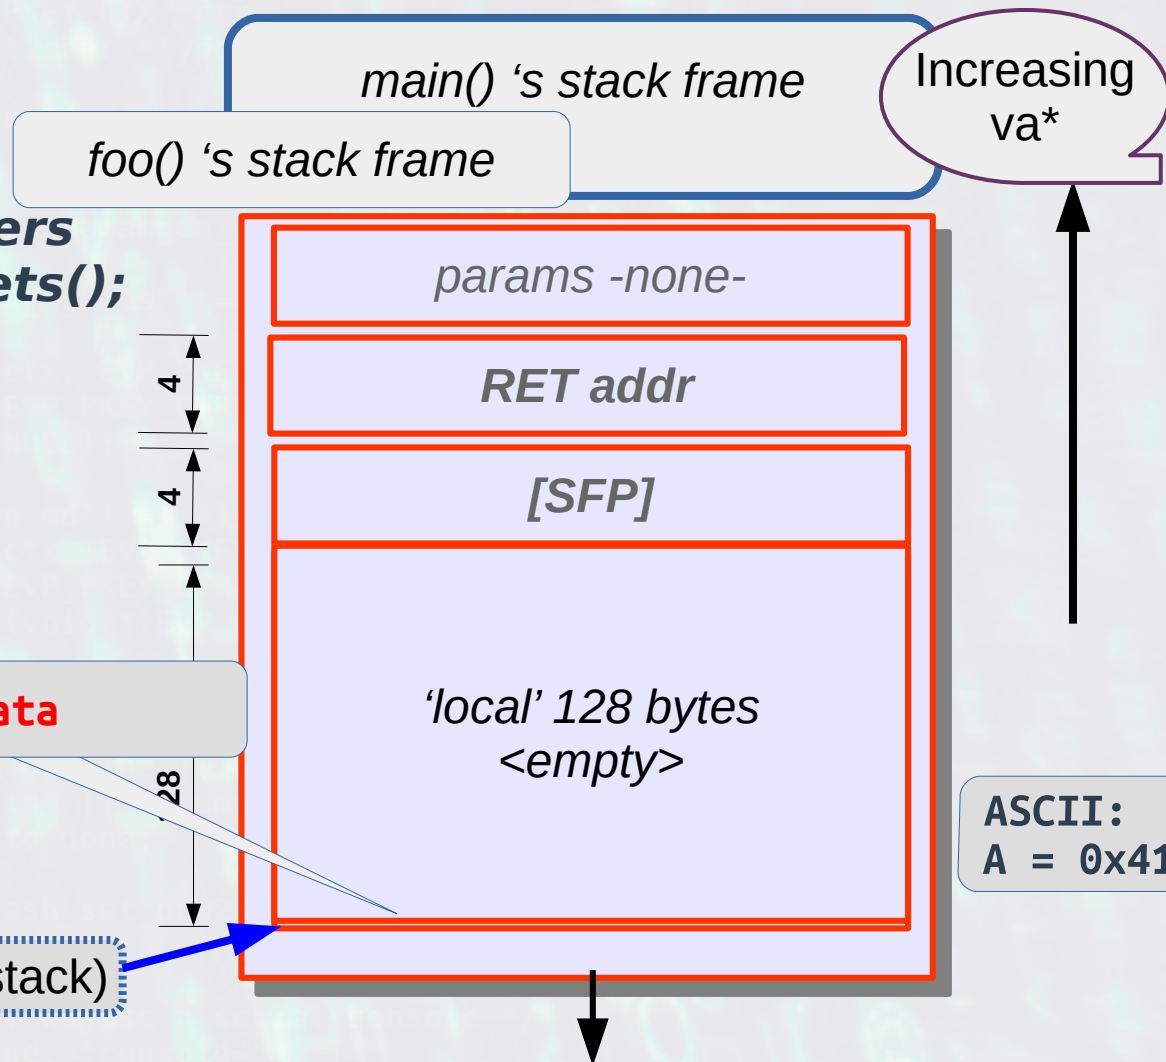
thus Overflowing it!

```
$ gcc getdata.c
[...]
```

136 chars
written into
the stack @ var 'local'

```
$ perl -e 'print "A"x136' | ./getdata
```

SP (top of the stack)



Buffer Overflow (BOF)

..... A Simple **BOF** – Action!

Okay, step by step,
Step 4 of 4 :

Input 128+4+4 = 136 characters

Why?

As the processor tries to return to the designated RET address, it attempts to access and execute code at virtual address 0x41414141 that's likely not there or is illegal

```
$ printf "AAAABBBBCCCCDDDD"
AAAABBBBCCCCDDDD$
```

```
$ perl -e 'print "A"x136' | ./getdata
```

Segmentation fault


```
$
```

SP (top of the stack)

It segfaults !

main() 's stack frame
foo() 's stack frame

Increasing va*



Buffer Overflow (BOF)

A Simple BOF / Where's the Problem?



So, where exactly is the problem / bug?

```
static void foo(void)
{
    char local[128];
    printf("Name: ")
    gets(local);
    [...]
}

void main()
{
    foo();
}
```

“Secure Programming for LINUX and UNIX HOWTO”
David Wheeler

Dangers in C/C++

C users must **avoid using dangerous functions that do not check bounds** unless they've ensured that the bounds will never get exceeded. Functions to avoid in most cases (or ensure protection) **include** the functions `strcpy(3)`, `strcat(3)`, `sprintf(3)` (with cousin `vsprintf(3)`), **and** `gets(3)`. These should be replaced with functions such as `strncpy(3)`, `strncat(3)`, `snprintf(3)`, and `fgets(3)` respectively, [...] The `scanf()` family (`scanf(3)`, `fscanf(3)`, `sscanf(3)`, `vscanf(3)`, `vsscanf(3)`, and `vfscanf(3)`) is often dangerous to use [...]

Buffer Overflow (BOF)

*A Simple BOF / **Dangerous?***




***A physical buffer overflow:
The Montparnasse
derailment of 1895***

Source:

“Secure Programming HOWTO”,

David Wheeler



Buffer Overflow (BOF)

A Simple BOF / Dangerous?



Okay, it crashed.

***So what?* ... you say**

**No danger, just a bug
to be fixed...**

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



Okay, it crashed.

So what? ... you say ...

No danger, just a bug to be fixed...

It IS DANGEROUS !!!

Why??



Buffer Overflow (BOF)

A Simple BOF / Dangerous?



Recall exactly *why* the process crashed (segfault-ed):

- The RETurn address was set to an incorrect/bogus/junk value (**0x41414141**)
- **Instead of just crashing the app, a hacker will carefully *craft* the RET address to a deliberate value – *code that (s)he wants executed!***
- **How exactly can this dangerous “*arbitrary code execution*” be achieved?**

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



Running the app like this:

```
$ perl -e 'print "A"x136' | ./getdata
```

which would cause it to “just” segfault.

But how about this:

```
$ perl -e 'print "A"x132 . "\x49\x8f\x04\x78"' | ./getdata
```

; where the address **0x498f0478** is a known location to code we want executed!

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload, or '**crafted buffer**' is:

payload = 128 A's + <SFP value> + <RET addr>
= AAAAA...AAAA + AAAA + 0x498f0478

- As seen, given a local buffer of 128 bytes, the *overflow* spills into the higher addresses of the stack
- In this case, the overflow is 4+4 bytes
- Which *overwrites* the
 - SFP (Saved Frame Pointer – essentially pointer to prev stack frame), and the
 - **RETurn address**, on the process stack
- ... thus causing control to be re-vectored to the RET address!
- *Thus, we have Arbitrary Code Execution (which could result in privilege escalation, a backdoor, etc)!*

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload or ‘crafted buffer’ can be used to deploy an attack in many forms:

- **Direct code execution:** executable machine code “injected” onto the stack, with the RET address arranged such that it points to this code
- **Indirect code execution:**
 - To internal program function(s)
 - To external library function(s)

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload or ‘crafted buffer’ can be deployed in many forms:

- Direct executable machine code “injected” onto the stack, with the RET address arranged such that it points to this code
 - What code?
 - Typically, a variation of the machine language for:
`setuid(0);`
`execl(“/bin/sh”, “sh”, (char *)0);`

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload or 'crafted buffer' can be deployed in many forms:

```
setuid(0);
```

```
execve("/bin/sh", argv, (char *)0);
```

- *In fact, no need to take the trouble to painstakingly build it, it's publicly available*
- Collection of shellcode!
<http://shell-storm.org/shellcode/>
 - Eg. 1 : `setuid(0); execve(/bin/sh,0)` for the IA-32:
<http://shell-storm.org/shellcode/files/shellcode-472.php>
- **Exploit-DB** (Offensive Security)
 - Or, use the Google Hacking Database ([GHDB](#), part of OffSec)

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload or 'crafted buffer' can be deployed in many forms:

- Eg. 2 (shell-storm):
adds a root user no-passwd to /etc/passwd
(84 bytes)

```
char shellcode[] =
```

```
"\x31\xc0\x31\xdb\x31\xc9\x53\x68\x73\x73\x77" "\x64\x68\x63\x2f\x70\x61\x  
x68\x2f\x2f\x65\x74"  
"\x89\xe3\x66\xb9\x01\x04\xb0\x05\xcd\x80\x89"  
"\xc3\x31\xc0\x31\xd2\x68\x6e\x2f\x73\x68\x68"  
"\x2f\x2f\x62\x69\x68\x3a\x3a\x2f\x3a\x68\x3a"  
"\x30\x3a\x30\x68\x62\x6f\x62\x3a\x89\xe1\xb2"  
"\x14\xb0\x04\xcd\x80\x31\xc0\xb0\x06\xcd\x80"  
"\x31\xc0\xb0\x01\xcd\x80";
```


Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload or 'crafted buffer' can be deployed in many forms:

- Indirect code execution:
 - To internal program function(s)
 - To external program function(s)
- Re-vector (forcibly change) the RET address such that control is vectored to an - **typically unexpected, out of the "normal" flow of control** - internal program function

<<

(Time permitting :-)

Demo of a BOF PoC on ARM Linux, showing precisely this

>>

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload or ‘crafted buffer’ can be deployed in many forms:

- Indirect code execution:
 - To internal program function(s)
 - To external library function(s)
- Revector (forcibly change) the RET address such control is vectored to an - typically unexpected, *out of the “normal” flow of control* - external library function

Buffer Overflow (BOF)

A Simple BOF / Dangerous?



The payload or ‘crafted buffer’ can be deployed in many forms:

- Re-vector (forcibly change) the RET address such that control is vectored to an - typically unexpected, out of the “normal” flow of control - external library function
- What if we re-vector control to a Std C Library (glibc) function:
 - Perhaps to, say, `system(char *cmd);`
 - Can setup the parameter (pointer to a command string) on the stack
 - *!!! Just think of the possibilities !!! - in effect, one can execute anything with the privilege of the hacked process*
 - *If root, then ... the system is compromised*
 - that’s pretty much exactly what the Ret2Libc hack / exploit is.

Modern OS Hardening Countermeasures



- A modern OS, like Linux, will / should implement a number of **countermeasures or “hardening” techniques** against vulnerabilities, and hence, potential exploits
- Reduces the attack surface
- ***Common Hardening Countermeasures include-***
 - 1) Using Managed Programming Languages
 - 2) Compiler Protection
 - 3) Library Protection
 - 4) Executable Space Protection
 - 5) [K]ASLR
 - 6) Better Testing

Modern OS Hardening Countermeasures

1) Using Managed Programming Languages

- *Programming in C/C++ is widespread and popular*
- *Pros- powerful, 'close to the metal', fast and effective code*
- *Cons-*
 - Programmer handles memory
 - Root Cause of many (most) memory bugs
 - Which lead to insecure exploitable software
- **A 'managed' language uses a framework (eg .NET) and/or a virtual machine construct**
- **Using a 'managed' language (Java, C#) greatly alleviates the burden of memory management from the programmer to the 'runtime'**
- **Reality -**
 - Many languages are *implemented* in C/C++
 - Real projects are usually a mix of managed and unmanaged code (eg. Android: Java @app layer + C/C++/JNI/DalvikVM @middleware + C/Assembly @kernel/drivers layers)
- **[Aside: is 'C' outdated? Nope; see the [TIOBE Index](#) for Programming languages]**

Modern OS Hardening Countermeasures



2) Compiler-level Protection

- Stack BOF Protection
 - Early implementations include
 - StackGuard (1997)
 - ProPolice (IBM, 2001)
 - GCC patches for stack-smashing protection
 - GCC
 - `-fstack-protector` flag (RedHat, 2005), and
 - `-fstack-protector-all` flag
 - `-fstack-protector-strong` flag (Google, 2012); gcc 4.9 onwards
 - Early in Android (1.5 onwards) – all Android binaries include this flag

Modern OS Hardening Countermeasures



2.1 Compiler-level Protection / Stack Protector GCC Flags

The **-fstack-protector-*<foo>*** gcc flags

From man gcc:

-fstack-protector

Emit extra code to check for buffer overflows, such as stack smashing attacks. This is done by adding a guard variable to functions with vulnerable objects. This includes functions that call "alloca", and functions with buffers larger than 8 bytes. The guards are initialized when a function is entered and then checked when the function exits. If a guard check fails, an error message is printed and the program exits.

-fstack-protector-all

Like **-fstack-protector** except that all functions are protected.

-fstack-protector-strong

Like **-fstack-protector** but includes additional functions to be protected --- those that have local array definitions, or have references to local frame addresses.

-fstack-protector-explicit

Like **-fstack-protector** but only protects those functions which have the "stack_protect" attribute.

Modern OS Hardening Countermeasures



2.1 Compiler-level Protection

- **From Wikipedia:**

“All [Fedora](#) packages are compiled with `-fstack-protector` since Fedora Core 5, and `-fstack-protector-strong` since Fedora 20.^{[19]cite_ref-20cite_ref-20[20]}

Most packages in [Ubuntu](#) are compiled with `-fstack-protector` since 6.10.^[21]

Every [Arch Linux](#) package is compiled with `-fstack-protector` since 2011.^[22]

All Arch Linux packages built since 4 May 2014 use `-fstack-protector-strong`.^[23]

Stack protection is only used for some packages in [Debian](#),^[24] and only for the [FreeBSD](#) base system since 8.0.^[25] ...”

- **How is the ‘`-fstack-protector<-xxx>`’ flag protection actually achieved?**
 - Typical stack frame layout:
[... local vars ...] [CTLI] [RET addr] ; where [CTLI] is control information (like the SFP)
 - A random value, called a **canary**, is placed by the compiler in the stack metadata, typically between the local variables and the RET address
 - **[... local vars ...] [canary] [CTLI] [RET addr]**

Modern OS Hardening Countermeasures



2.1 Compiler-level Protection

How is the ‘-fstack-protector<-xxx>’ flag protection actually achieved?
[contd.]

- Before a function returns, the canary is checked (by instructions inserted by the compiler into the function epilogue)

[... local vars ...] [**canary**] [CTLI] [RET addr]

- If the canary has changed, it's determined that an attack is underway (it might be an unintentional bug too), and the process is aborted (if this occurs in kernel-mode, the Linux kernel panics!)
- The overhead is considered minimal
- *[Exercise: try a BOF program. (Re)compile it with -fstack-protector gcc flag and retry (remember, requires >= gcc-4.9)]*

Modern OS Hardening Countermeasures



2.2 Compiler-level Protection

(Some) mitigation against a [format-string attack](#):

- Use the GCC flags **-Wformat-security** **-Werror=format-security**
- Realize that it's a GCC warning, nothing more
- From man gcc:
“-Wformat-security

If -Wformat is specified, also warn about uses of format functions that represent possible security problems. At present, this warns about calls to "printf" and "scanf" functions where the format string is not a string literal and there are no format arguments, **as in "printf (foo);". This may be a security hole if the format string came from untrusted input and contains %n."**

- Android
 - Oct 2008: disables use of “%n” format specifier (%n: init a var to # of chars printed before the %n specifier; can be used to set a variable to an arbitrary value)
 - 2.3 (Gingerbread) onwards uses the -Wformat-security and the -Werror=format-security GCC flags for all binaries

Modern OS Hardening Countermeasures



2.3 Compiler-level Protection

Using GCC `_FORTIFY_SOURCE`

- Lightweight protection against BOF in typical libc functions
- Requires GCC ver ≥ 4.0
- Provides **wrappers** around the following ‘typically dangerous’ functions:
 - `memcpy`, `mempcpy`, `memmove`, `memset`, `strcpy`, `stpcpy`, `strncpy`, `strcat`, `strncat`, `sprintf`, `vsprintf`, `snprintf`, `vsnprintf`, `gets`
- Must be used in conjunction with Optimization `[-On]` directive:
`-On -D_FORTIFY_SOURCE=n ; (n \geq 1)`
- Eg. `gcc prog.c -O2 -D_FORTIFY_SOURCE=2 -o prog -Wall <...>`
- [More details.](#)

Modern OS Hardening Countermeasures

2.4 Compiler-level Protection

- **RELRO – Read-Only Relocation**
 - Marks the program ELF binary headers Read-Only (RO) once symbol resolution is done at app launch
 - Thus any attack to change / redirect functions at run-time by modifying linkage is eclipsed
 - Achieved by compiling with the linker options:
 - Partial RELRO : `-Wl,-z,relro` : lazy-binding is still possible
 - Full RELRO : `-Wl,-z,relro,-z,now` : (process-specific) GOT and PLT marked RO as well, lazy-binding impossible
 - Android v4.4.1 onwards
- **Use checksec.sh!**

```
$ ./checksec.sh --file /bin/ps
RELRO      STACK CANARY      NX      PIE      RPATH      RUNPATH      FILE
Full RELRO Canary found  NX enabled  PIE enabled  No RPATH    No RUNPATH   /bin/ps
$ ./checksec.sh --file /opt/teamviewer/tv_bin/teamviewerd
RELRO      STACK CANARY      NX      PIE      RPATH      RUNPATH      FILE
Partial RELRO  Canary found  NX enabled  No PIE      No RPATH    No RUNPATH   /opt/teamviewer/tv_bin/teamviewerd
$
```


Modern OS Hardening Countermeasures

2.4 Compiler-level Protection

Compiler Instrumentation
– **Sanitizers** or UB
(Undefined Behavior)
Checkers (Google)

Dynamic Analysis

Run-time
instrumentation added by
GCC to programs to check
for UB and detect
programming errors.

<foo>Sanitizer : compiler instrumentation based family of tools ; where **<foo>** = *Address* | *Kernel* | *Thread* | *Leak* | *UndefinedBehavior*

Tool (click for documentation)	Purpose	Short Name	Environment Variable	Supported Platforms
AddressSanitizer	memory error detector	ASan	ASAN_OPTIONS [1]	x86, ARM, MIPS (32- and 64-bit of all), PowerPC64
KernelSanitizer		KASan	-	4.0 kernel: x86_64 only (and ARM64 from 4.4)
ThreadSanitizer	data race detector	TSan	TSAN_OPTIONS [2]	Linux x86_64 (tested on Ubuntu 12.04)
LeakSanitizer	memory leak detector	LSan	LSAN_OPTIONS [3]	Linux x86_64
UndefinedBehaviorSanitizer	undefined behavior detector	UBSan	-	i386/x86_64, ARM, Aarch64, PowerPC64, MIPS/MIPS64
UndefinedBehaviorSanitizer for Linux Kernel			-	Compiler: gcc 4.9.x; clang[+]

Modern OS Hardening Countermeasures



2.4 Compiler-level Protection

- **<foo>Sanitizer**
 - **Address Sanitizer (ASan)**
 - Kernel Sanitizer (KASAN)
 - Thread Sanitizer (TSan)
 - Leak Sanitizer
 - Undefined Behavior Sanitizer (UBSan)
 - UBSan for kernel
- **Enable by GCC switch : -fsanitize=<foo>**
; <foo>=[[kernel]-address | thread | leak | undefined]
- **Address Sanitizer (ASan)**
 - ASan: “a programming tool that detects memory corruption bugs such as buffer overflows or accesses to a dangling pointer (use-after-free). AddressSanitizer is based on compiler instrumentation and directly-mapped shadow memory. AddressSanitizer is currently implemented in Clang (starting from version 3.1[1]) and GCC (starting from version 4.8[2]). On average, the instrumentation increases processing time by about 73% and memory usage by 340%.[3]”
 - “**Address sanitizer is nothing short of amazing**; it does an excellent job at detecting nearly all buffer over-reads and over-writes (for global, stack, or heap values), use-after-free, and double-free. It can also detect use-after-return and memory leaks” - D Wheeler, “Heartbleed”
 - Usage: just compile with the GCC flag: **-fsanitize=address**

Modern OS Hardening Countermeasures



2.5 Compiler-level Protection

“The Stack is Back”, Jon Oberheide - a must-read slide deck!

Real-world stack overflow based exploits:

[CVE-2010-3848](#)

[CVE-2010-3850](#)

Econet-sendmsg: VLA (var len array) on the stack with attacker controlled length!

[Demo:

- Achieve adjacent kstacks
- Process #2 goes to sleep
- Stack overflow in process #1
- Overwrite return address on process #2 kernel stack
- Process #2 wakes up
- Process #2 returns to attacker control address
- Privilege escalation payload executed!

<https://jon.oberheide.org/files/half-nelson.c>

]

Mitigation: the new STACKLEAK feature!

“Trying to get STACKLEAK into the kernel”, LWN, Sept 2018

STACKLEAK merged in 4.20 (only in 5.0?) Aug 2018 [commit].

Modern OS Hardening Countermeasures

3) Libraries

- BOF exploits – how does one attack?
- By studying real running apps, looking for a weakness to exploit (enumeration)
 - f.e. the infamous libc gets() and similar functions in [g]libc!
- It's mostly by exploiting these common memory bugs that an exploit can be crafted and executed
- Thus, it's **Really Important** that we developers re-learn: **Must Avoid** using std lib functions which are not bounds-checked
 - gets, sprintf, strcpy, scanf, etc
 - Replace gets with fgets (or better still with getline / getdelim); similarly for snprintf, strncpy, sprintf, etc
 - s/strfoo/strnfoo

Modern OS Hardening Countermeasures

3) Libraries

- Best to make use of “safe” libraries, especially for string handling
- Obviously, a major goal is to prevent security vulnerabilities
- Examples include
 - [The Better String Library](#)
 - [Safe C Library](#)
 - [Simple Dynamic String library](#)
 - [Libsafe](#)
 - Also see:
[Ch 6 “Library Solutions in C/C++ Library Solutions in C/C++”, Secure Programming for UNIX and Linux HOWTO, D Wheeler](#)
- **Source** - Cisco Application Developer Security Guide

“... In recent years, web-based vulnerabilities have surpassed traditional buffer overflow attacks both in terms of absolute numbers as well as the rate of growth. The most common forms of web-based attacks, such as cross-site scripting (XSS) and SQL injection, can be mitigated with proper input validation.

Cisco strongly recommends that you incorporate the [Enterprise Security API \(ESAPI\) Toolkit](#) from the Open Web Application Security Project ([OWASP](#)) for input validation of web-based applications. ESAPI comes with a set of well-defined security API, along with ready-to-deploy reference implementations.”

Modern OS Hardening Countermeasures



4) Executable Space Protection

- The most common attack vector
 - Inject shellcode onto the stack (or heap), typically via a BOF vuln
 - Arrange to have the shellcode execute, thus gaining privilege (or a backdoor)
- Modern processors have the ability to ‘mark’ a page with an NX (No eXecute) bit
 - So if we ensure that all pages of data regions like the stack, heap, BSS, etc are marked as NX, then the shellcode holds no danger!
 - The typical BOF (‘stack smashing’) attack relies on memory being readable, writeable and executable (rwx)
- Key Principle: **W^X pages**
 - LSMs: opt-in feature of the kernel
 - LSMs do incorporate W^X mechanisms

Modern OS Hardening Countermeasures



4) Executable Space Protection

- Linux kernel
 - Supports the NX bit from v2.6.8 onwards
 - On processors that have the hardware capability
 - Includes x86, x86_64 and x86_64 running in 32-bit mode
 - x86_32 requires PAE to support NX
 - (However) For CPUs that do not natively support NX, 32-bit Linux has software that emulates the NX bit, thus protecting non-executable pages
 - Check for NX hardware support (on x86[_64] Linux):
`grep -w nx -q /proc/cpuinfo && echo "Yes" || echo "Nope"`
 - A commit by Kees Cook (v2.6.38) ensures that even if NX support is turned off in the BIOS, that is ignored by the OS and protection remains

Modern OS Hardening Countermeasures

4) Executable Space Protection

- Ref: https://en.wikipedia.org/wiki/NX_bit
- (More on) Processors supporting the NX bit
 - Intel markets it as XD (eXecute Disable) ; AMD as ‘EVP’ - Enhanced Virus Protection
 - MS calls it DEP (Data Execution Prevention) ; ARM as XN – eXecute Never
 - *Android: As of Android 2.3 and later, architectures which support it have non-executable pages by default, including non-executable stack and heap.[1][2][3]*
- ARMv6 onwards (new PTE format with XN bit) ; [PowerPC, Itanium, Alpha, SunSparc, etc, too support NX]
- Intel **SMEP** – Supervisor Mode Execution Prevention – bit in CR4 (ARM equivalent: PXN/PAN)
 - When set, when in Ring 0, MMU faults when trying to execute a page in Ring 3
 - Prevents the “usual” kernel exploit vector: map some shellcode in userland, exploit some kernel bug to overwrite kernel memory to point to it, and get it to trigger
 - PaX solves this via PAX_UDEREF
 - “SMEP: What is It, and How to Beat It on Linux”, Dan Rosenberg
- Intel **SMAP** – Supervisor Mode Access Prevention
 - When set, when in Ring 0, MMU faults when trying to access (r|w|x) a usermode page
 - In Linux since 3.8

Modern OS Hardening Countermeasures



5) ASLR – Address Space Layout Randomization

- NX (or DEP) protects a system by not allowing arbitrary code execution on non-text pages (stack/heap/data/BSS/etc, generically, on W^X pages)
- But it **cannot** protect against attacks that invoke *legal code* – like [g]libc functions, system calls (as they're in a valid text segment and are thus marked as r-x in their respective PTE entries)
- In fact, this is the attack vector for what is commonly called **Ret2Libc and ROP**-style (ROP = Return Oriented Programming) attacks
- How can *these* attacks be prevented (or at least mitigated)?
 - ASLR : by *randomizing* the layout of the process VAS (virtual address space), an attacker cannot know (or guess) in advance the location (virtual address) of glibc code, system call code, etc
 - Hence, attempting to launch this attack usually causes the process to just crash and the attack fails
 - ASLR in Linux from early on (2005; CONFIG_RANDOMIZE_BASE), and KASLR from 3.14 (2014); KASLR is only enabled *by default* in recent 4.12 Linux kernels.

Modern OS Hardening Countermeasures



5) ASLR – Address Space Layout Randomization

- Note though:
 - (K)ASLR is a *statistical* protection not an absolute one; it just adds an additional layer of difficulty (depending on the # of random bits available; curr 9 bits used) for an attacker, but does not inherently prevent attacks in any way
 - Also, even with full ASLR support, a particular process may *not* have it's VAS randomized
 - Why? Because this requires compile-time support (within the binary executable too): the binary must be built as a Position Independent Executable (PIE) [`-no-pie -mforce-no-pic -fno-stack-protector` etc]
- Process ASLR – turned On by compiling source with the `-fPIE` and `-pie` gcc flags
- **Information leakage** (for eg. a known kernel pointer value) can completely compromise the ASLR schema (example)
- Recent: a Perl script to detect 'leaking' kernel addresses added in 4.14 (commit; TC Harding)
 - leaking_addresses: add 32-bit support : commit 4.17-rc1 29 Jan 2018
Suggested-by: Kaiwan N Billimoria <kaiwan.billimoria@gmail.com> ☺
Signed-off-by: Tobin C. Harding <me@tobin.cc>

Modern OS Hardening Countermeasures



6) Better Testing

- Of course, most Q&A teams (as well as conscientious developers) will devise, implement and run an impressive array of test cases for the given product or project
- However, it's usually the case that most of these fall into the *positive* test-cases bracket (check that the test yields the desired outcome)
- This approach will typically fail to find bugs and vulnerabilities that an attacker probes for
 - We have to adopt an “attacker” mindset (*“set a thief to catch a thief”*)
 - We need to develop an impressive array of thorough **negative test-cases** (which check whether the program/device-under-test fails correctly and gracefully)

Modern OS Hardening Countermeasures



6) Better Testing

- A typical example:
the user is to pass a simple integer value:
 - Have test cases been written to check that it's within designated bounds?
 - Both positive and negative (integer overflow – IOF - bugs are heavily exploited! ;
[see SO: How is integer overflow exploitable?](#))
- From [OWASP](#): “Arithmetic operations cause a number to either grow too large to be represented in the number of bits allocated to it, or too small. This could cause a positive number to become negative or a negative number to become positive, resulting in unexpected/dangerous behavior.”
- Blog article: ['Integer Overflow', sploitfun](#)

Modern OS Hardening Countermeasures



6) Better Testing

- *Food for thought*

```
ptr = calloc(var_a*var_b, sizeof(int));
```

- What if it overflows??
 - Did you write a **validity check** for the parameter to calloc?
 - Old libc bug- an IOF could result in a much smaller buffer being allocated via calloc() ! (which could then be a good BOF attack candidate)
- Static analysis could / should catch a bug like this
- Dynamic analysis - take Valgrind's MemCheck tool (of course, Valgrind will only work on dynamic memory, not compile-time memory)
- The <foo>Sanitizer tools

Modern OS Hardening Countermeasures



6) Better Testing

- **IOF** (Integer Overflow)
 - Google wrote a *safe_iop* (*integer operations*) library for Android (from first rel)
 - However, as of Android 4.2.2, it appears to be used in a very limited fashion and is out-of-date too
- **Fuzzing**
 - “Fuzz testing or **fuzzing** is a software testing technique used to discover coding errors and security loopholes in software, operating systems or networks by inputting massive amounts of random data, called fuzz, to the system in an attempt to make it crash.” Source

Modern OS Hardening Countermeasures



6) *Better Testing*

- Mostly-positive testing is practically useless for security-testing
- Thorough Negative Testing is a MUST
- *Fuzzing*
 - especially effective in finding security-related bugs
 - Bugs that cause a program to crash (in the normal case)
 - Eg. Trinity – fuzzing tool used for kernel (syscall) testing

Concluding Remarks




- Experience shows that having several hardening techniques in place is *far superior* to having just one or two
- *Depth-of-Defense is critical*
- For example, take ASLR and NX (or XN):
 - Only NX, no ASLR: security bypassed via ROP-based attacks
 - Only ASLR, no NX: security bypassed via code injection techniques like stack-smashing, [or heap spraying](#)
 - Both full ASLR and NX: (more) difficult to bypass by an attacker

Concluding Remarks



Linux kernel – security patches into mainline

- Not so simple; the proverbial “tip of the iceberg”
- As far as security and hardening is concerned, projects like GRSecurity / PaX, KSP and OpenWall have shown what can be regarded as the “right” way forward
- A cool tool for checking kernel security / hardening status (from GRSec): **paxtest** (also lynis, checksec.sh, etc)
- However, the reality is that there continues to be resistance from the kernel community to merging in similar patchsets
- Why? Some legitimate reasons-
 - Info hiding - many apps / debuggers rely on pointers, information from procfs, sysfs, debugfs, etc
 - Debugging – breakpoints into code - don't work with NX on
 - Boot issues on some processors when NX used (being solved now)
- More info available: *Making attacks a little harder, LWN, Nov 2010*



Concluding Remarks



FYI :: Basic principle of attack

First, a program with an exploitable vulnerability – local or remote - must be found. This process is called *Reconnaissance / footprinting / enumeration*. (Dynamic approach- attackers will often ‘fuzz’ a program to determine how it behaves; static- use tools to disassemble/decompile (objdump,strings,IDA Pro,etc) the program and search for vulnerable patterns. Use vuln scanners).

[Quick Tip: Check out *nmap*, [Exploit-DB](#), the [GHDB](#) (Google Hacking Database) and the [Metasploit](#) pen(etration)-testing framework].

A **string containing shellcode** is passed as input to the vulnerable program. It overflows a buffer (**a BOF**), causing the shellcode to be executed (arbitrary code execution). The shellcode provides some means of access (a backdoor, or simply a direct shell) to the target system for the attacker. If *kernel* code paths can be revectorred to malicious code in userspace, gaining root is now trivial!

Stealth- the target system should be unaware it's been attacked (log cleaning, hiding).

Concluding Remarks

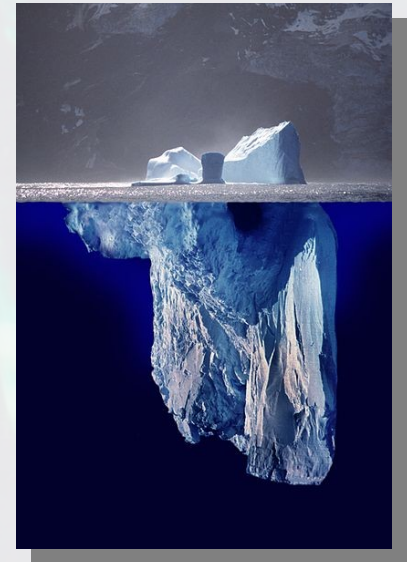


ADVANCED-

Defeat protections?

- **ROP (Return Oriented Programming)**
- **Defeats ASLR, NX**
 - Not completely; modern Linux PIE executables and library PIC code
- **Uses “gadgets” to alter and control PC execution flow**
- **A gadget is an existing piece of machine code that is leveraged to piece together a sequence of statements**
 - it's a non-linear programming technique!
 - Each gadget ends with a :
 - X86: 'ret'
 - RISC (ARM): `pop {rX, ..., pc}`
- **Sophisticated, harder to pull off**
- **But do-able!**

...



Questions?

Thank You!

kaiwanTECH
<http://kaiwantech.in>

Contact Info

Kaiwan N Billimoria

kaiwan@kaiwantech.com
kaiwan.billimoria@gmail.com

Author: “Hands-On System Programming with Linux: Explore Linux system programming interfaces, theory, and practice”, Packt, Oct 2018.

[Buy on Amazon.](#)

[LinkedIn public profile](#)

[My Tech Blog](#) [do follow!]

[GitHub page](#)

4931, Highpoint IV, 45 Palace Road, Bangalore 560001.

